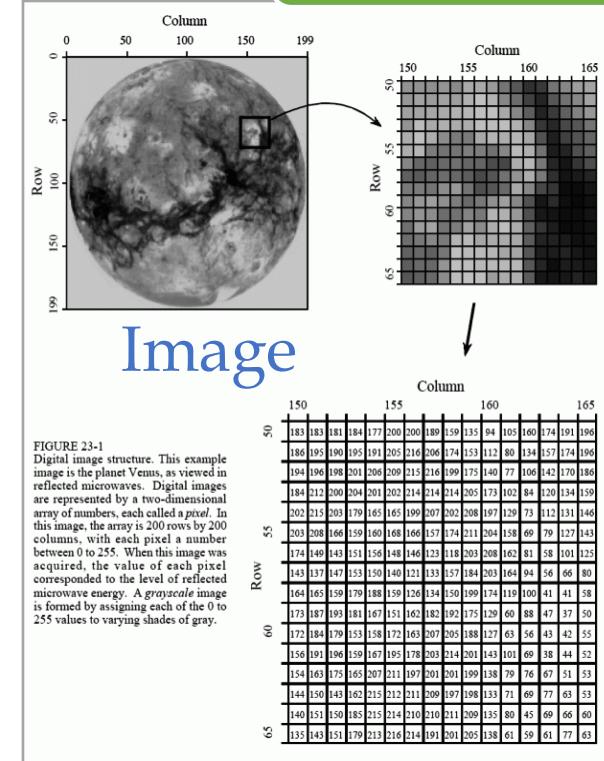
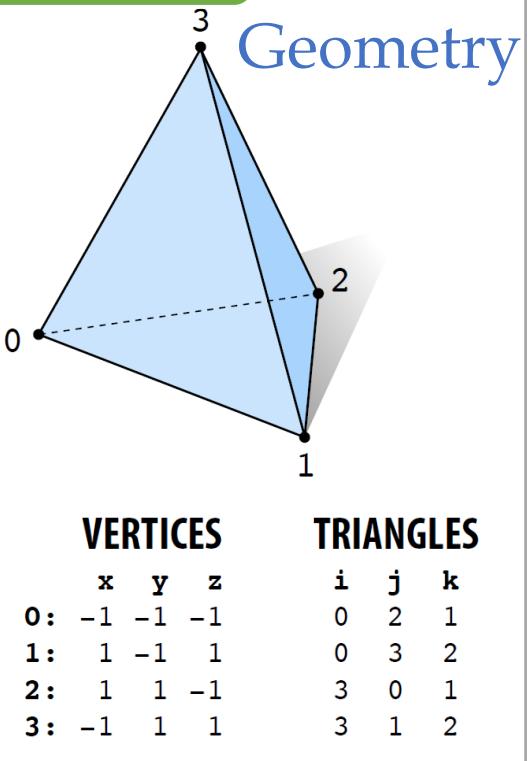


Geometry Reconstruction

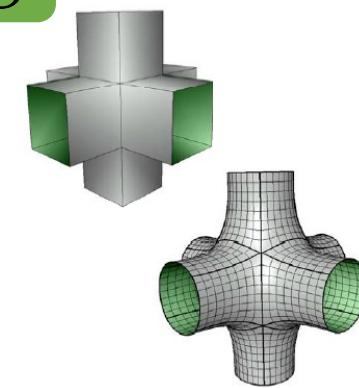
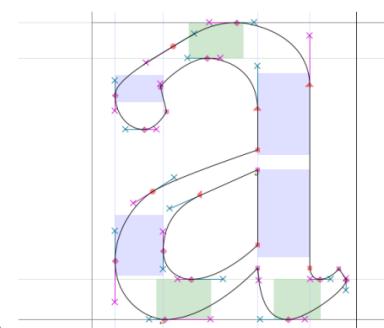
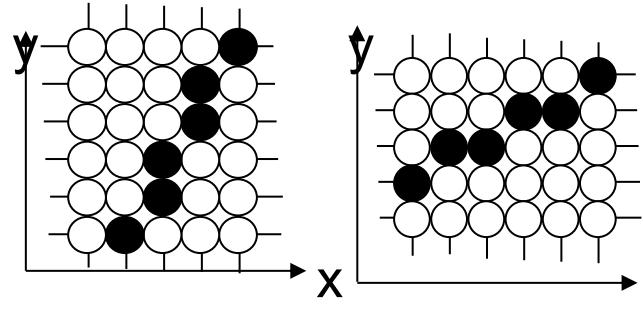
Representation



Image

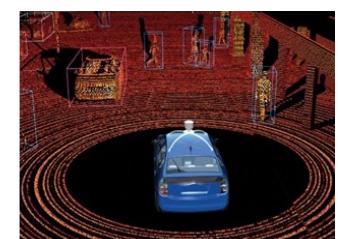
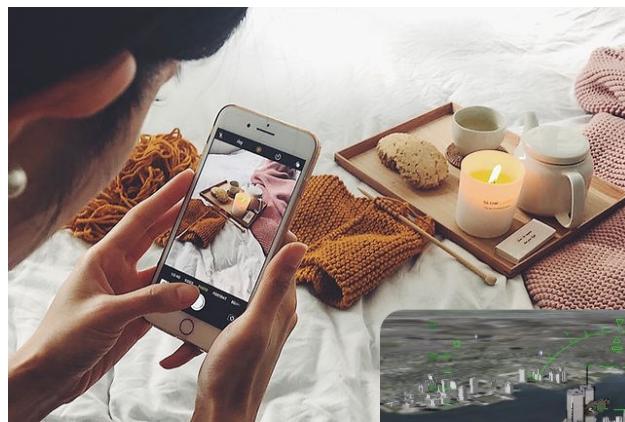


Editing



"Lotso Poses" by Daniel Arriaga
Toy Story 3, 2010
Pencil on paper

Geometry



Capture & Reconstruction

Reconstruction

- Data: 3D point clouds
 - From Lidars and RGB-D cameras
 - From computer vision algorithms such as triangulation, bundle adjustment, and deep learning

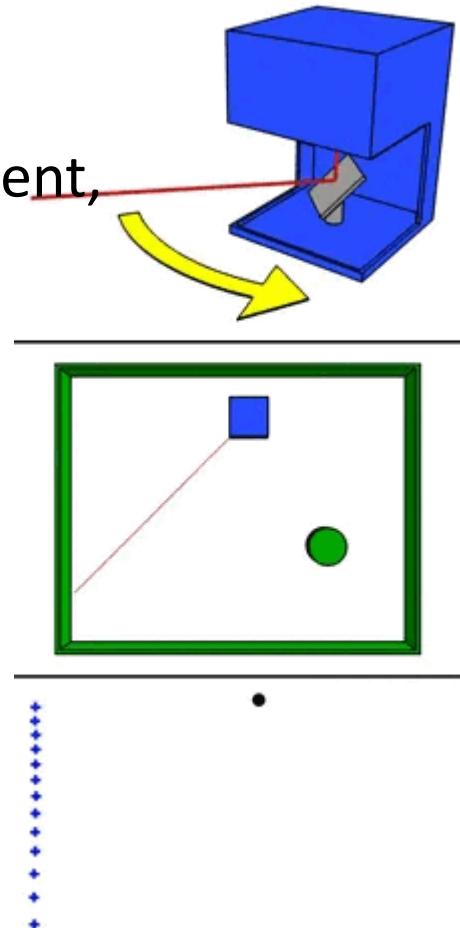
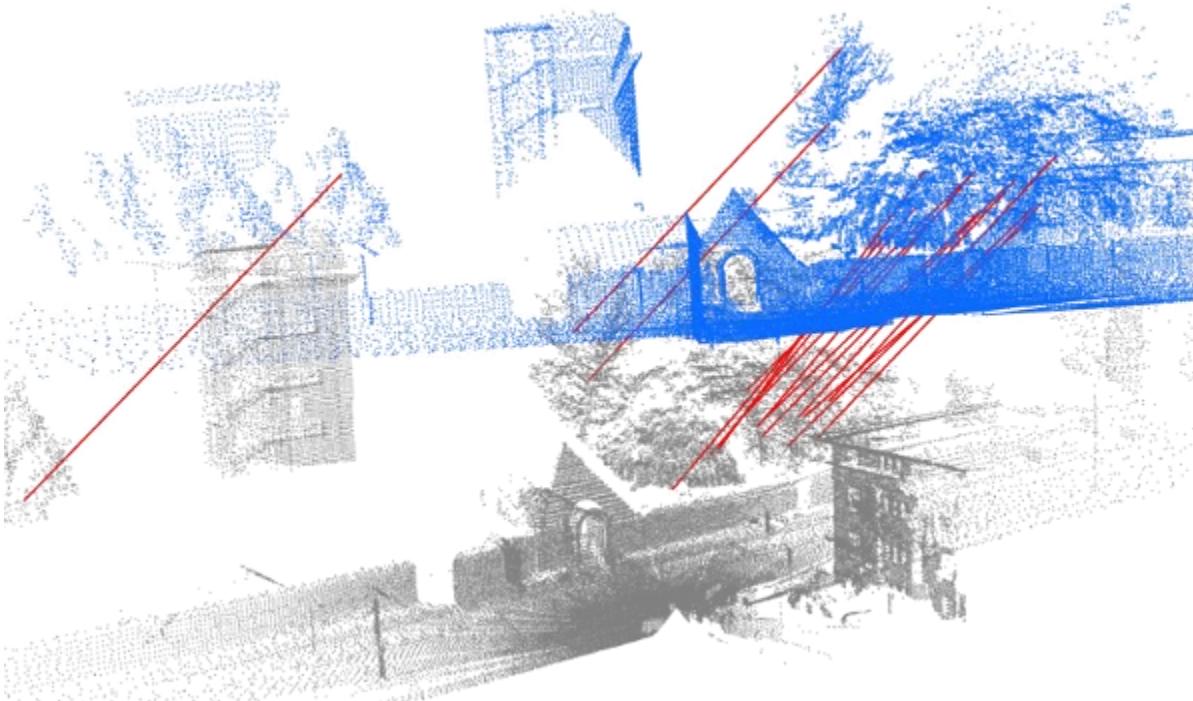


Figure: point clouds generate from buildings in Peking University.

Reconstruction

- After getting 3D point clouds
 - Registration



Figure

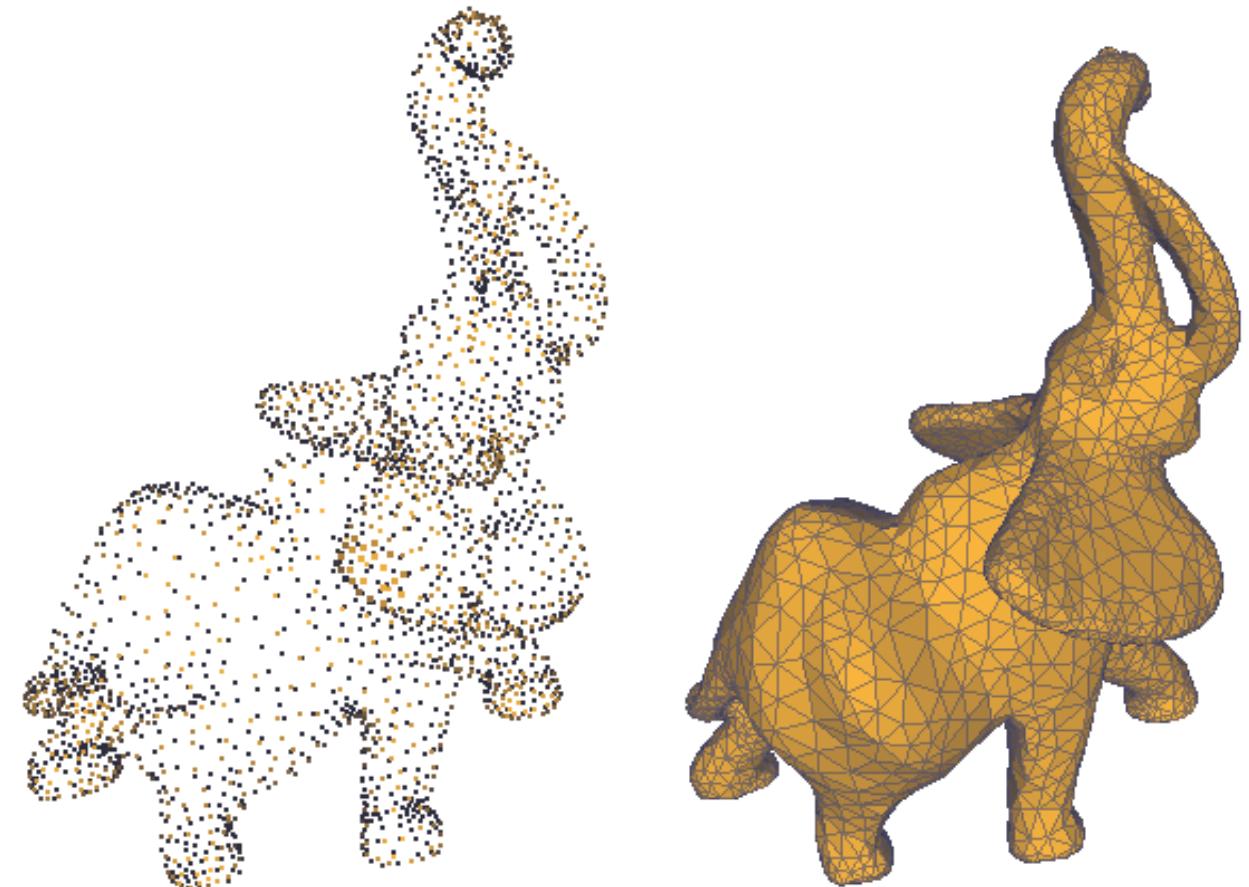
- (a) Left: data from two 3D scans of the same environment need to be aligned using point set registration.
(b) Right: data registered successfully using a variant of iterative closest point.



Courtesy: Martin Holzkothen, Michael Korn

Reconstruction

- After getting 3D point clouds
 - Surface Reconstruction
 - Motivation:
 - 1) Effective rendering of the mo
 - 2) Computational analysis
 - 3) other geometry processings: parameterization, morphing, blending etc..

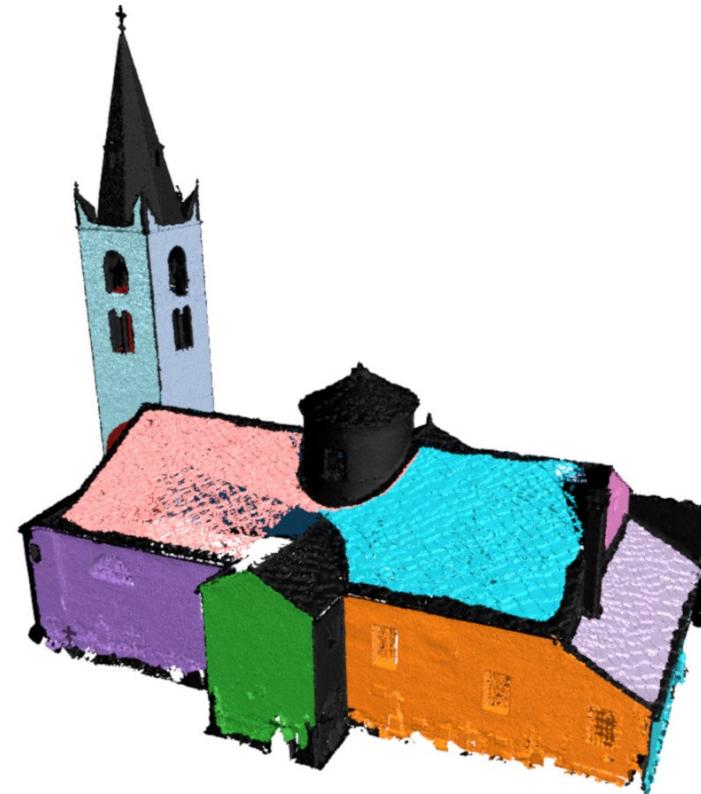


Courtesy: Jiju Peethambaran and Ramanathan Muthuganapathy

Figure: before and after surface reconstruction (triangulation).

Reconstruction

- After getting 3D point clouds
 - Model Fitting
 - Plane
 - Sphere
 - Cube
 - ...
 - Usually by RANSAC
 - RANDom SAmple Consensus



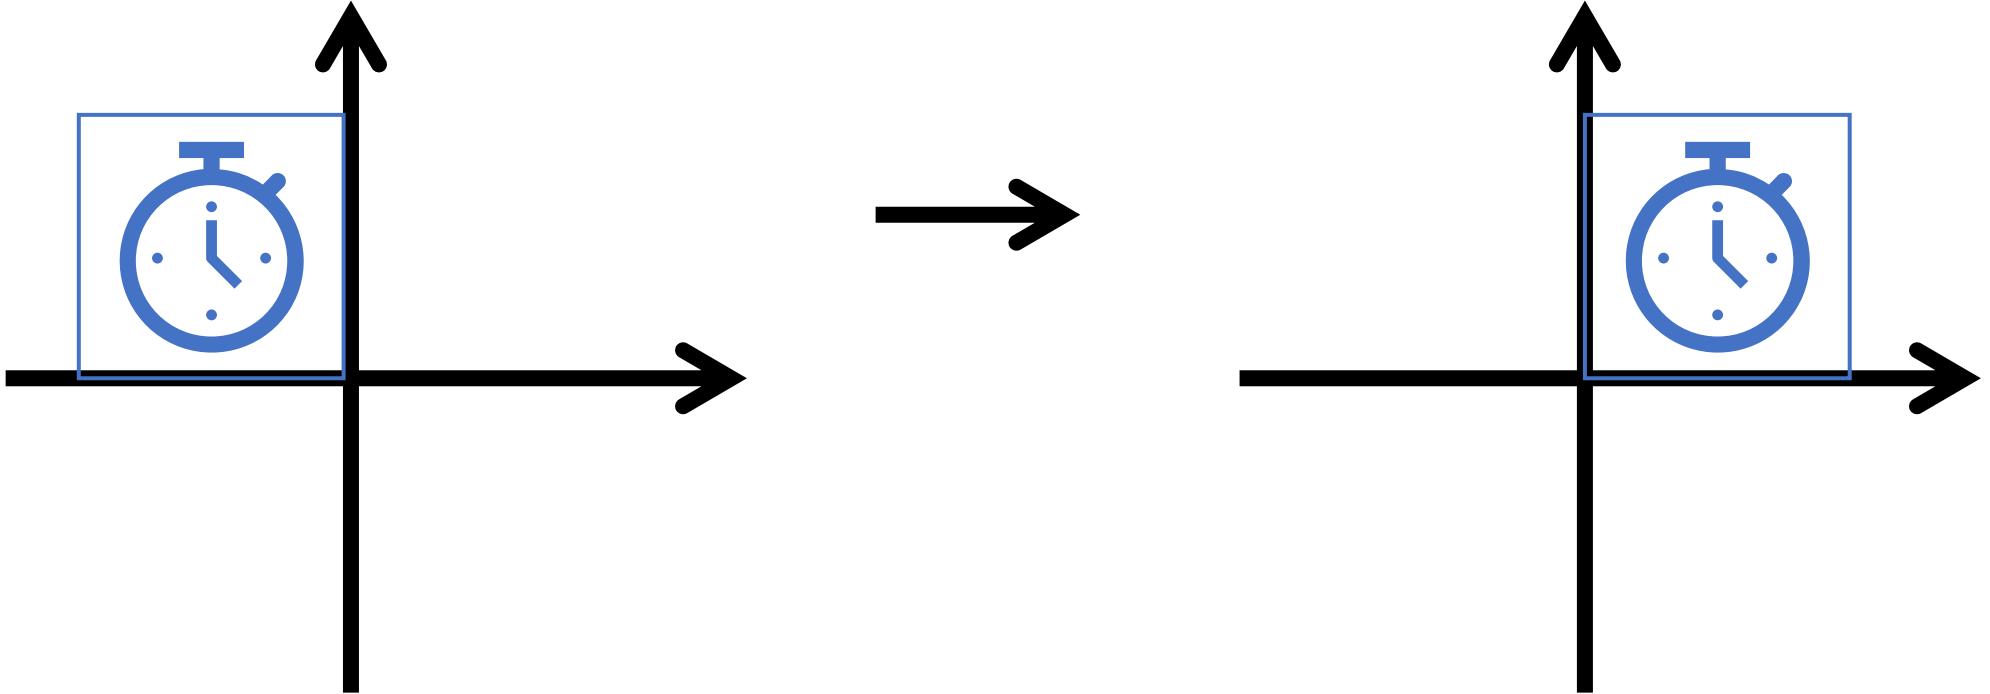
Courtesy: <https://github.com/STORM-IRIT/Plane-Detection-Point-Cloud>

Figure: the plane detection of the model of a building

Outline

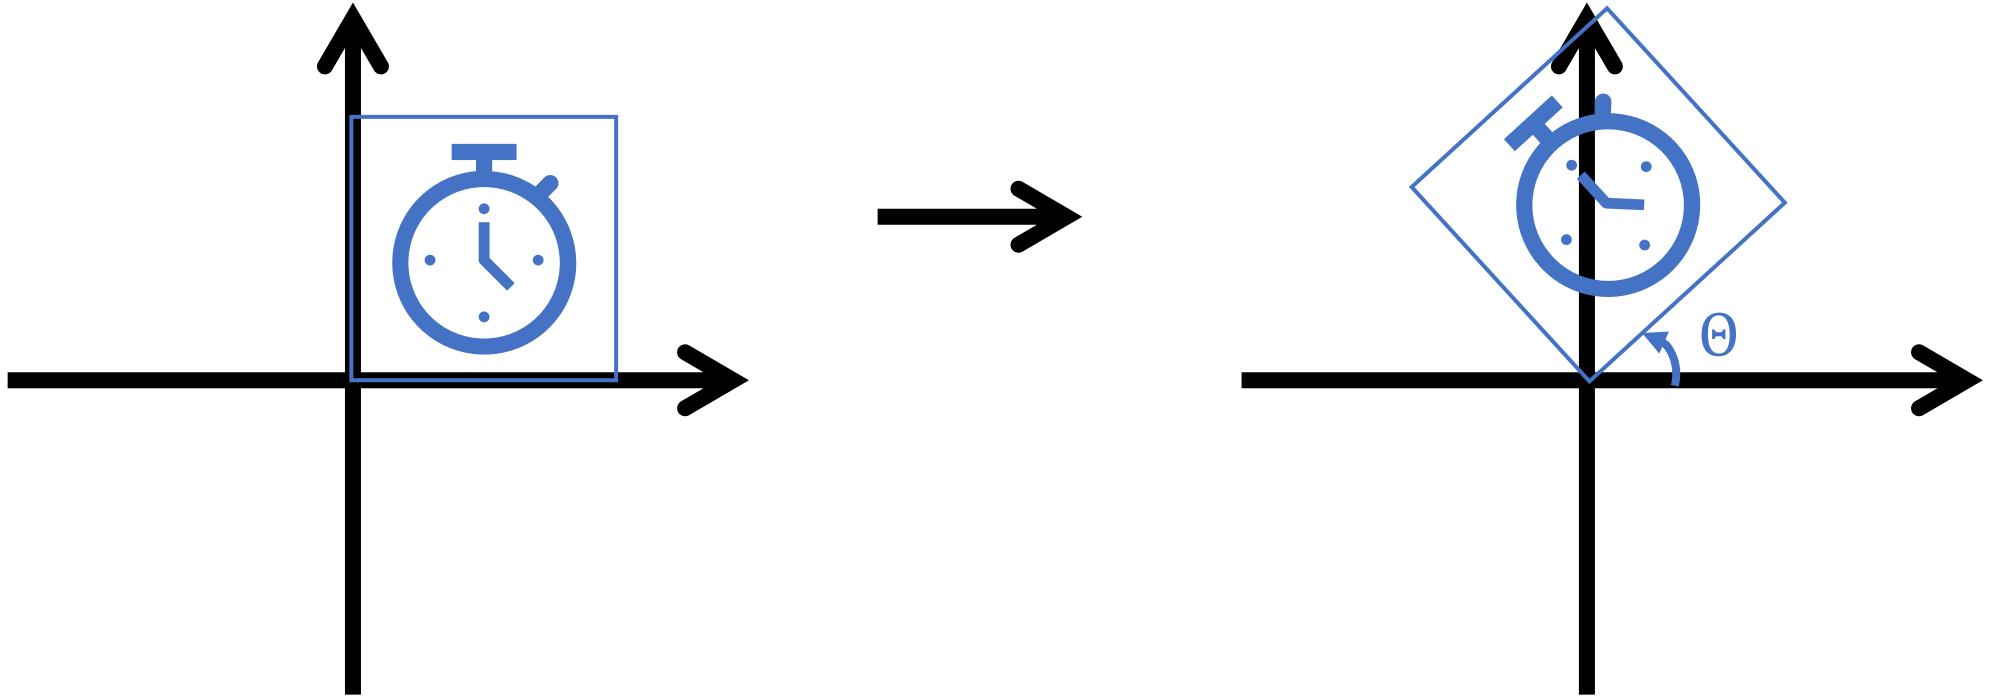
- Some Basics: Translation and Rotation
- Registration
 - Iterative Closest Point (ICP)
- Surface Reconstruction
 - Delaunay Triangulation
 - Poisson Surface Reconstruction
- Model Fitting
 - RANSAC

2D translation



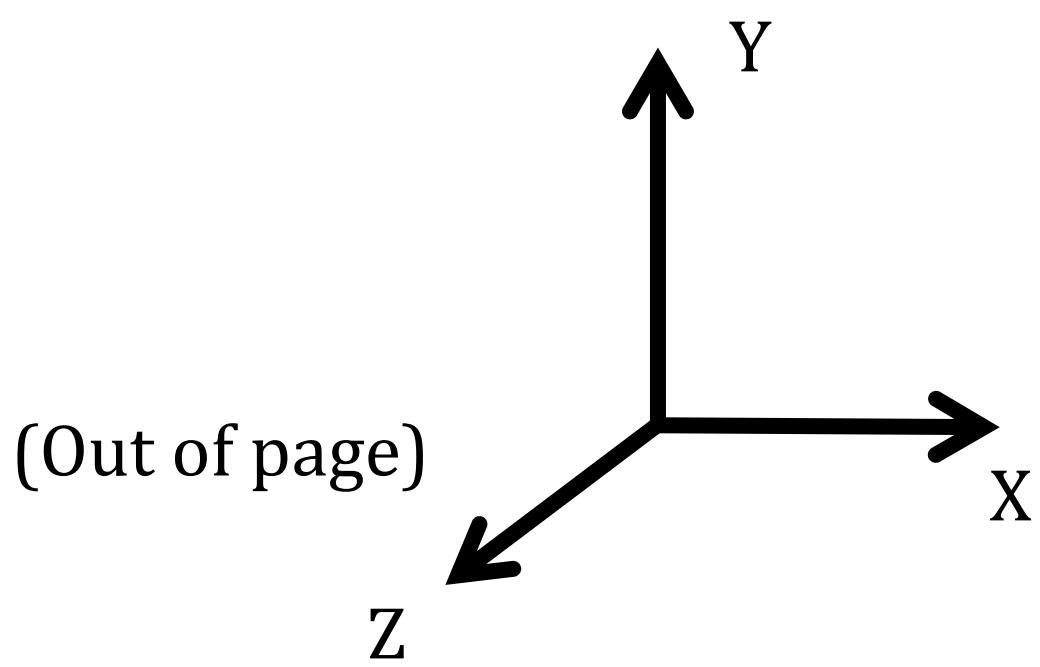
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} t_x \\ t_y \end{bmatrix} + \begin{bmatrix} x \\ y \end{bmatrix}$$

2D rotation

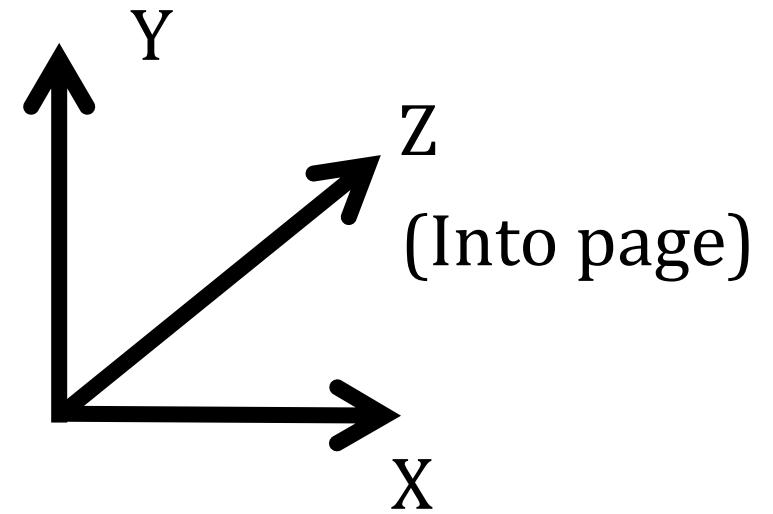


$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

From 2D to 3D



Right Hand

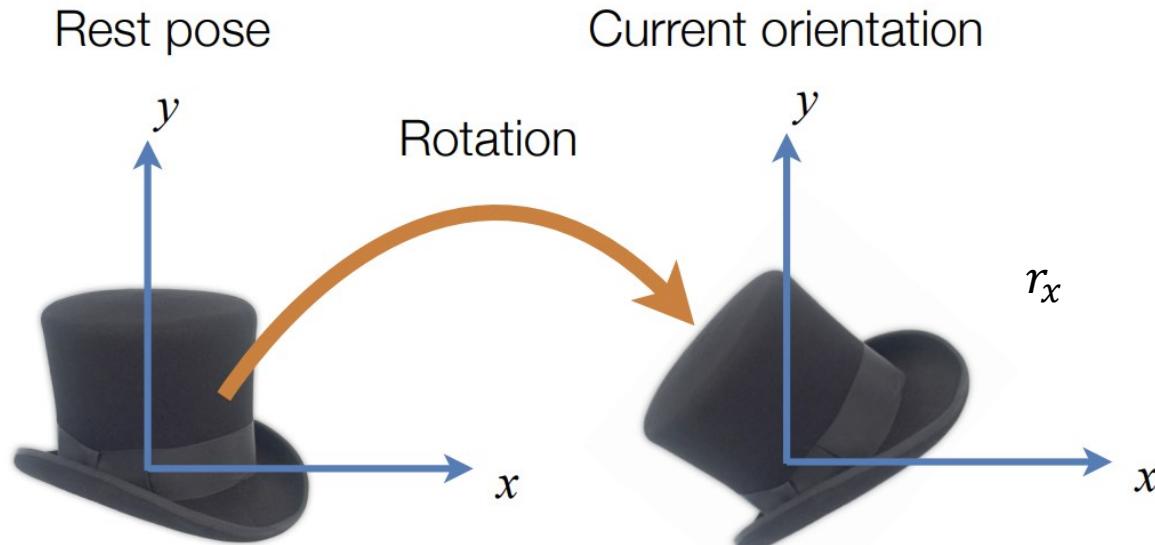


Left Hand

Z-axis determined from X and Y by cross product: $Z=X\times Y$

3D Rotation

- Rotation Matrix

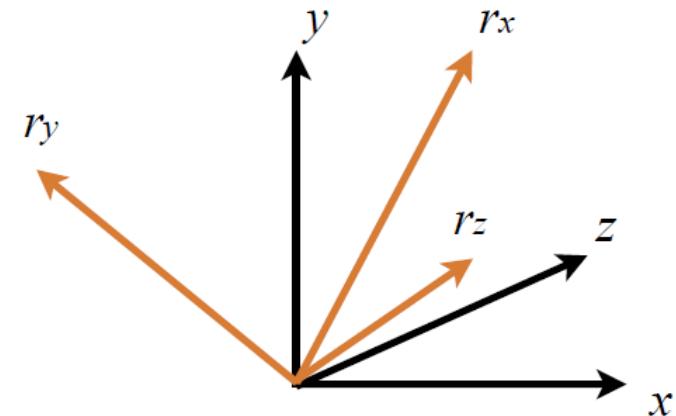


r_x, r_y, r_z are normalized orthogonal vectors

$$\mathbf{R}^{-1} = \mathbf{R}^T$$

Hint: it is easy to prove that $\mathbf{R}\mathbf{R}^T = \mathbf{I}$

$$\begin{aligned}\mathbf{p}' &= \mathbf{R}\mathbf{p} = \begin{bmatrix} r_{x0} & r_{y0} & r_{z0} \\ r_{x1} & r_{y1} & r_{z1} \\ r_{x2} & r_{y2} & r_{z2} \end{bmatrix} \begin{pmatrix} p_x \\ p_y \\ p_z \end{pmatrix} \\ &= r_x p_x + r_y p_y + r_z p_z\end{aligned}$$



Axis and Angles & Rotation Matrix

- Around Cartesian Axis

2D:

$$\text{Rot}_\alpha = \begin{pmatrix} \cos\alpha & -\sin\alpha \\ \sin\alpha & \cos\alpha \end{pmatrix}$$

3D:

$$R_{x,\phi} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \cos(\phi) & -\sin(\phi) \\ 0 & \sin(\phi) & \cos(\phi) \end{pmatrix}$$

$$R_{y,\phi} = \begin{pmatrix} \cos(\phi) & 0 & \sin(\phi) \\ 0 & 1 & 0 \\ -\sin(\phi) & 0 & \cos(\phi) \end{pmatrix}$$

$$R_{z,\phi} = \begin{pmatrix} \cos(\phi) & -\sin(\phi) & 0 \\ \sin(\phi) & \cos(\phi) & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

Axis and Angles & Rotation Matrix

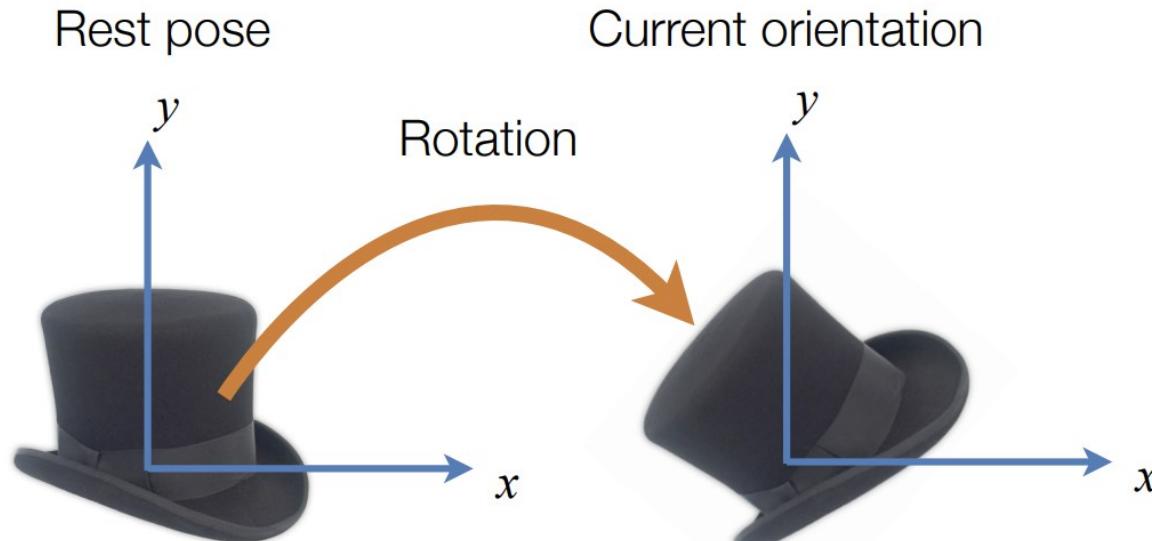
- Around Arbitrary Axis

$$R_{a,\phi} = \begin{pmatrix} c + (1 - c)a_x^2 & (1 - c)a_xa_y - sa_z & (1 - c)a_xa_z - sa_y \\ (1 - c)a_xa_y - sa_z & c + (1 - c)a_y^2 & (1 - c)a_ya_z - sa_x \\ (1 - c)a_xa_z - sa_y & (1 - c)a_ya_z - sa_x & c + (1 - c)a_z^2 \end{pmatrix}$$

with $c = \cos(\phi)$, and $s = \sin(\phi)$

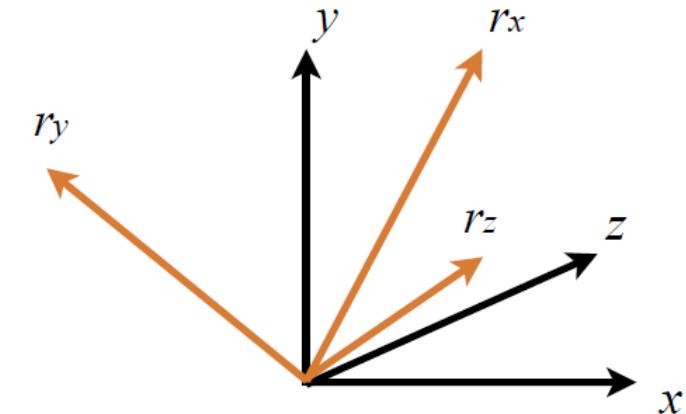
Inverse Matrix of a Rotation Matrix

- Rotation Matrix



$$\mathbf{R}^{-1} = \mathbf{R}^T$$

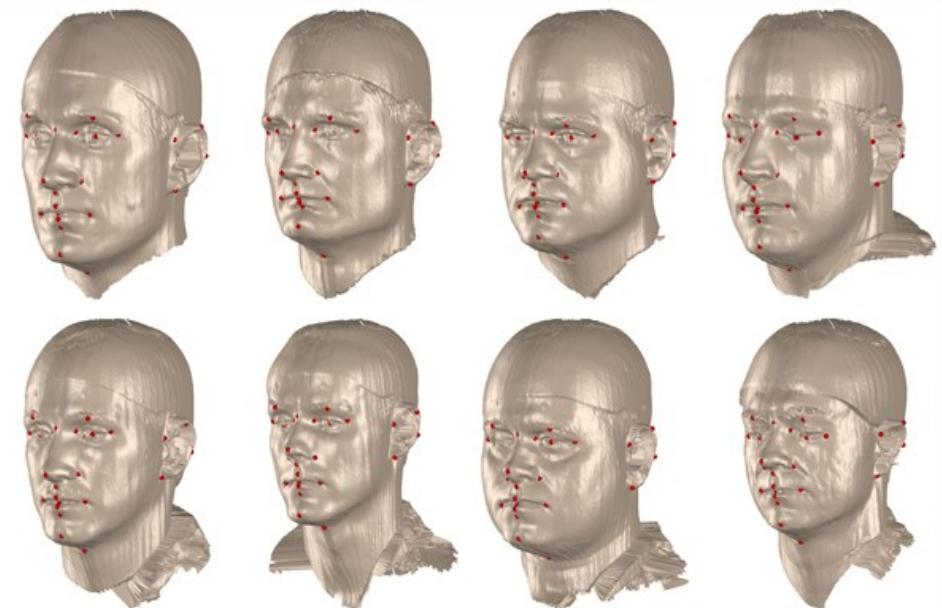
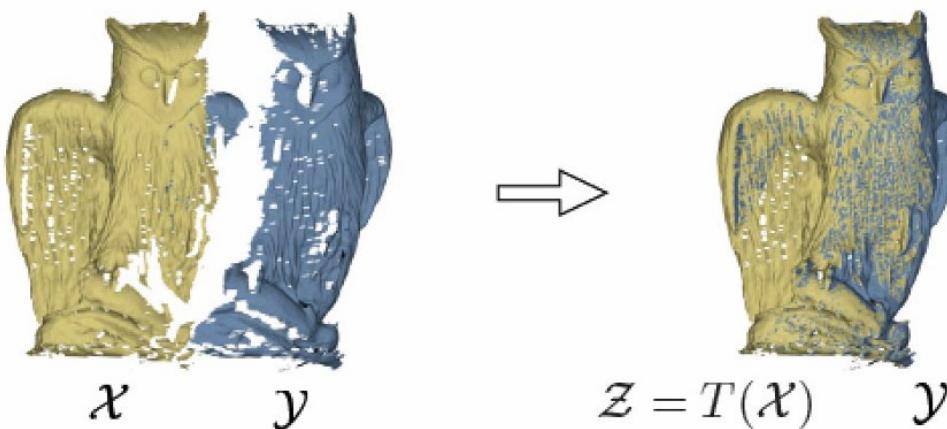
$$\begin{aligned}\mathbf{p}' &= \mathbf{R}\mathbf{p} = \begin{bmatrix} r_{x0} & r_{y0} & r_{z0} \\ r_{x1} & r_{y1} & r_{z1} \\ r_{x2} & r_{y2} & r_{z2} \end{bmatrix} \begin{pmatrix} p_x \\ p_y \\ p_z \end{pmatrix} \\ &= r_x p_x + r_y p_y + r_z p_z\end{aligned}$$



Registration

- Definition: transform one model to another based on their partially overlapping features

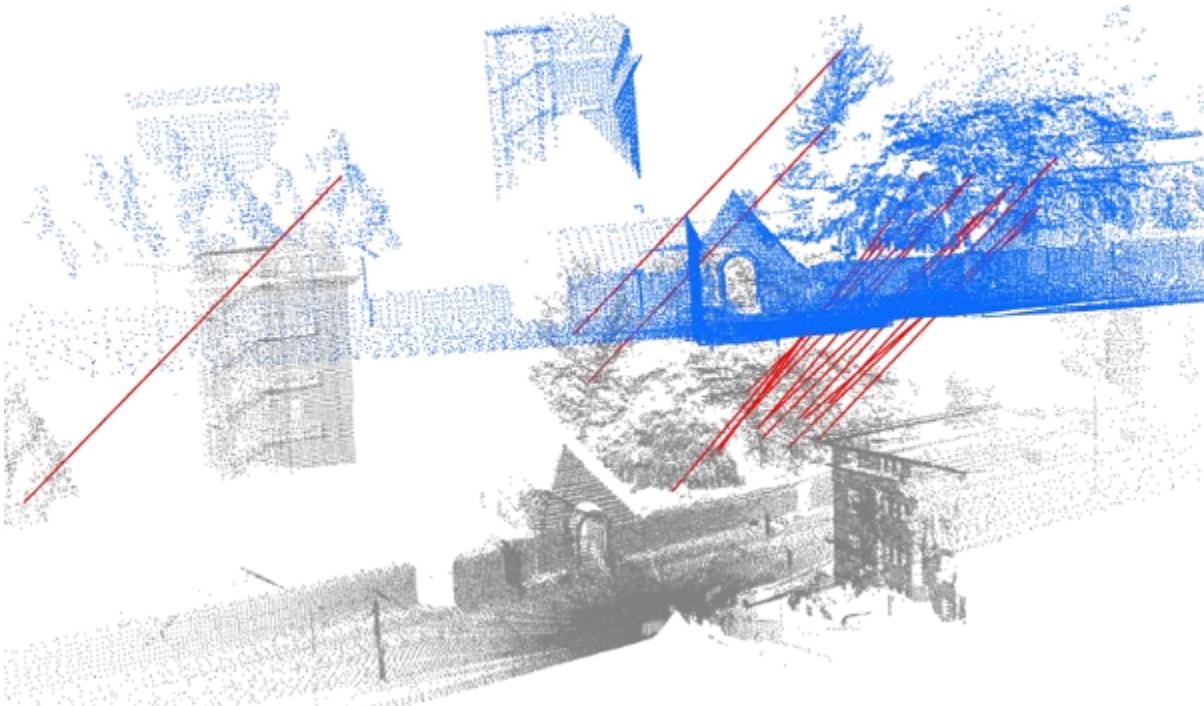
- Automated annotation
- Tracking and motion analysis
- Shape and data comparison



Courtesy: Schneider and Eisert

Registration

- We need do alignment.
 - Alignment: find a correct transformation

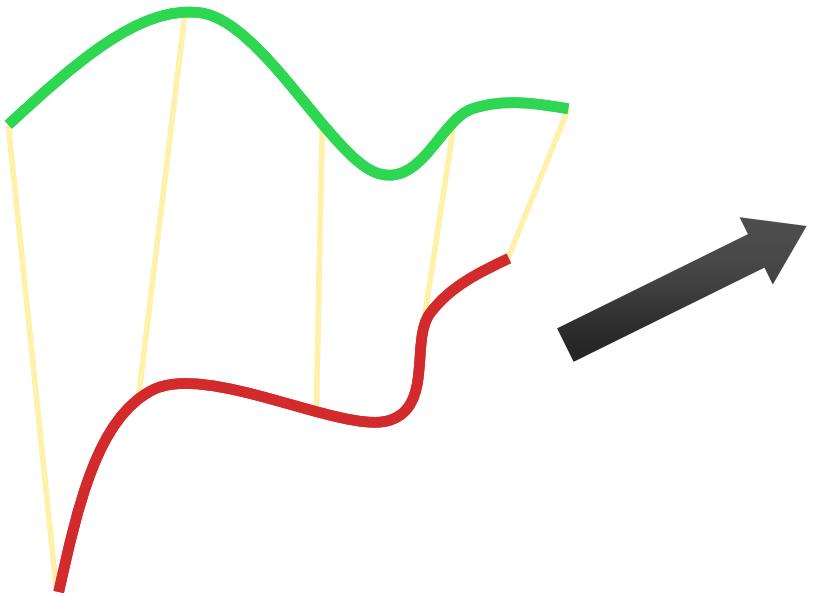


Courtesy: Martin Holzkothen, Michael Korn

- Figure
- (a) Left: data from two 3D scans of the same environment need to be aligned using point set registration.
 - (b) Right: data registered successfully using a variant of iterative closest point.

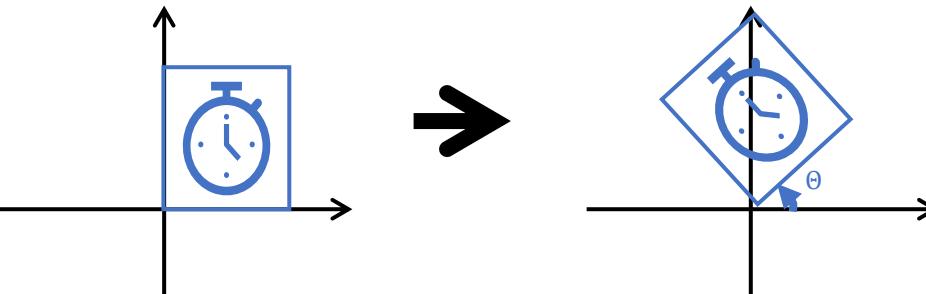
Alignment

- If rotation/translation are known, we can easily find correct correspondences



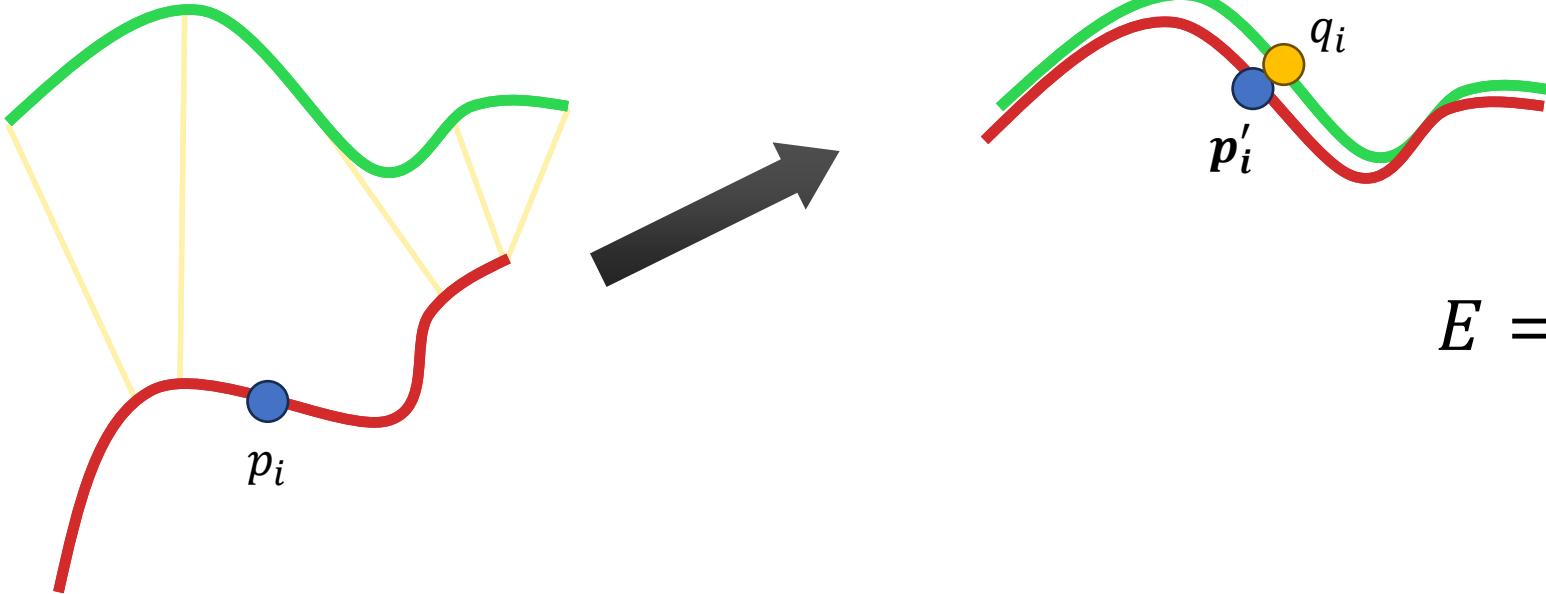
$$p'_i = Rp_i + t$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$



Alignment

- How to find correspondences:
 - User input? Feature detection? Signatures?
- ICP: Iterative Closest Point
assume closest points correspond

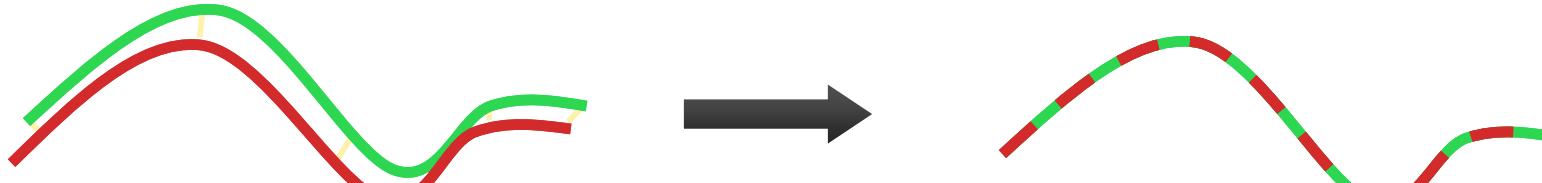


$$p'_i = R^* p_i + t^*$$

$$E = \sum |R^* p_i + t^* - q_i|^2$$

Alignment

- ... and iterate to find alignment
 - Iterative Closest Points (ICP) [Besl & McKay 92]
 - Converges if starting position “close enough”



ICP

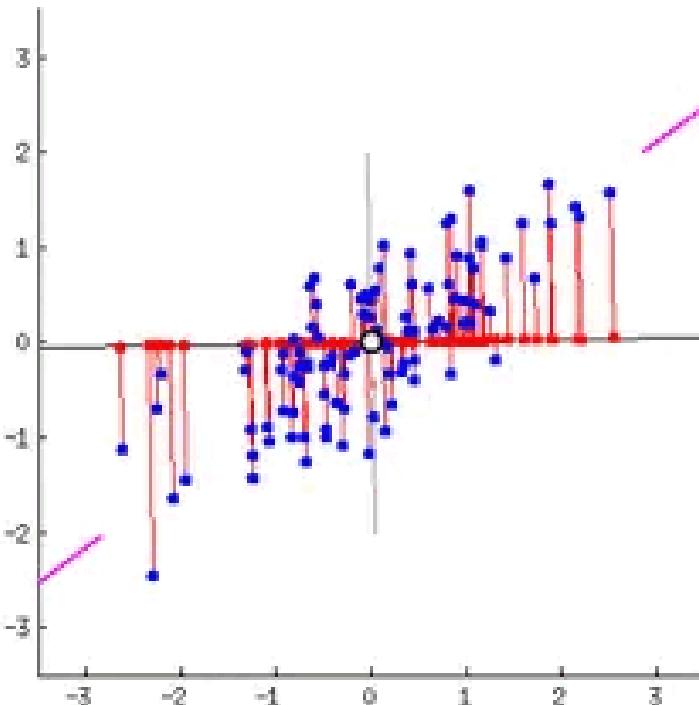
- 1. **Initialize** transformation R, t by **PCA** (Principle Component Analysis)
- 2. **Match** each $p'_i = Rp_i + t$ to closest point q_i on other scan
- 3. **Reject** pairs with distance $> k$ times median (outliers)
- 4. Construct **error function**:

$$E = \sum |Rp_i + t - q_i|^2$$

- 5. **Minimize** the error by **SVD**
- 6. **Loop** step 2-5 until the error is small enough

ICP: Initialize by PCA

- PCA: Principal Component Analysis
 - It can be used to compute axes
- Consider a set of points p_1, \dots, p_n with centroid location c



Find the orthogonal axes
Each axis is a “Principal Component”
Principal Component:

- direction with the largest variance,
- direction with the second-largest variance,
- ...
- direction with the least variance

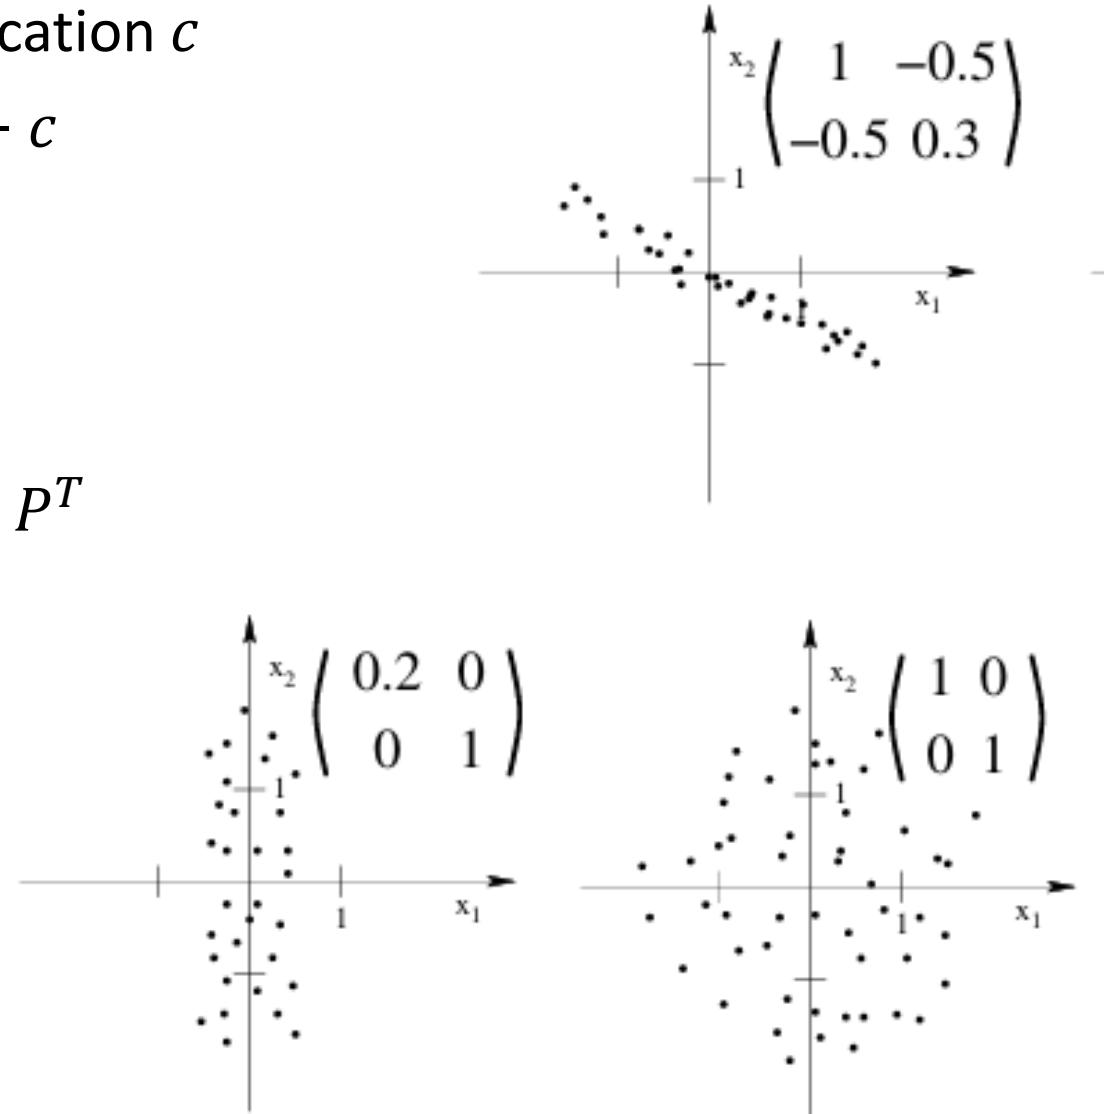
PCA by diagonalizing the covariance matrix

- Consider a set of points p_1, \dots, p_n with centroid location c
- Let P be a matrix whose i-th column is vector $p_i - c$

$$P = \begin{pmatrix} p_{1x} & c_x & p_{2x} & c_x & \dots & p_{nx} & c_x \\ p_{1y} & c_y & p_{2y} & c_y & \dots & p_{ny} & c_y \\ p_{1z} & c_z & p_{2z} & c_z & \dots & p_{nz} & c_z \end{pmatrix}$$

- Build the covariance matrix (协方差矩阵) : $M = P \times P^T$

$$\text{cov}(X) = \frac{1}{m} XX^T = \begin{bmatrix} \text{cov}(x, x) & \text{cov}(x, y) & \text{cov}(x, z) \\ \text{cov}(y, x) & \text{cov}(y, y) & \text{cov}(y, z) \\ \text{cov}(z, x) & \text{cov}(z, y) & \text{cov}(z, z) \end{bmatrix}$$



PCA by diagonalizing the covariance matrix

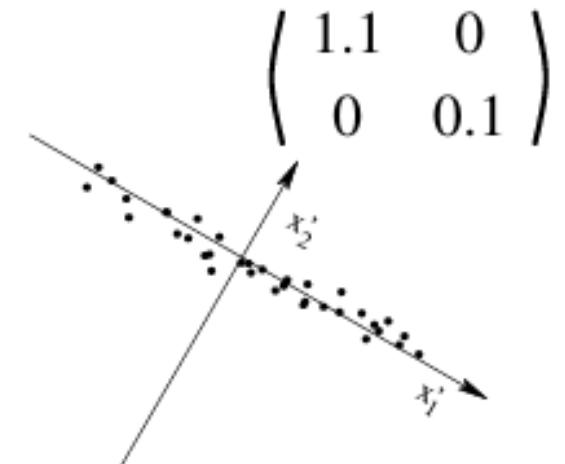
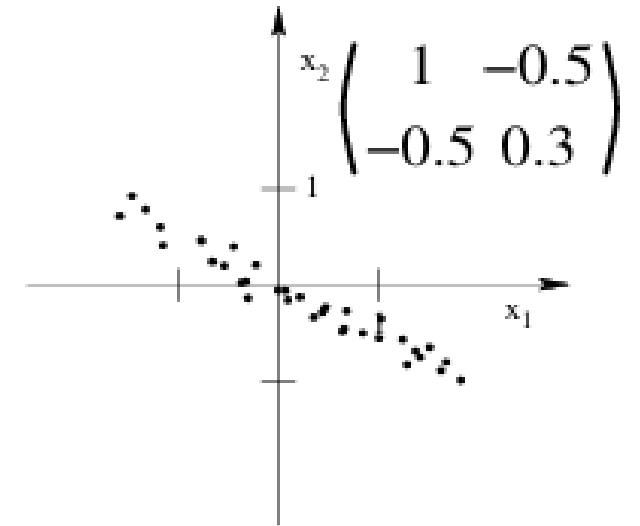
- Consider a set of points p_1, \dots, p_n with centroid location c
- Let P be a matrix whose i-th column is vector $p_i - c$

$$P = \begin{pmatrix} p_{1x} & c_x & p_{2x} & c_x & \dots & p_{nx} & c_x \\ p_{1y} & c_y & p_{2y} & c_y & \dots & p_{ny} & c_y \\ p_{1z} & c_z & p_{2z} & c_z & \dots & p_{nz} & c_z \end{pmatrix}$$

- Build the covariance matrix (协方差矩阵) : $M = P \times P^T$

$$\text{cov}(X) = \frac{1}{m} XX^T = \begin{bmatrix} \text{cov}(x, x) & \text{cov}(x, y) & \text{cov}(x, z) \\ \text{cov}(y, x) & \text{cov}(y, y) & \text{cov}(y, z) \\ \text{cov}(z, x) & \text{cov}(z, y) & \text{cov}(z, z) \end{bmatrix}$$

- The diagonalization of M:



PCA by diagonalizing the covariance matrix

- Consider a set of points p_1, \dots, p_n with centroid location c
- Let P be a matrix whose i-th column is vector $p_i - c$

$$P = \begin{pmatrix} p_{1x} & c_x & p_{2x} & c_x & \dots & p_{nx} & c_x \\ p_{1y} & c_y & p_{2y} & c_y & \dots & p_{ny} & c_y \\ p_{1z} & c_z & p_{2z} & c_z & \dots & p_{nz} & c_z \end{pmatrix}$$

- Build the covariance matrix (协方差矩阵) : $M = P \times P^T$

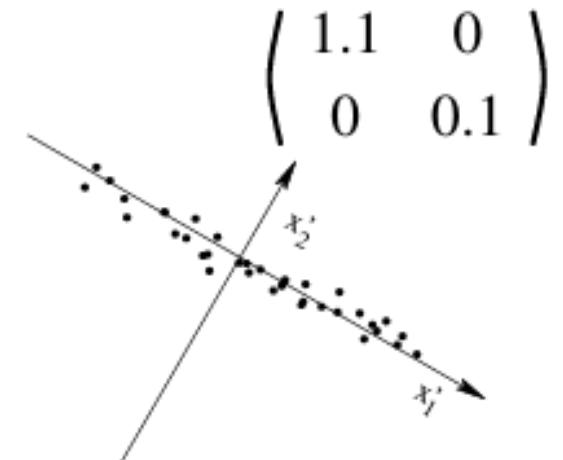
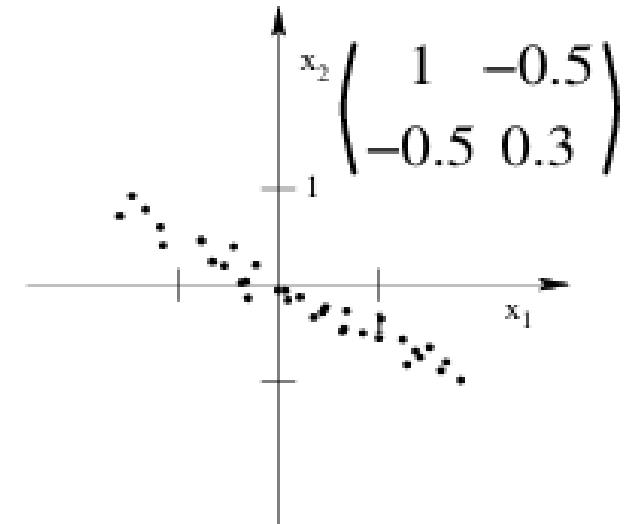
$$\text{cov}(X) = \frac{1}{m} XX^T = \begin{bmatrix} \text{cov}(x, x) & \text{cov}(x, y) & \text{cov}(x, z) \\ \text{cov}(y, x) & \text{cov}(y, y) & \text{cov}(y, z) \\ \text{cov}(z, x) & \text{cov}(z, y) & \text{cov}(z, z) \end{bmatrix}$$

- The diagonalization of M:

$$M = V \Lambda V^T$$

V is the matrix composed of the eigenvectors.

Λ is a diagonal matrix containing the corresponding eigenvalues.



PCA by diagonalizing the covariance matrix

Expressing Variance

First, we can express the variance of the projection in terms of the covariance matrix. The covariance matrix Σ of the dataset \mathbf{X} is given by:

$$\Sigma = \frac{1}{n}(\mathbf{X}^T \mathbf{X})$$

The variance of the projected data can be rewritten as:

$$\text{Var}(\mathbf{X}\mathbf{w}) = \mathbf{w}^T \Sigma \mathbf{w}$$

Using Lagrange Multipliers

We need to maximize $\mathbf{w}^T \Sigma \mathbf{w}$ subject to the constraint that \mathbf{w} is a unit vector:

$$\mathbf{w}^T \mathbf{w} = 1$$

We use Lagrange multipliers, introducing a multiplier λ :

$$\mathcal{L}(\mathbf{w}, \lambda) = \mathbf{w}^T \Sigma \mathbf{w} - \lambda(\mathbf{w}^T \mathbf{w} - 1)$$

Conclusion

Eigenvectors of the covariance matrix represent the directions of maximum variance in the data, while the corresponding eigenvalues indicate the variance in those directions, from the highest to the lowest.

Taking Derivatives

To find the extrema, we take the derivative of \mathcal{L} with respect to \mathbf{w} and λ , and set them to zero:

- Derivative with respect to \mathbf{w} :

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = 2\Sigma\mathbf{w} - 2\lambda\mathbf{w} = 0$$

This simplifies to:

$$\Sigma\mathbf{w} = \lambda\mathbf{w}$$

- Derivative with respect to λ :

$$\frac{\partial \mathcal{L}}{\partial \lambda} = -(\mathbf{w}^T \mathbf{w} - 1) = 0$$

Eigenvalue Problem

The equation $\Sigma\mathbf{w} = \lambda\mathbf{w}$ shows that \mathbf{w} is an eigenvector of the covariance matrix Σ , and λ is the corresponding eigenvalue.

ICP: Initialize by PCA

- PCA: Principal Component Analysis
 - It can be used to compute axes
- Consider a set of points p_1, \dots, p_n with centroid location c
- Let P be a matrix whose i-th column is vector $p_i - c$

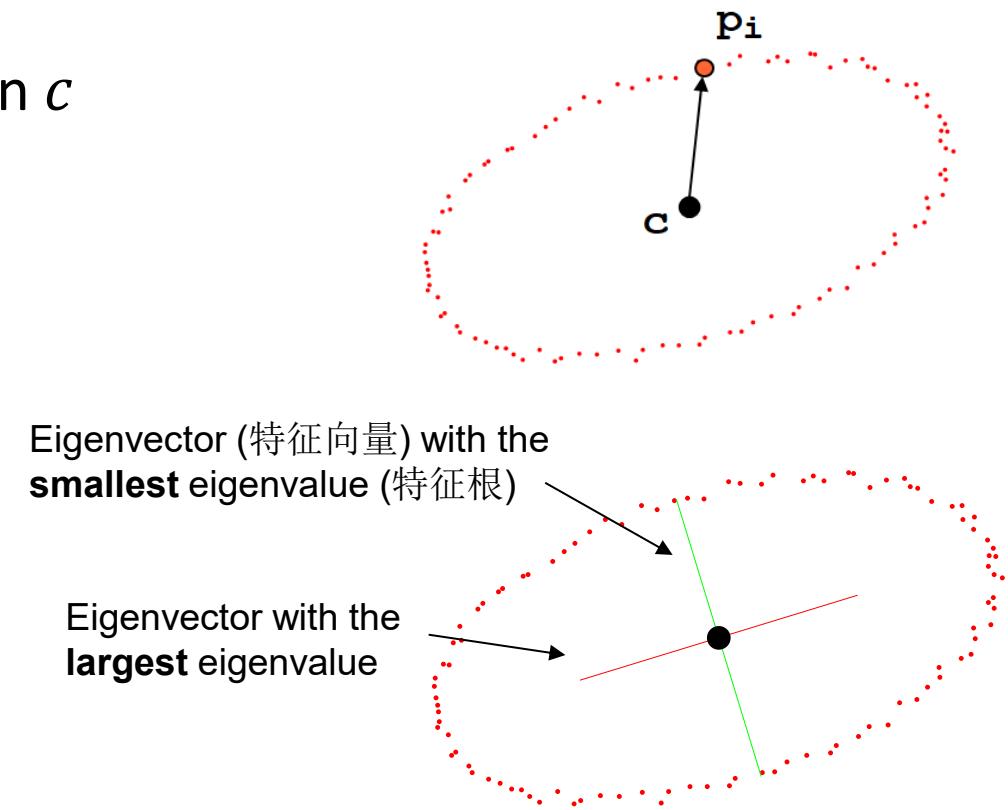
$$P = \begin{pmatrix} p_{1x} & c_x & p_{2x} & c_x & \dots & p_{nx} & c_x \\ p_{1y} & c_y & p_{2y} & c_y & \dots & p_{ny} & c_y \\ p_{1z} & c_z & p_{2z} & c_z & \dots & p_{nz} & c_z \end{pmatrix}$$

- Build the covariance matrix (协方差矩阵) : $M = P \times P^T$

$$\text{cov}(X) = \frac{1}{m} XX^T = \begin{bmatrix} \text{cov}(x, x) & \text{cov}(x, y) & \text{cov}(x, z) \\ \text{cov}(y, x) & \text{cov}(y, y) & \text{cov}(y, z) \\ \text{cov}(z, x) & \text{cov}(z, y) & \text{cov}(z, z) \end{bmatrix}$$

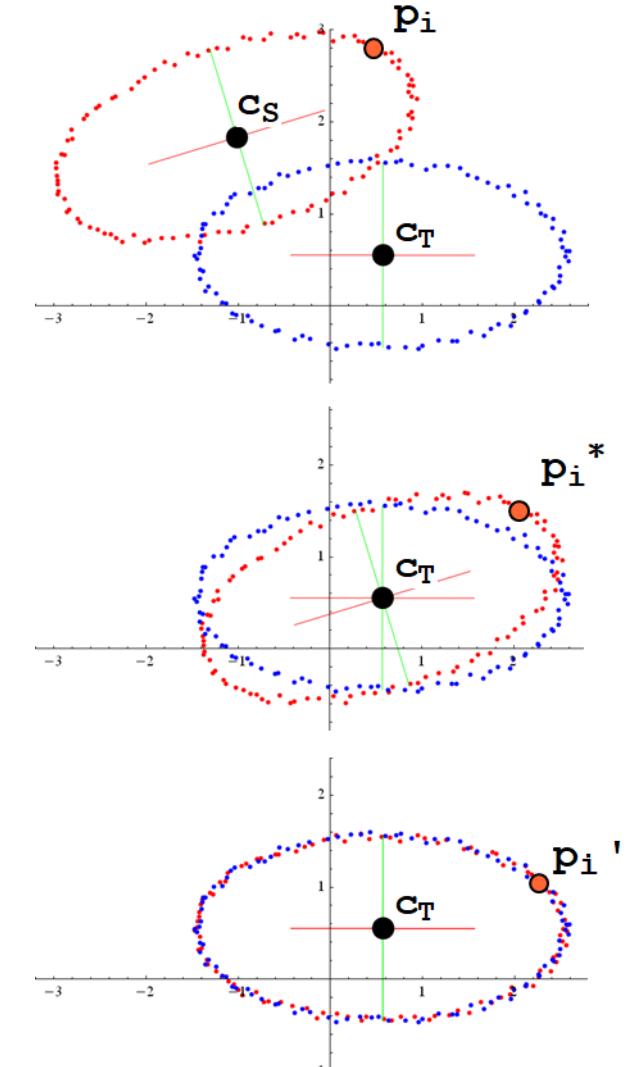
- Eigenvectors of M represent principal directions of shape variation
 - The eigenvectors form orthogonal axes (2 vectors in 2D; 3 vectors in 3D)
 - Note they are “un-signed”: lacking an orientation.

<https://graphics.stanford.edu/courses/cs233-20-spring/ReferencedPapers/LectureNotes-PCA.pdf>



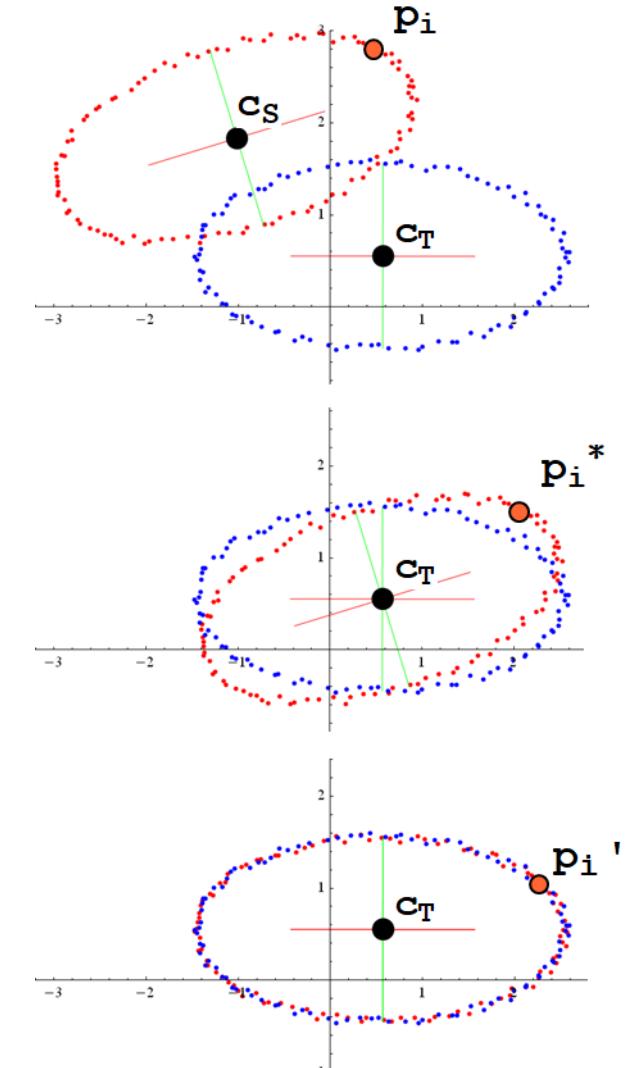
ICP: Initialize by PCA

- After finding axes, how to do alignment?
- PCA-based alignment
 - Let c_S, c_T be centroids of source and target.
- First, translate source to align c_S with c_T :
$$p_i^* = p_i + (c_T - c_S)$$
- Next, find rotation R that aligns two sets of PCA axes
$$p_i' = c_T + R \times (p_i^* - c_T)$$
- Combined:
$$p_i' = c_T + R \times (p_i - c_S)$$



ICP: Initialize by PCA

- After finding axes, how to do alignment?
- PCA-based alignment
 - Let c_S, c_T be centroids of source and target.
- First, translate source to align c_S with c_T :
$$p_i^* = p_i + (c_T - c_S)$$
- Next, find rotation R that aligns two sets of PCA axes
$$p_i' = c_T + R \times (p_i^* - c_T)$$
- Combined:
$$p_i' = c_T + R \times (p_i - c_S)$$
- Then we get the initial transformation:
$$R = R, t = c_T - R \times c_S$$



ICP

- 1. **Initialize** transformation R, t by PCA
- 2. **Match** each to closest point $p'_i = Rp_i + t$ on other scan
- 3. **Reject** pairs with distance $> k$ times median
- 4. Construct **error function**:

$$E = \sum |Rp_i + t - q_i|^2$$

- 5. Minimize the error by SVD
- 6. Loop step 2-5 until the error is small enough

ICP: Minimize by SVD

- Now we have pairs (p_i, q_i) from two scans
- How to find a better R, t to minimize error function?
- Use Singular Value Decomposition (SVD/奇异值分解):
 - Let P be a matrix whose i-th column is vector $p_i - c_S$
 - Let Q be a matrix whose i-th column is vector $q_i - c_T$
 - Forming the cross-covariance matrix (互协方差矩阵)

$$M = P \times Q^T$$

- Computing SVD

$$M = U \times \Sigma \times V^T$$

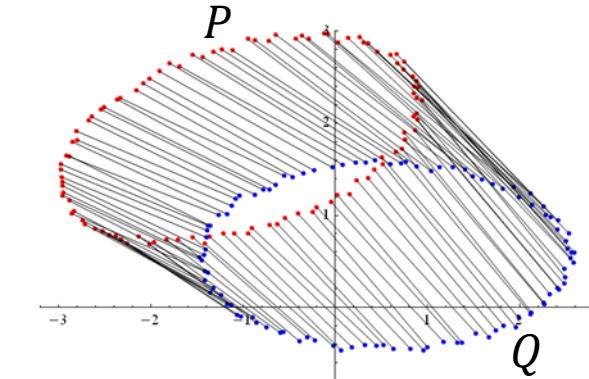
- The rotation matrix is

$$R = V \times U^T \quad (\text{hint: consider } M = (RP) \times Q^T)$$

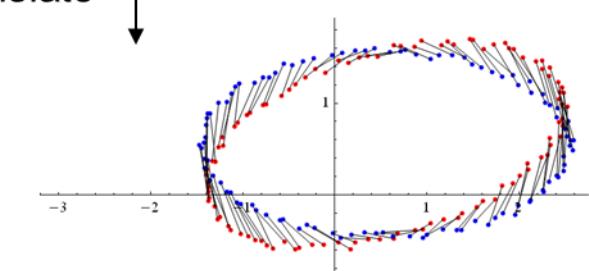
- Update the translation after rotate the source:

$$t = c_T - R \times c_S$$

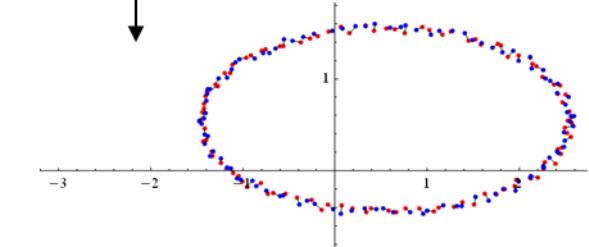
$$M = \begin{bmatrix} \text{Cov}(p_x, q_x) & \text{Cov}(p_x, q_y) & \text{Cov}(p_x, q_z) \\ \text{Cov}(p_y, q_x) & \text{Cov}(p_y, q_y) & \text{Cov}(p_y, q_z) \\ \text{Cov}(p_z, q_x) & \text{Cov}(p_z, q_y) & \text{Cov}(p_z, q_z) \end{bmatrix}$$



Translate



Rotate



ICP: Minimize by SVD

- The Orthogonal Procrustes problem (正交普鲁克问题) [1966]
https://en.wikipedia.org/wiki/Orthogonal_Procrustes_problem

The **orthogonal Procrustes problem**^[1] is a **matrix approximation problem** in **linear algebra**. In its classical form, one is given two **matrices** A and B and asked to find an **orthogonal matrix** Ω which most closely maps A to B .^{[2][3]}
Specifically, the orthogonal Procrustes problem is an **optimization problem** given by

$$\begin{aligned} & \underset{\Omega}{\text{minimize}} && \|\Omega A - B\|_F \\ & \text{subject to} && \Omega^T \Omega = I, \end{aligned}$$

- Algorithm Kabsch-Umeyama [1976]
https://en.wikipedia.org/wiki/Kabsch_algorithm
matlab, c++, python

- Least-Squares Rigid Motion Using SVD

ICP

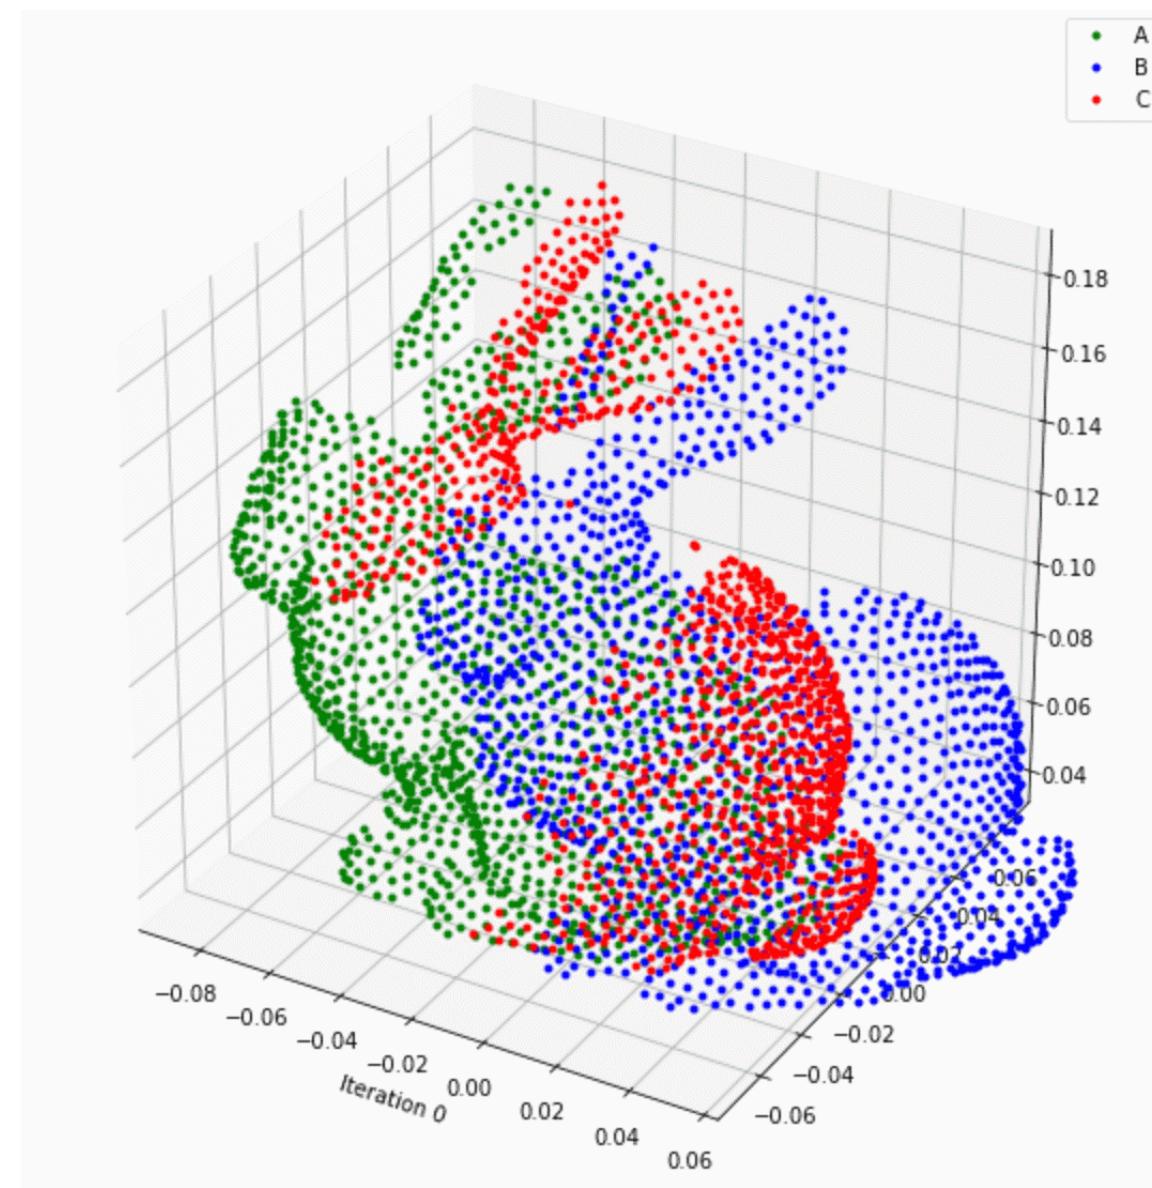
- 1. **Initialize** transformation R, t by PCA
- 2. **Match** each to closest point $p'_i = Rp_i + t$ on other scan
- 3. **Reject** pairs with distance $> k$ times median
- 4. Construct **error function**:

$$E = \sum |Rp_i + t - q_i|^2$$

- 5. **Minimize** the error by SVD
- 6. **Loop** step 2-5 until the error is small enough

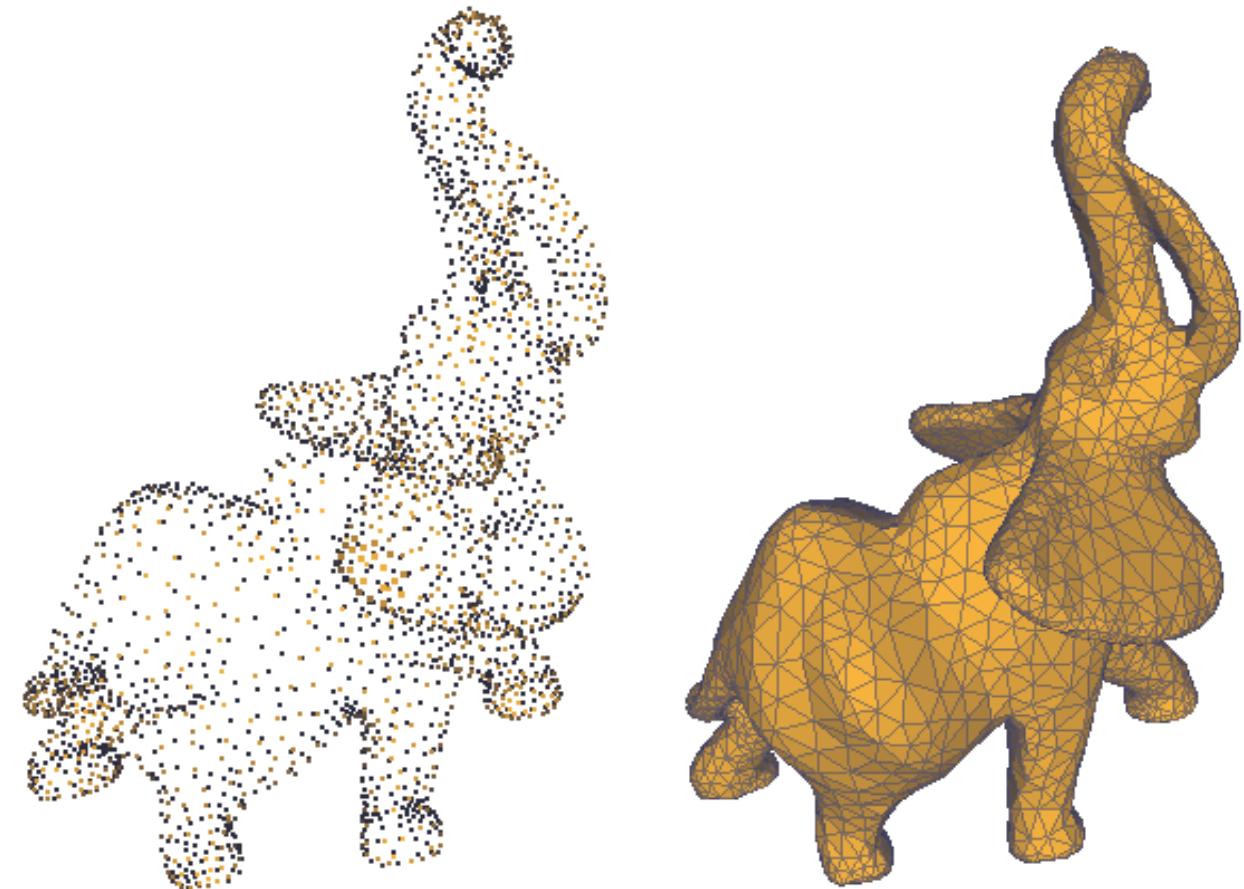
ICP

- Visualization of an ICP variant
 - <https://laempty.github.io/pypoints/tutorials/icp.html#References>
- Also, there are many variants of ICP. They are more robust and more efficient. Refer to:
 - https://gfx.cs.princeton.edu/proj/iccv05_course/iccv05_icp_gr.ppt



Surface Reconstruction

- Motivation:
 - 1) Effective rendering of the model
 - 2) Computational analysis
 - 3) Other geometry processing: parameterization, morphing, blending etc..
- Methods:
 - CAD: manually construct
 - Algorithm: automatically generate
 - mainly from point cloud

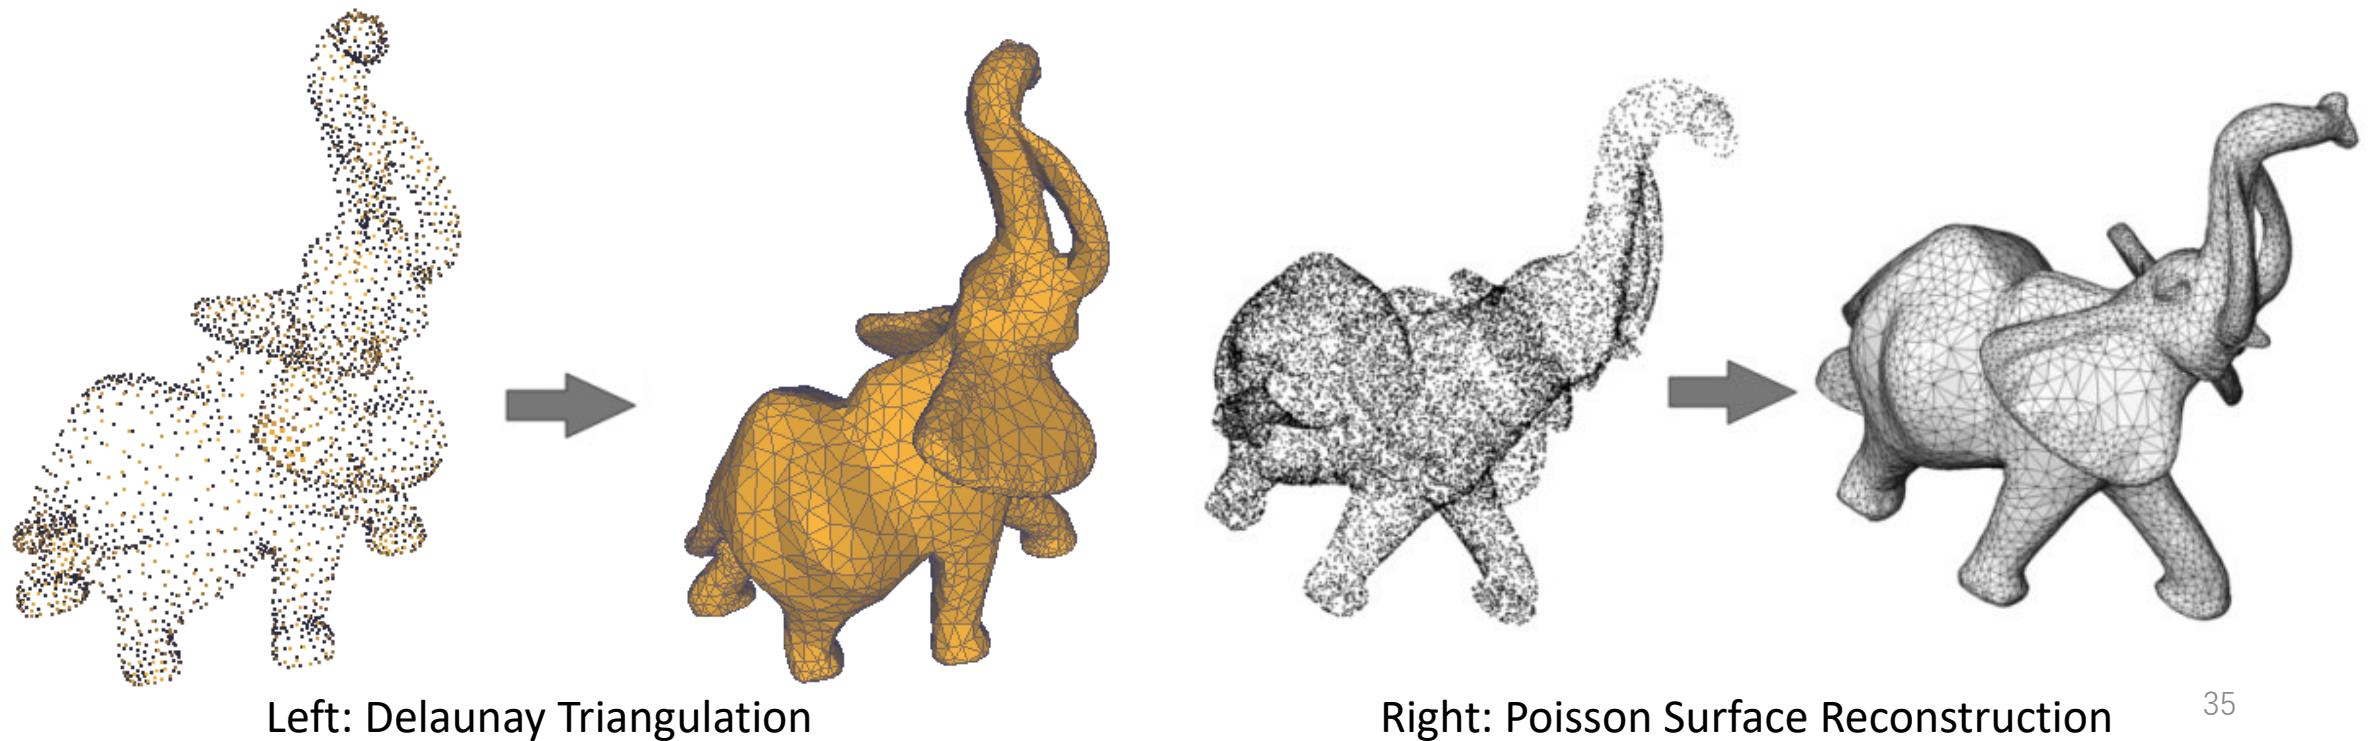


Courtesy: Jiju Peethambaran and Ramanathan Muthuganapathy

Figure: before and after surface reconstruction (triangulation).

Surface Reconstruction Algorithm

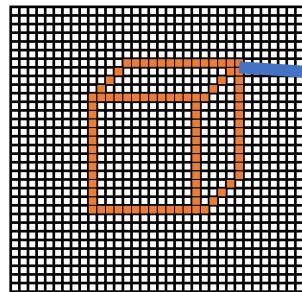
- Classic surface reconstruction from point cloud:
 - Directly: Triangulation
 - Delaunay Triangulation
 - Indirectly: Implicit Surfaces + Marching Cubes
 - Poisson Surface Reconstruction



How do points divide space?

- On Images (Grids):

- Pixels

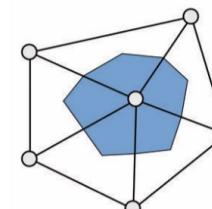
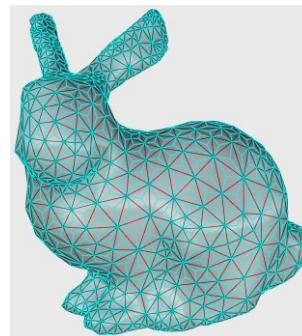


(x, y) - position

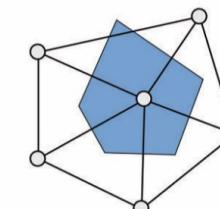
255	0	0
-----	---	---

- On Meshes:

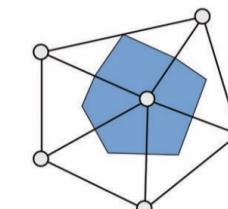
- Local Averaging Region



Barycentric cell



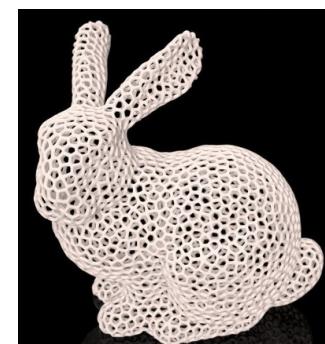
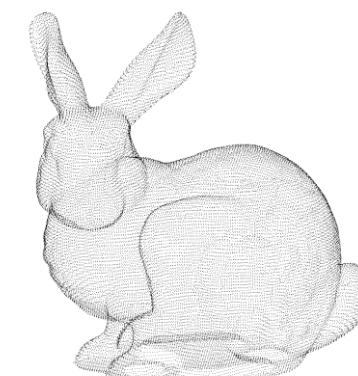
Voronoi cell



Mixed Voronoi cell

- On Point Clouds:

- Voronoi cells (voronoi diagram)



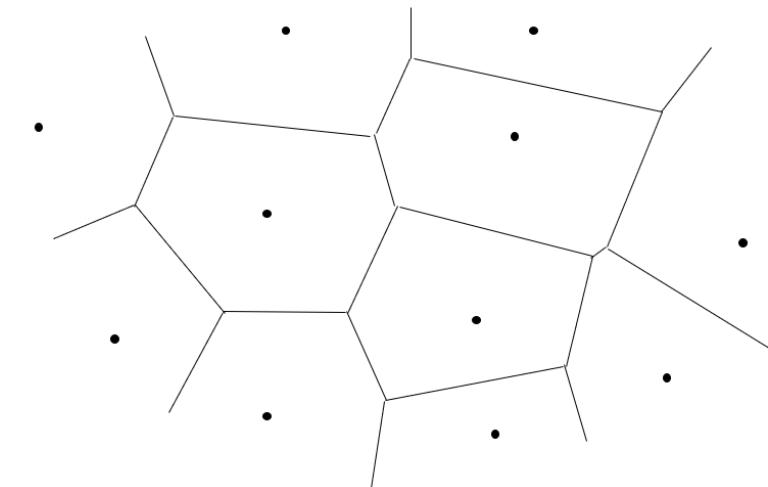
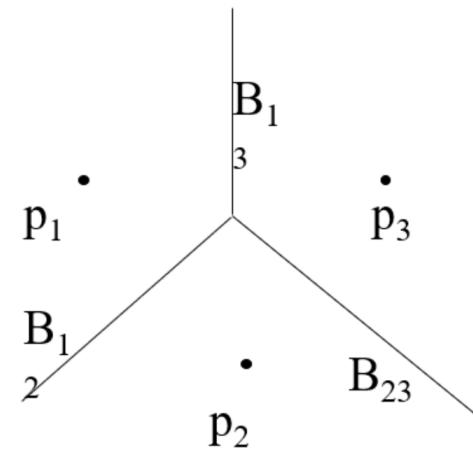
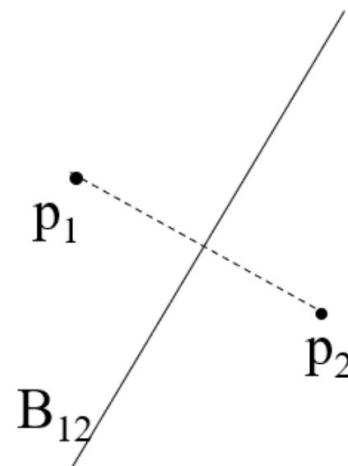
voronoi diagram

Equivalence Space Division

2D Voronoi diagram

- Voronoi diagram is a partition of a plane into regions close to each of a given set of objects.
- All those points assigned to p_i from the **Voronoi region** $V(p_i)$ consists of all the points at least as close to p_i as to any other site:

$$V(p) = \{x \mid d(p_i - x) \leq d(p_j - x), \forall j \neq i\}$$



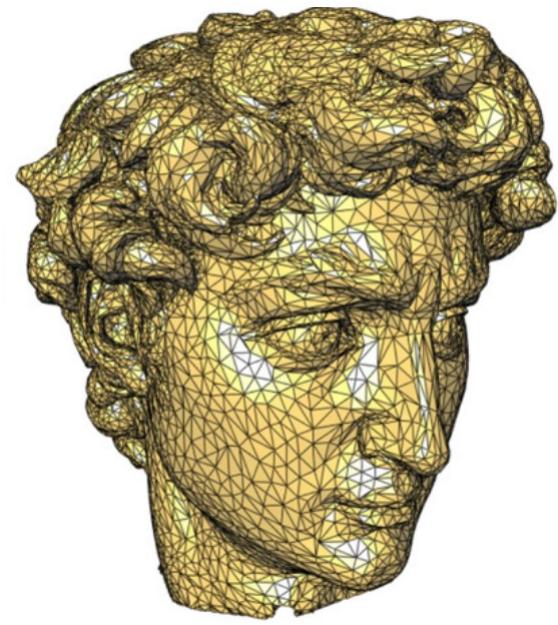
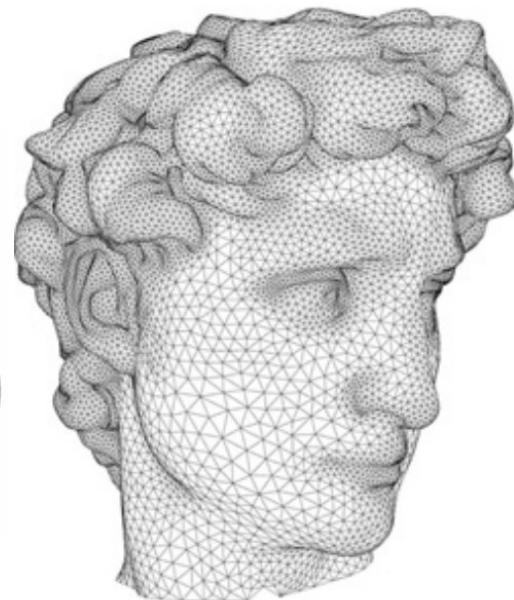
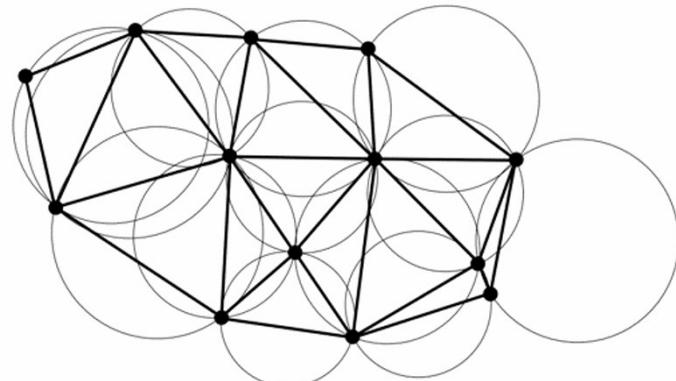
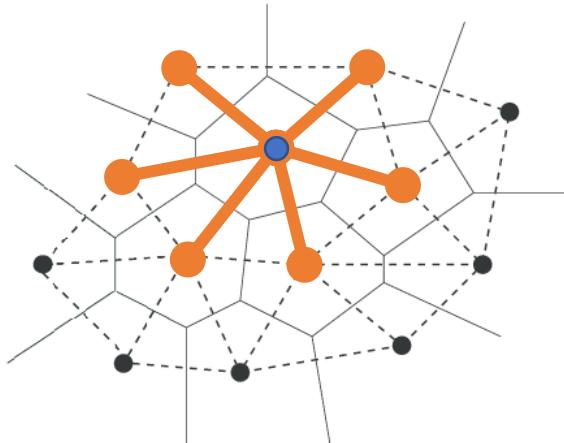
Voronoi diagram

- A Voronoi tessellation emerges by radial growth from seeds outward.



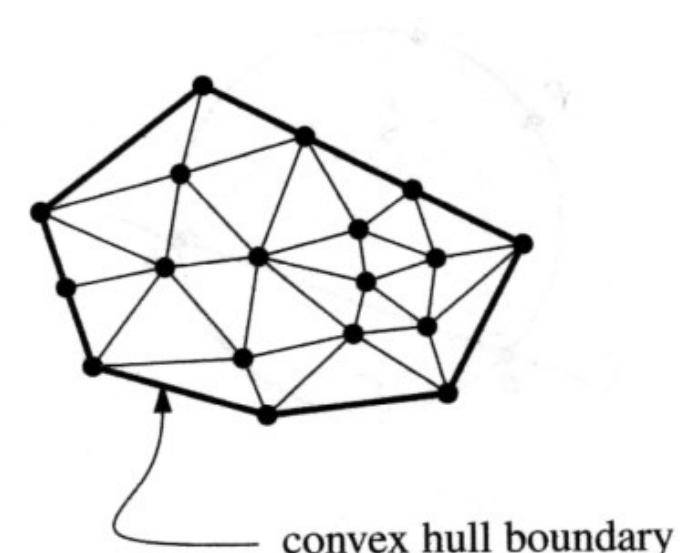
Delaunay triangulation

- The geometric dual of the Voronoi diagram. There is a line segment connecting P_i and P_j in a Delaunay triangulation if and only if the Voronoi Diagram Regions of P_i and P_j share the same edge.
- Property:
For each triangle of the triangulation, the circumcircle of that triangle is empty of all other sites.
Angle Optimal Triangulations



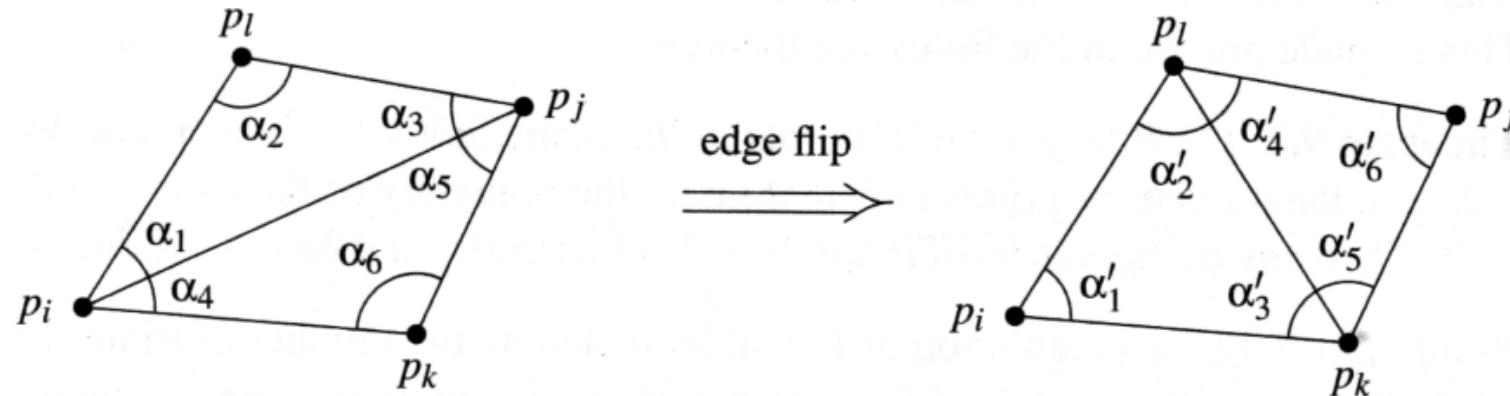
Delaunay Triangulation in 2D

- Delaunay Triangulation has a good property: **Angle Optimal Triangulations**
- What is **triangulation**?
 - No additional edge can be added to connect two vertices without causing intersections with existing edges.
- What is **angle optimal**?:
 - Create **angle vector** $A(T)$ of the sorted angles of triangulation T :
$$A(T) = (\alpha_1, \alpha_2, \dots, \alpha_{3m}), \quad \alpha_1 \leq \alpha_2 \leq \dots \leq \alpha_{3m}$$
 - $A(T)$ is larger than $A(T')$ if and only if there exists an i such that:
$$\alpha_j = \alpha'_j, \forall j < i \text{ and } \alpha_i > \alpha'_i$$
 - **angle optimal** has the largest angle vector $A(T)$.
 - **Equivalent to maximizing the minimum angle**:
the triangulation with the largest angle vector $A(T)$ is the one that maximizes the minimum angle.



Delaunay Triangulation in 2D

- For example, consider two adjacent triangles of T :



- If the two triangles form a convex quadrilateral, we could have an alternative triangulation by performing an **edge flip** on their shared edge.
- And, if:

$$\min_{1 \leq i \leq 6} \alpha_i < \min_{1 \leq i \leq 6} \alpha'_i$$

- We can get a **better** triangulation by such an edge flip.
- And we call such an edge p_ip_j by **illegal edge**.

Delaunay Triangulation in 2D

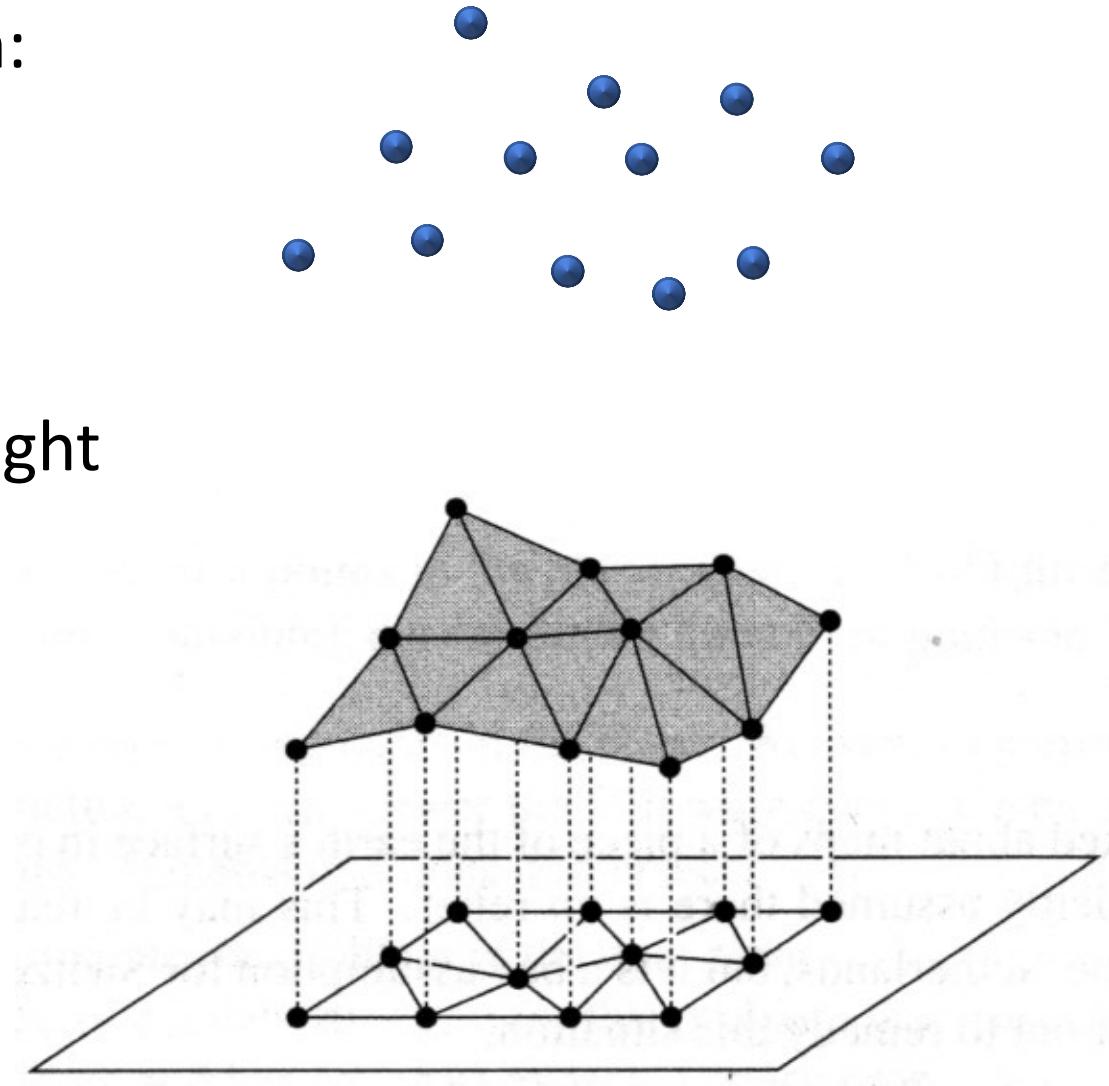
- By the idea of “edge flip”, we can give a naïve algorithm to get a “best” triangulation:
 - 1. Compute a triangulation of input points P .
 - 2. Flip illegal edges of this triangulation until all edges are legal.
- Algorithm will terminate because there is a finite number of triangulations.
- The final “best” triangulation can be called **Delaunay Triangulation**.
- Delaunay Triangulation has many interesting properties including maximizing minimum angle.
- The **Delaunay triangulation** is the straight-line dual of the **Voronoi Diagram**.
Note: The Delaunay edges don't have to cross their Voronoi duals.

Delaunay Triangulation in 2D

- However, such a naïve algorithm is too slow.
- There is a series of developed methods to perform efficient Delaunay Triangulation using divide and conquer algorithm.
- Here is some 2D visualization demo:
 - <https://cartography-playground.gitlab.io/playgrounds/triangulation-delaunay-voronoi-diagram/>
 - <https://travellermap.com/tmp/delaunay.htm>
- Most importantly, it can be extended to arbitrary dimensions (3D or dD).

Delaunay Triangulation in 2D

- An Application of 2D Delaunay Triangulation:
Terrains
- Set of data points $A \subset R^2$
- Height $f(p)$ defined at each point p in A
- How can we most naturally approximate height
of points not in A ?
- Option: Linear Interpolation?
 - Determine a **triangulation** of $A \subset R^2$
 - Then raise points to desired height
 - Now, we can get heights of points not in A



Borrowed from Glenn Eguchi, MIT 46

Delaunay Triangulation in 3D

- Extend to a 3D point cloud: set of points $A \subset R^3$
- Attribute $f(p)$ defined at each point p in A
 - The attribute can be color, UV, depth (from a certain view), ...
- How to know the attribute of points not in A ?
 - The point cloud is sparse.
 - But we want a continuous $f(p)$.
- Option: Bilinear Interpolation!
 - Via triangulation in 3D
 - That's why the **Mesh** is such a powerful 3D representation



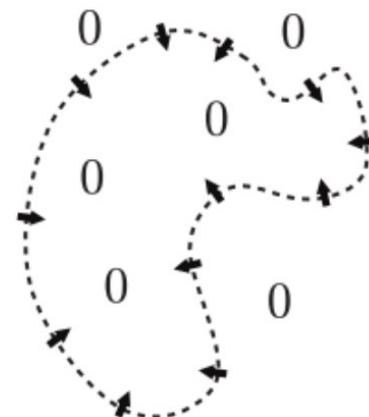
Courtesy: Jiju Peethambaran and Ramanathan Muthuganapathy

Poisson Surface Reconstruction

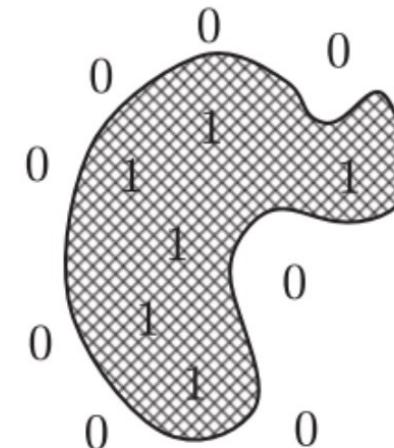
- The input is oriented points \vec{V} (points with normals) sampled from a shape M .
- Poisson Surface Reconstruction construct an implicit surface first.
- χ_M is indicator function of M .
- We want to get $\nabla \chi_M$ to represent the surface ∂M



Oriented points
 \vec{V}



Indicator gradient
 $\nabla \chi_M$



Indicator function
 χ_M



Surface
 ∂M

Poisson Surface Reconstruction

- According to Poisson Equation:

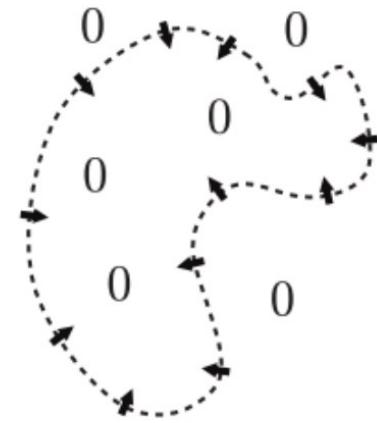
$$\nabla \cdot (\nabla \chi) = \nabla \cdot \vec{V} \Leftrightarrow \Delta \chi = \nabla \cdot \vec{V}$$

- For 3D, we can construct an octree and solve it by PDE
- Then we get the implicit representation of surface.



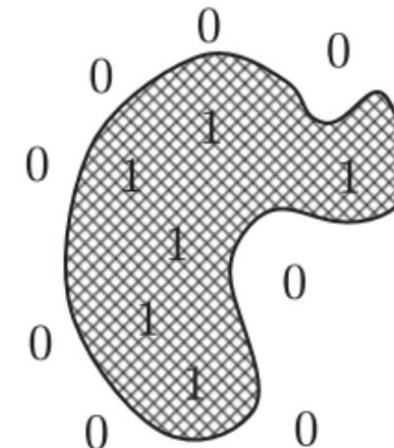
Oriented points

$$\vec{V}$$



Indicator gradient

$$\nabla \chi_M$$



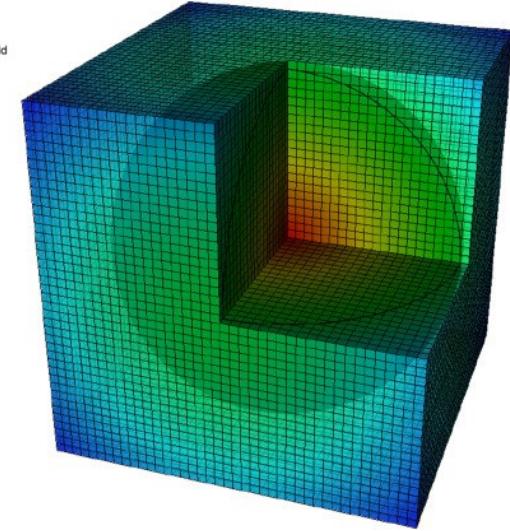
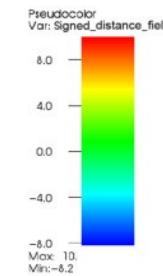
Indicator function

$$\chi_M$$



Surface

$$\partial M$$



Poisson Surface Reconstruction

- Finally we get the triangle mesh by [Marching Cubes](#).
- An example:

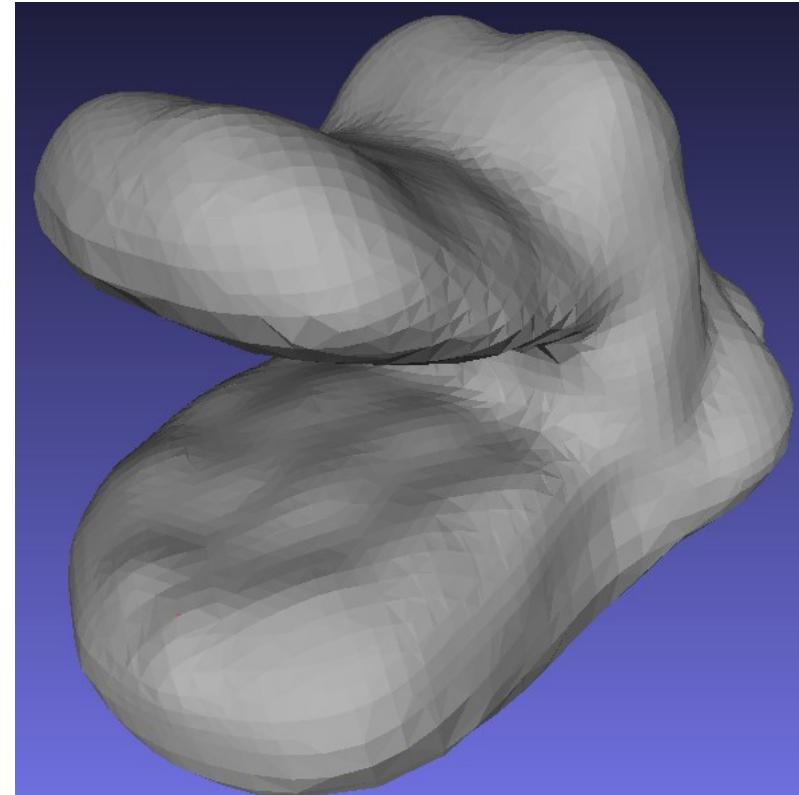
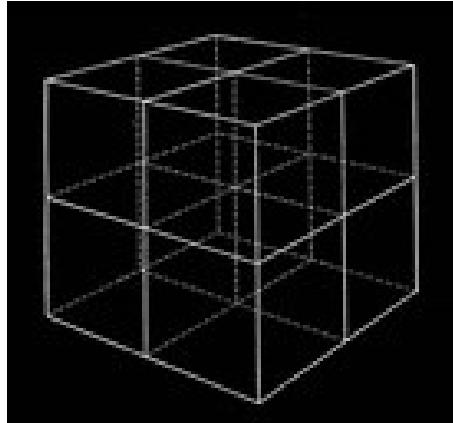
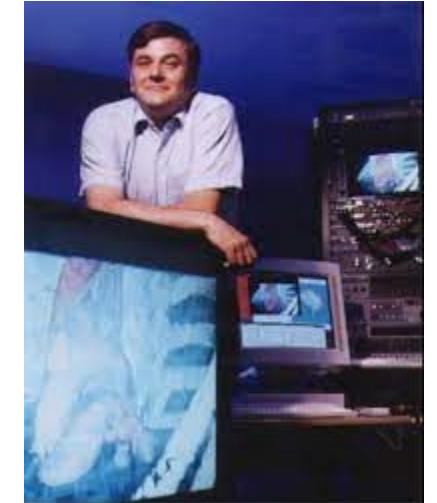


Figure: (a) On the left is the point cloud. (b) On the right is the reconstructed mesh.



Marching Cubes:

A High Resolution 3D Surface Construction Algorithm



William E. Lorensen

Harvey E. Cline

General Electric Company

Corporate Research and Development, SIGGRAPH 1987

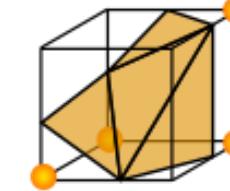
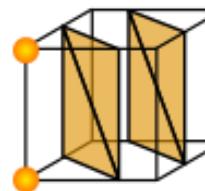
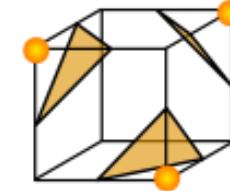
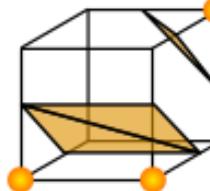
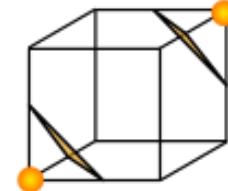
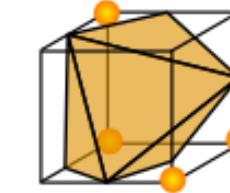
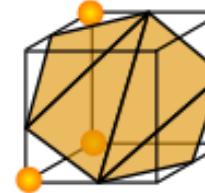
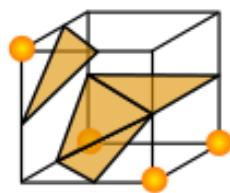
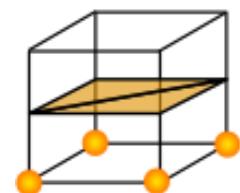
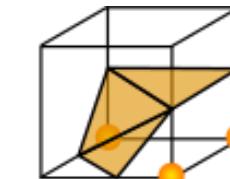
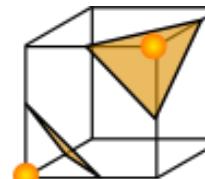
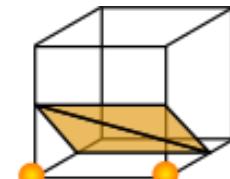
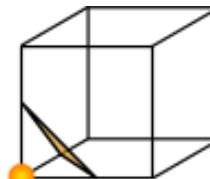
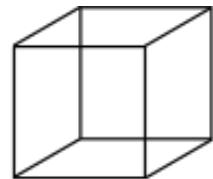
- W. Lorensen and H. Cline. “Marching cubes: A High Resolution 3D Surface Construction Algorithm”, Proceedings of SIGGRAPH 1987, pages 163-169, 1987.
- The Visible Human: http://www.nlm.nih.gov/research/visible/visible_human.html
- Marching Cubes Demo/Tutorial: <http://users.polytech.unice.fr/~lingrand/MarchingCubes/accueil.html>

Step 1: Surface Intersection

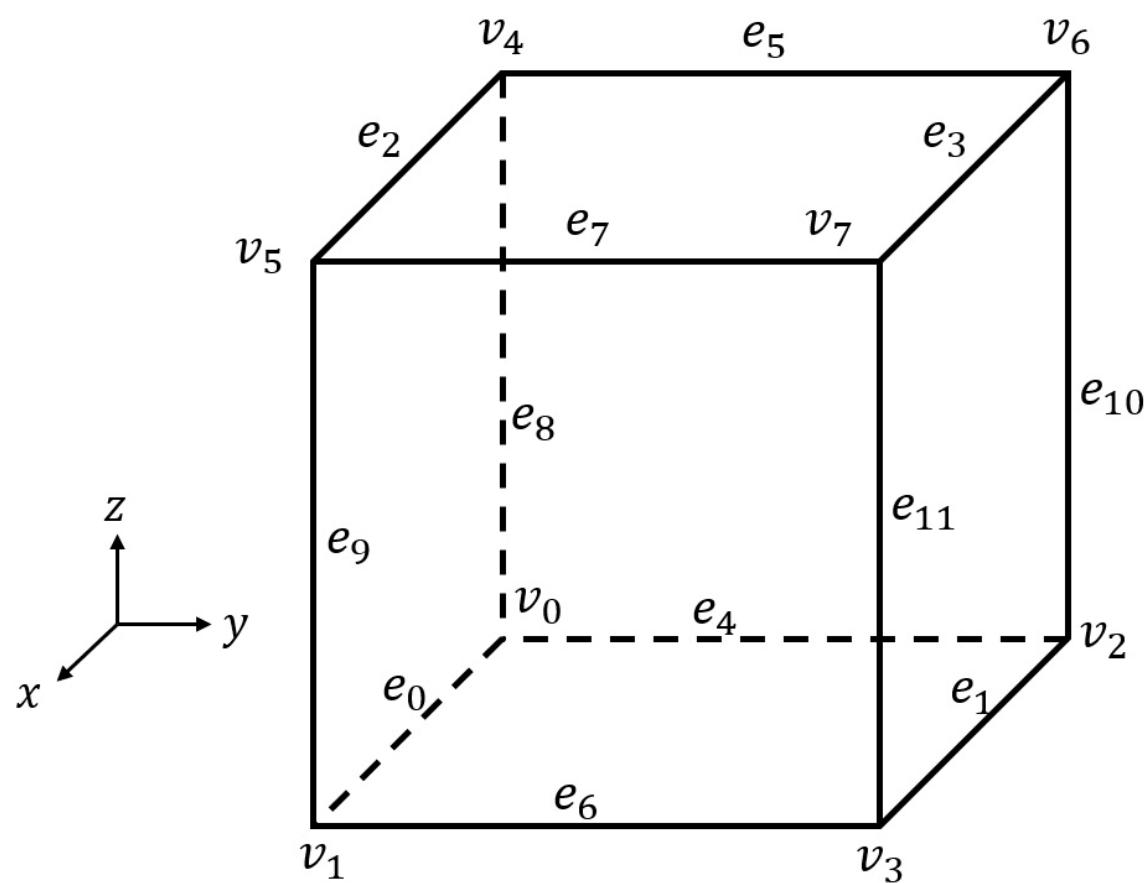
- Given user specified value T_U , binary vertex assignment: $(p(i, j, k) \geq T_U? 1 : 0)$
 - Set cube vertex to value of 1 if the data value at that vertex exceeds (or equals) the value of the surface we are constructing
 - Otherwise, set cube vertex to 0
- If a vertex = 1 then it is “inside” the surface
- If a vertex = 0 then it is “outside”
- Any cube with vertices of both types is “intersected” by the surface.

Step 2 : Triangulation

- For each cube, we have 8 vertices with 2 possible states each (inside or outside).
- This gives us 2^8 possible patterns = 256 cases.
- Enumerate cases to create a LUT (LookUp Table)
- Use symmetries to reduce problem from 256 to 15 cases.



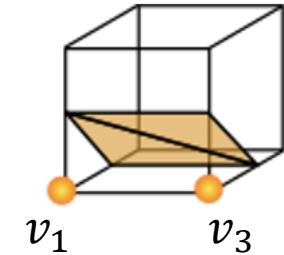
Step 2 : Triangulation



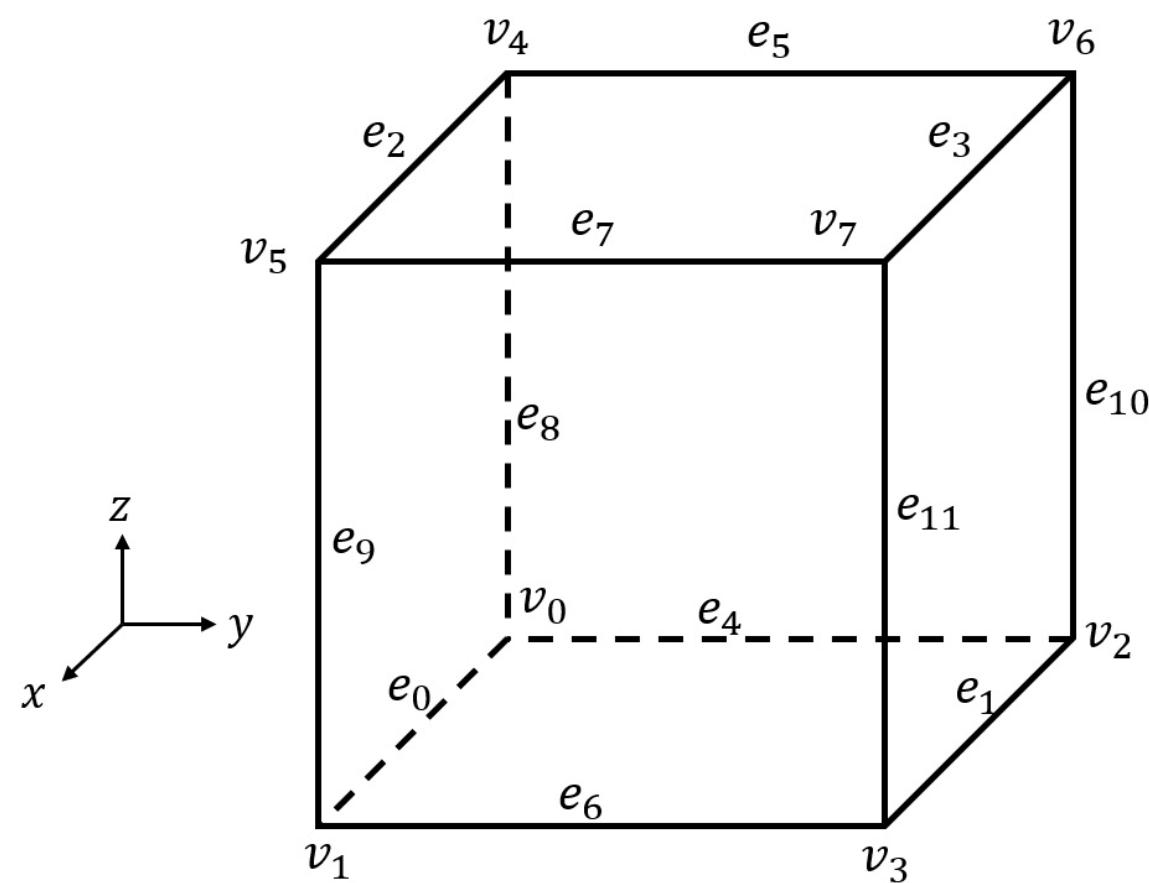
Vertex bit mask:

v_7	v_6	v_5	v_4	v_3	v_2	v_1	v_0
-------	-------	-------	-------	-------	-------	-------	-------

- 2.1 Use vertex bit mask to create an index for each case based on the state of the vertexes.
- Example:
- If v_1 and v_3 are “outside” the surface and thus our index would be 245 (1111-0101 for binary).



Step 2 : Triangulation



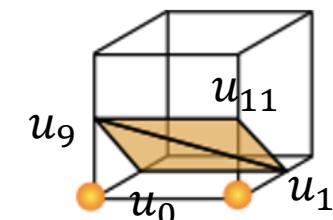
Vertex bit mask:

v_7	v_6	v_5	v_4	v_3	v_2	v_1	v_0
-------	-------	-------	-------	-------	-------	-------	-------

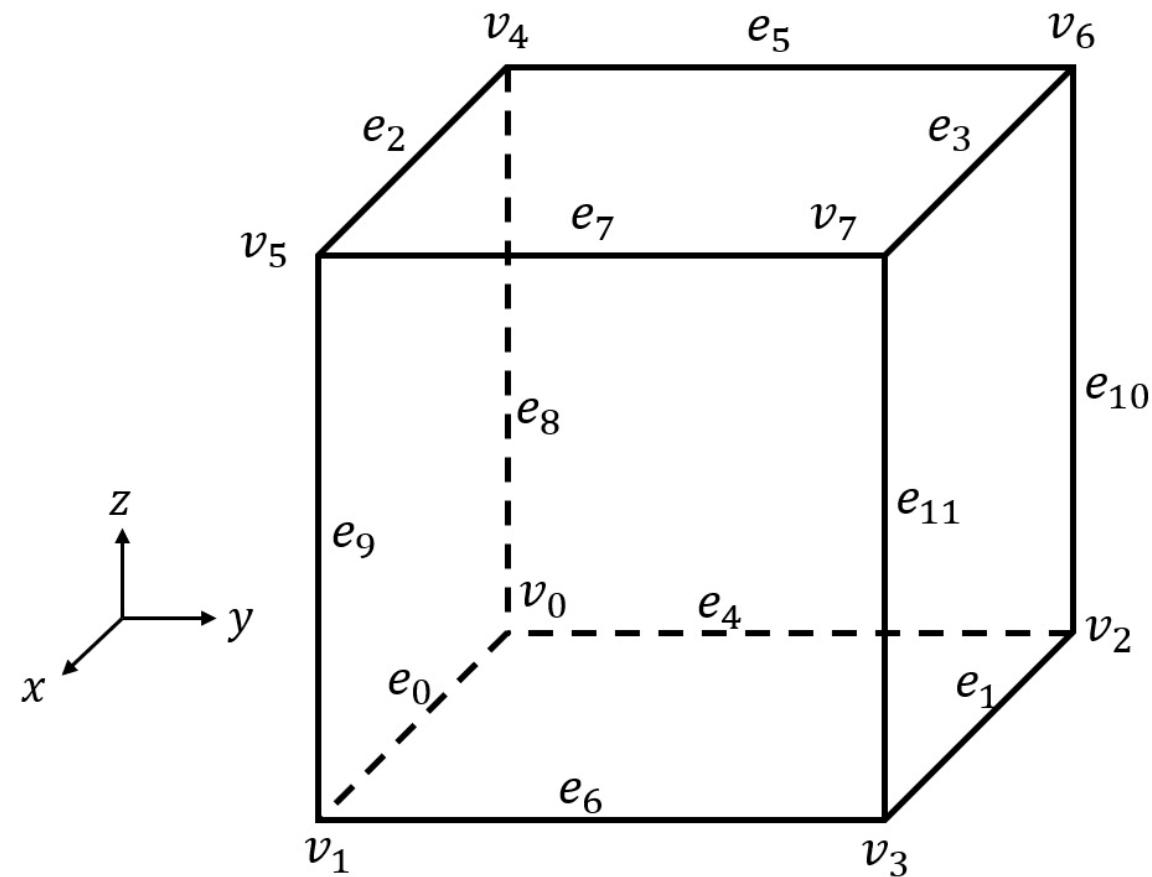
Edge State:

e_{11}	e_{10}	e_9	e_8	e_7	e_6	e_5	e_4	e_3	e_2	e_1	e_0
----------	----------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

- 2.2 Using the index to tell which edge the surface intersects, we can then linearly interpolate the surface intersection along the edge.
- Example:
- We search edge state by index 245, getting code 2563 (1010-0000-0011 for binary) which means e_{11} , e_9 , e_1 and e_0 are intersected by the surface.
- And we insert intersection vertices on e_{11} , e_9 , e_1 and e_0 , generating u_{11} , u_9 , u_1 and u_0 .



Step 2 : Triangulation



Vertex bit mask:

v_7	v_6	v_5	v_4	v_3	v_2	v_1	v_0
-------	-------	-------	-------	-------	-------	-------	-------

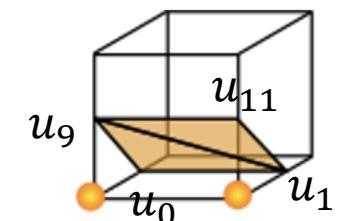
Edge State:

e_{11}	e_{10}	e_9	e_8	e_7	e_6	e_5	e_4	e_3	e_2	e_1	e_0
----------	----------	-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

Edge Order:

o_0, o_1, o_2	o_3, o_4, o_5	\dots
-----------------	-----------------	---------

- 2.3 Using the index to tell how the intersection vertices are connected to form a triangle facet.
- Example:
- And we insert intersection vertices on e_{11} , e_9 , e_1 and e_0 , generating u_{11} , u_9 , u_1 and u_0 .
- Finally we search edge order by index 245, getting array $[[0, 1, 11], [9, 0, 11]]$, which means there are two triangle facets to generate: $u_0-u_1-u_{11}$ and $u_9-u_0-u_{11}$.



Step 3 : Surface normals

- To calculate surface normal, we need to determine gradient vector \vec{G} (derivative of the density function).
- To estimate the gradient vector at the surface of interest, we first estimate the gradient vectors at the vertices and interpolate the gradient at the intersection.
- The gradient at cube vertex (i, j, k) , is estimated using central differences along the three coordinate axes by:

$$G_x(i, j, k) = \frac{D(i + 1, j, k) - D(i - 1, j, k)}{\Delta x}$$

$$G_y(i, j, k) = \frac{D(i, j + 1, k) - D(i, j - 1, k)}{\Delta y}$$

$$G_z(i, j, k) = \frac{D(i, j, k + 1) - D(i, j, k - 1)}{\Delta z}$$

$D(i, j, k)$ is the density at pixel (i, j) in slice k .

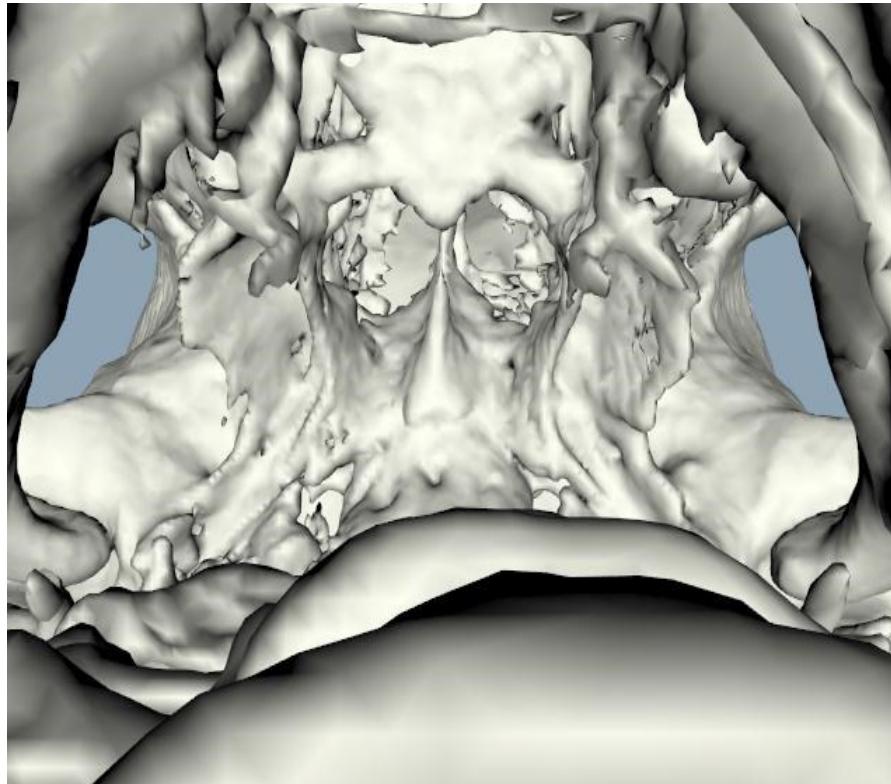
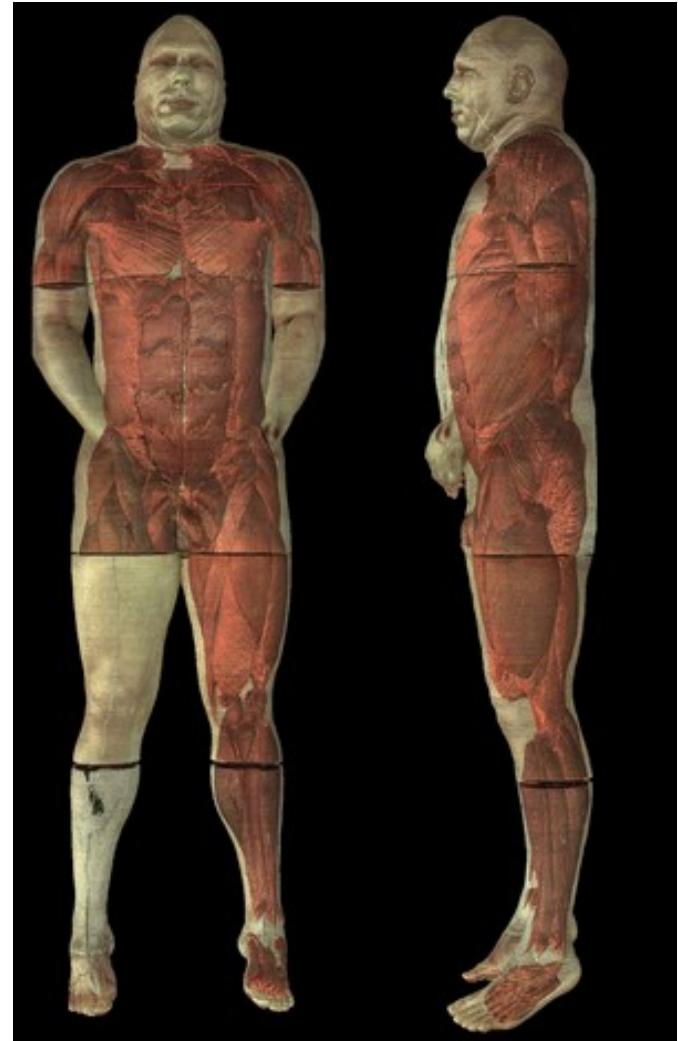
$\Delta x, \Delta y, \Delta z$ are lengths of the cube edges

Step 3 : Surface normals

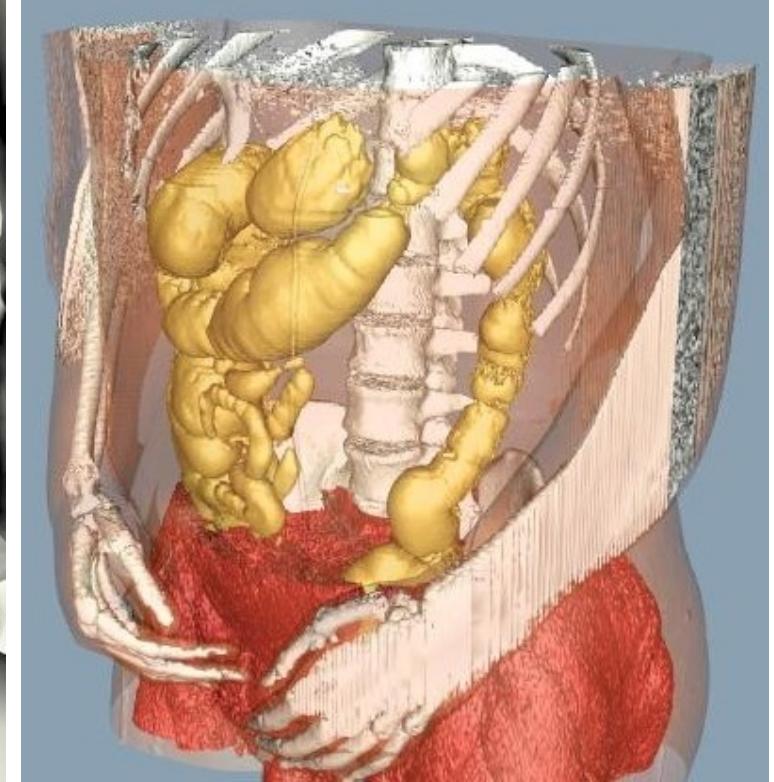
- Dividing the gradient by its length produces the unit normal at the vertex required for rendering.
- Then the algorithm linearly interpolates this normal to the point of intersection.

$$\vec{N}(i, j, k) = \frac{\vec{G}(i, j, k)}{\|\vec{G}(i, j, k)\|_2}$$

Results –The Visible Man

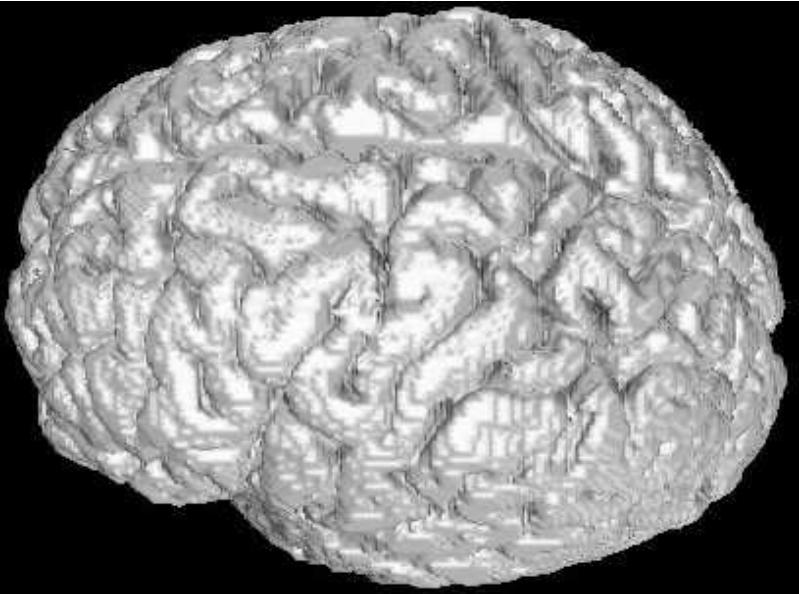


Inside skeleton view

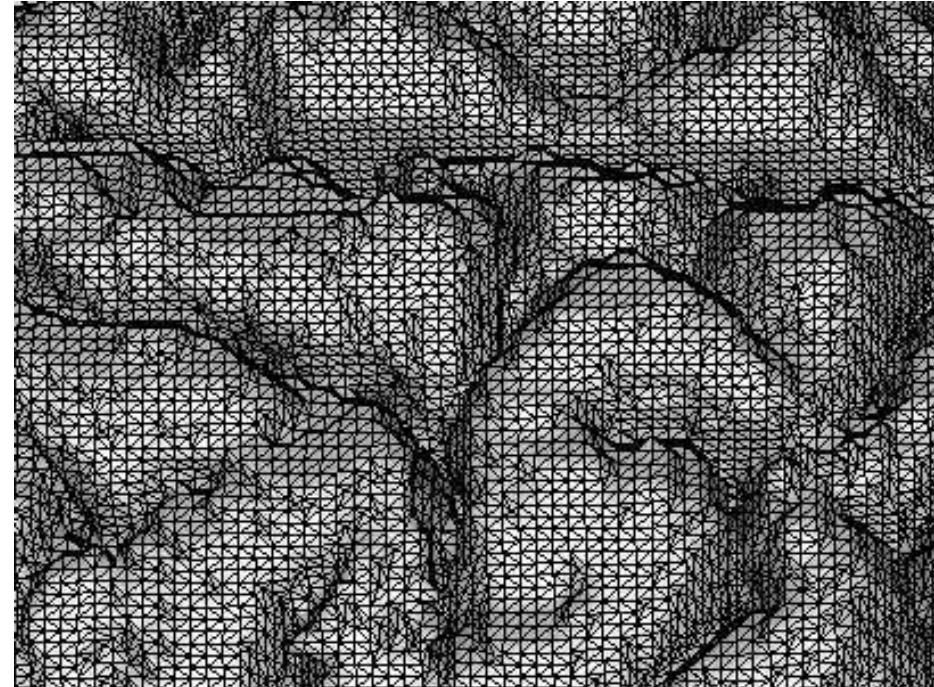


Torso / bowels

Results



Human brain surface reconstructed by using marching cubes (128,984 vertices and 258,004 triangles)



Magnified display of brain surface

Algorithm Summary

1. Scan 2 slices and create cube
 2. Calculate index for cube based on vertices
 3. Use index to lookup list of edges intersected
 4. Use densities to interpolate edge intersections
 5. Calculate unit normal at each edge vertex using central differences.
Interpolate normal to each triangle vertex
 6. Output the triangle vertices and vertex normals
 7. March to next position and repeat.
- Enhancements:
 - Take advantage of pixel-to-pixel, line-to-line, and slice-to-slice coherence by keeping previous calculations.
 - Thousands more

Model Fitting

- Motivation
 - How do we fit models (i.e., a parametric representation of data that's smaller than the data) to data?
- For example, given such a point cloud:
 - How to find a plane?
 - A 3D plane can be described as:
 - $Ax + By + Cz = D$
 - But how to know the parameters?

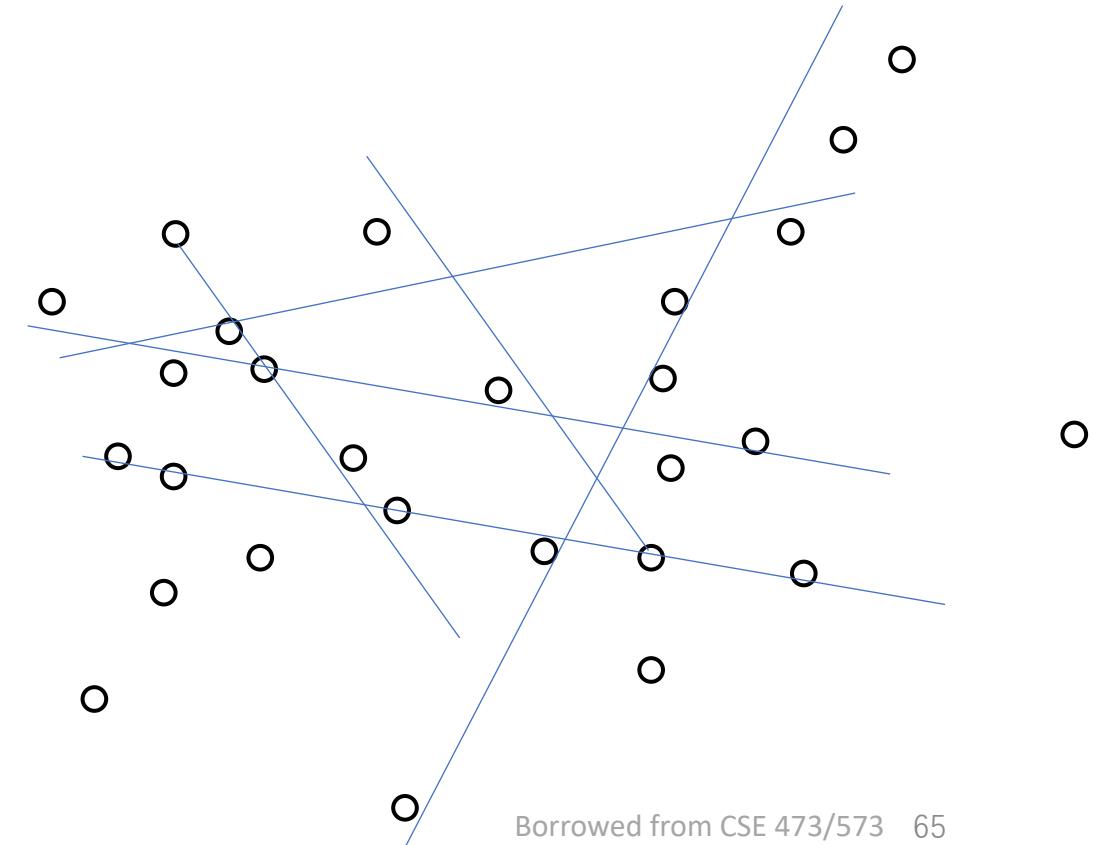


Courtesy: <https://github.com/STORM-IRIT/Plane-Detection-Point-Cloud>

Figure: the plane detection of the model of a building

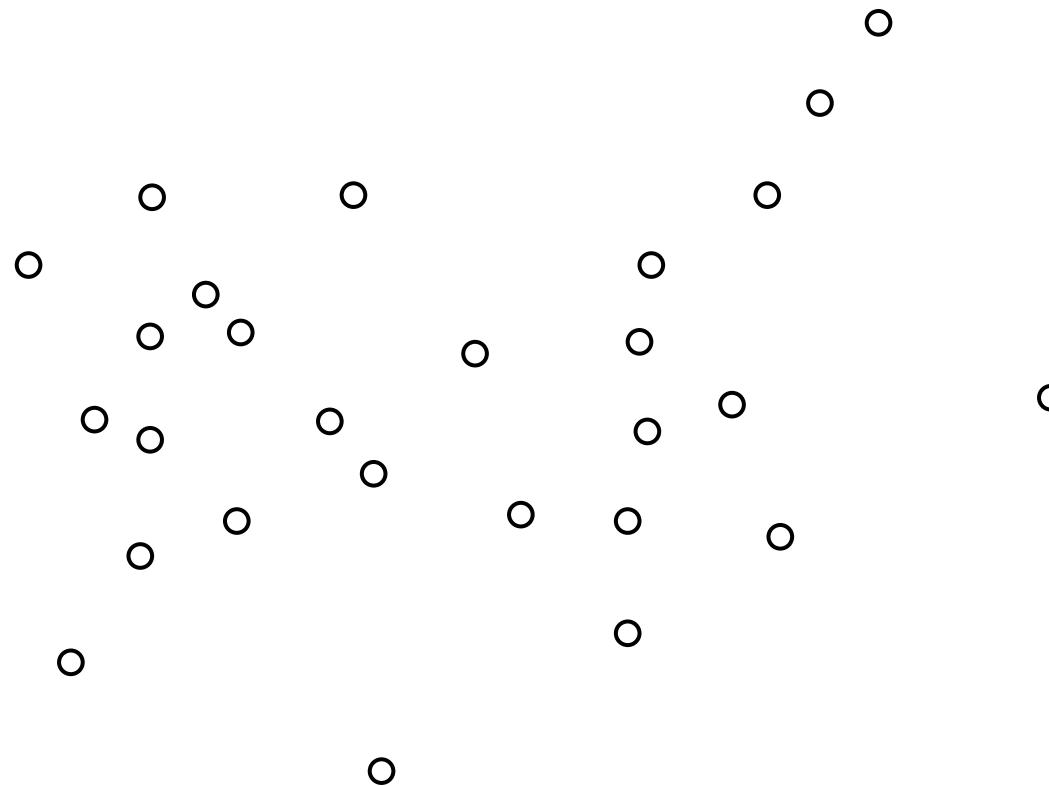
Model Fitting

- Let's consider a simple case: line fitting.
- Given a set of 2D points, how to fit the best possible line to these points?
- Brute Force Search - 2^N possibilities!
- Not feasible.
- Better Strategy?



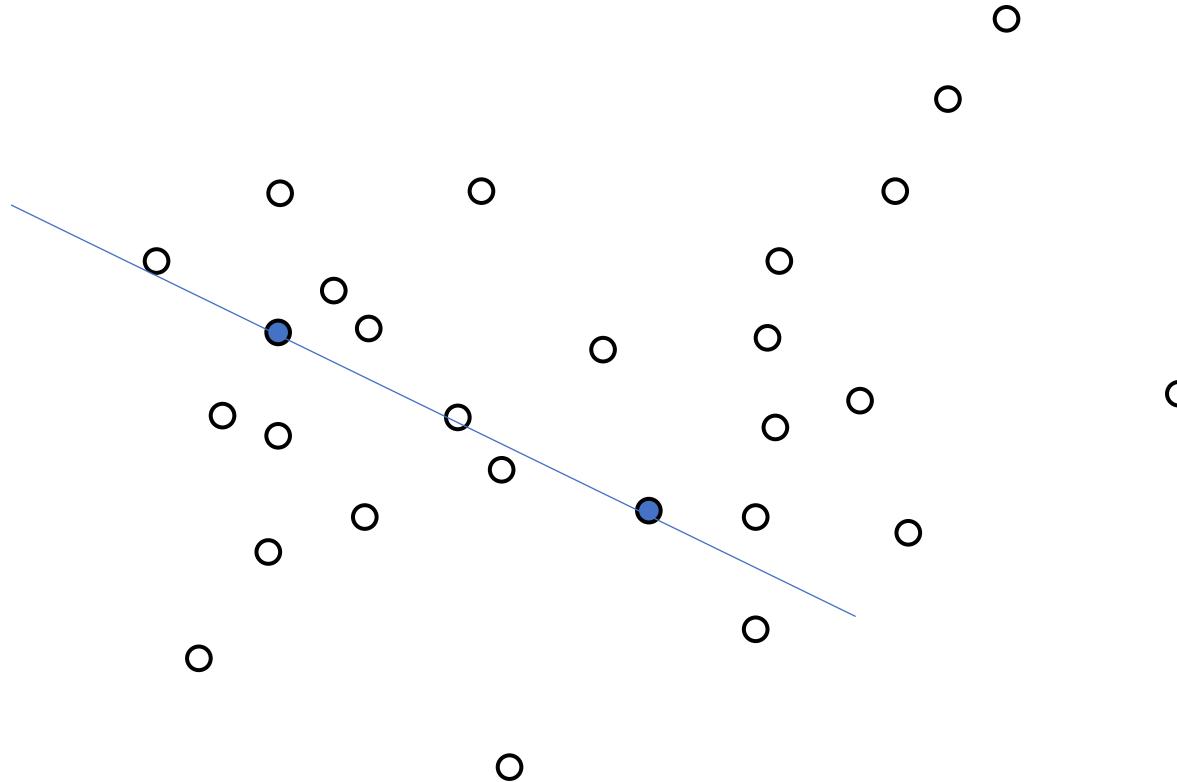
RANSAC

- Random Search – Much Faster!



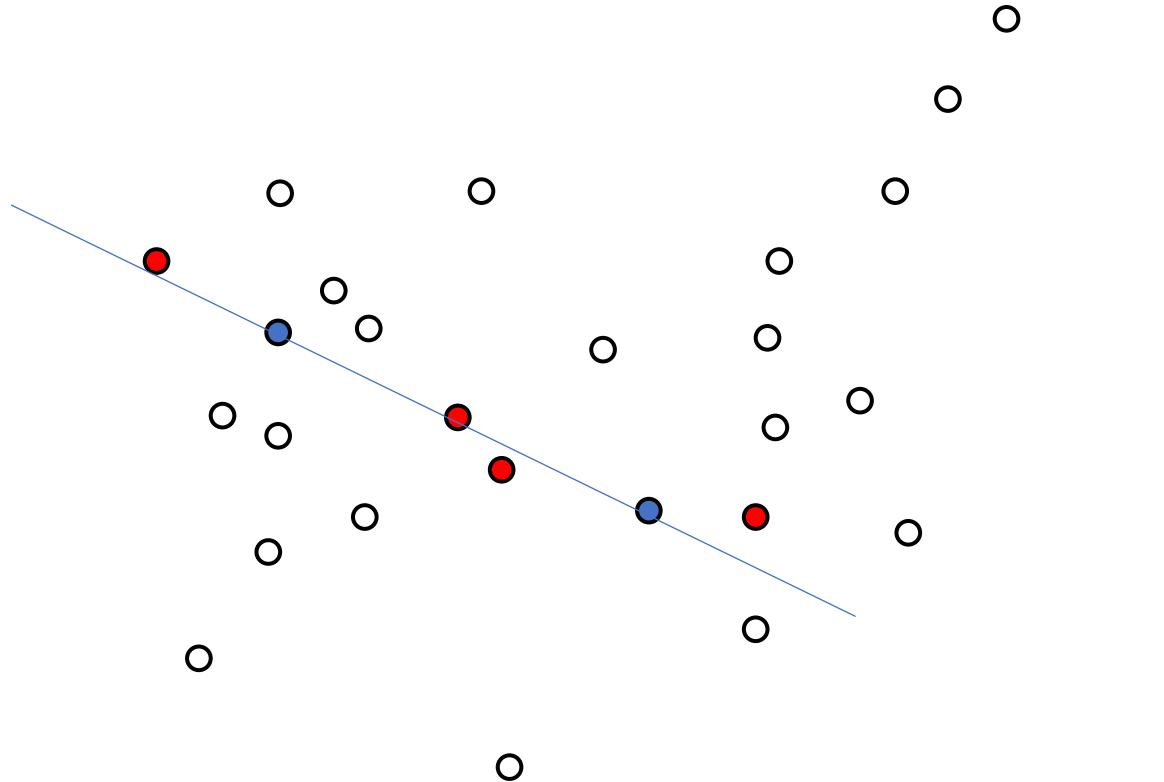
RANSAC

- See how RANSAC works.
- Iteration 1:



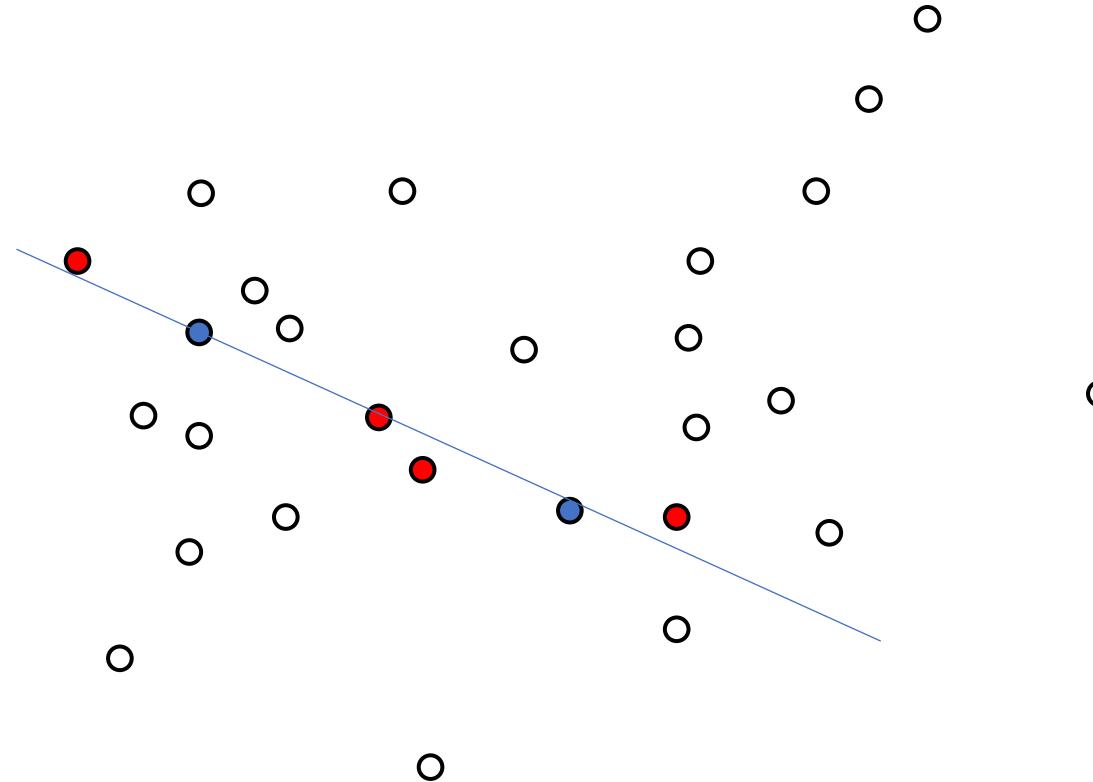
RANSAC

- See how RANSAC works.
- Iteration 1:



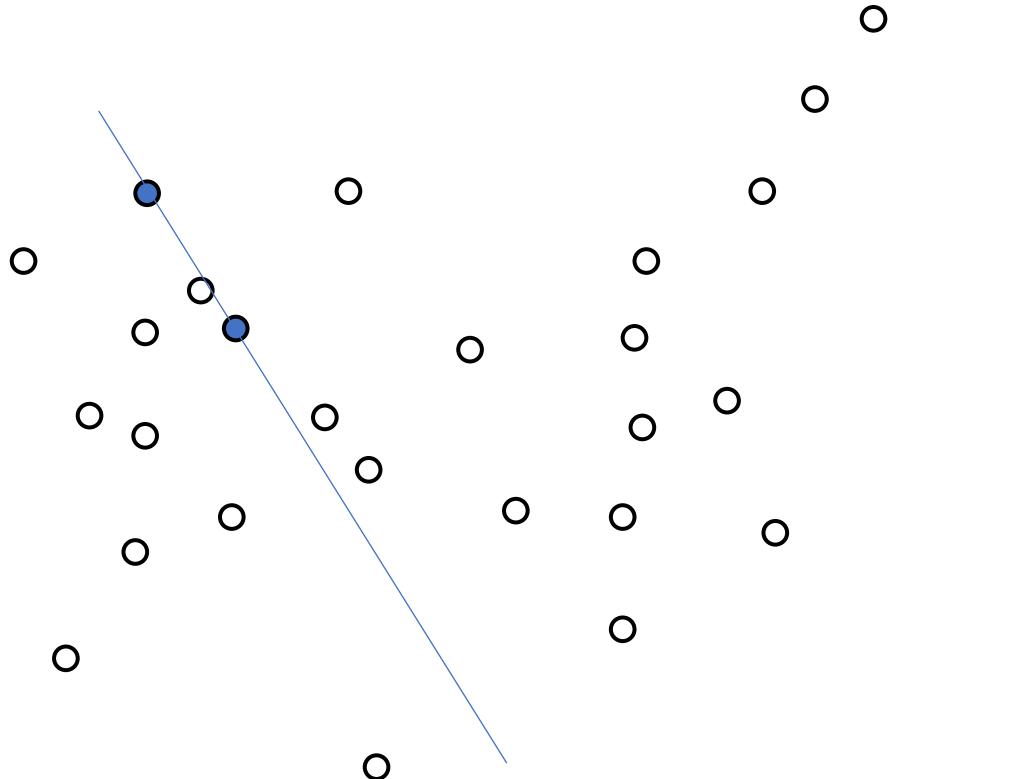
RANSAC

- See how RANSAC works.
- Iteration 1:



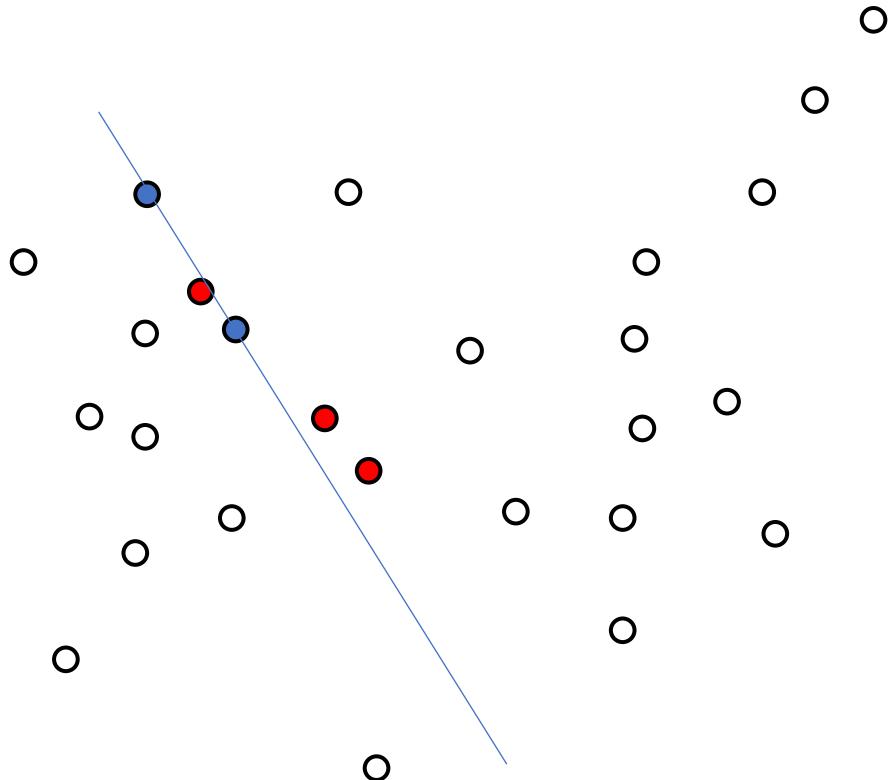
RANSAC

- See how RANSAC works.
- Iteration 2:



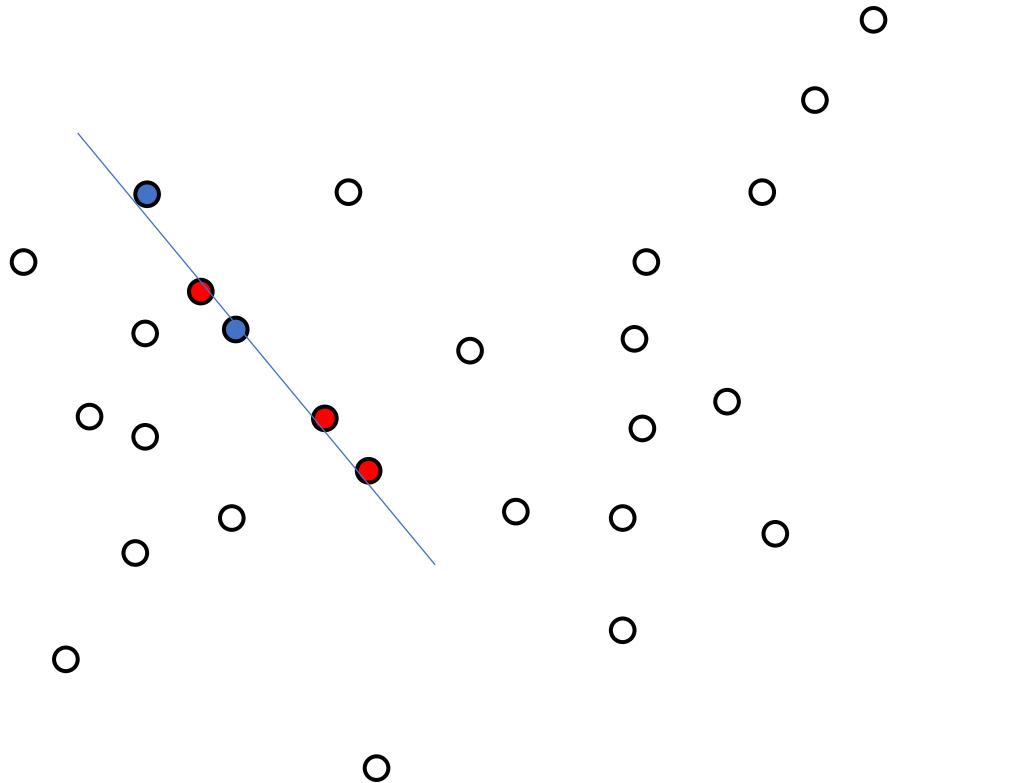
RANSAC

- See how RANSAC works.
- Iteration 2:



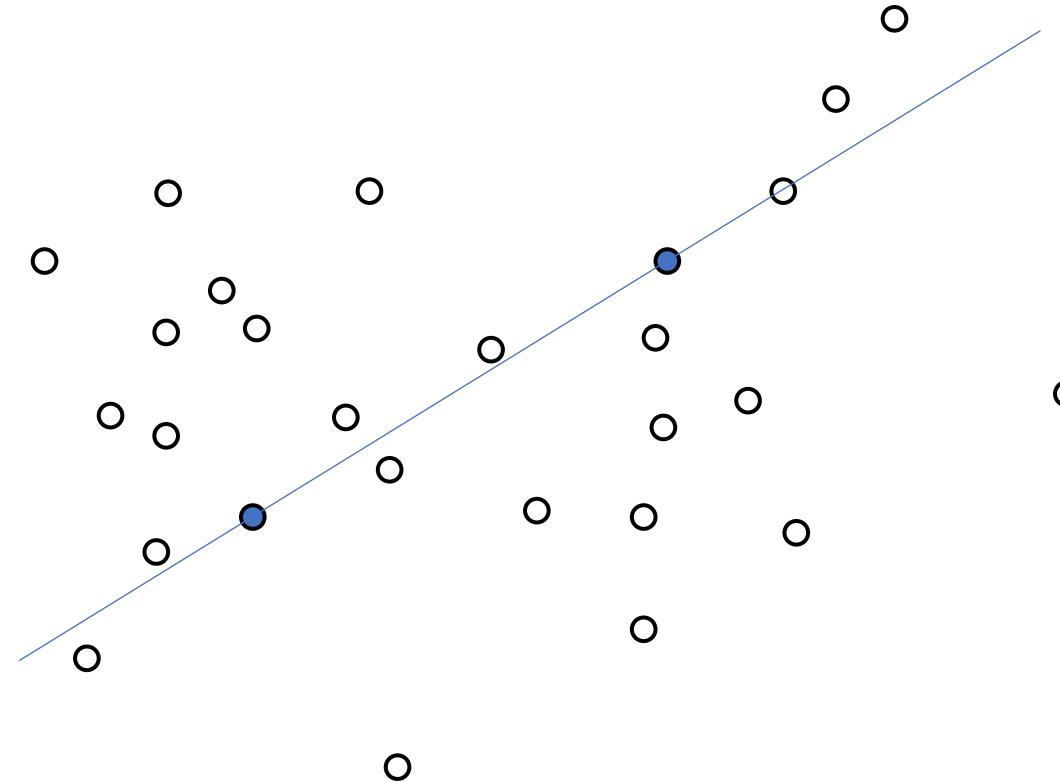
RANSAC

- See how RANSAC works.
- Iteration 2:



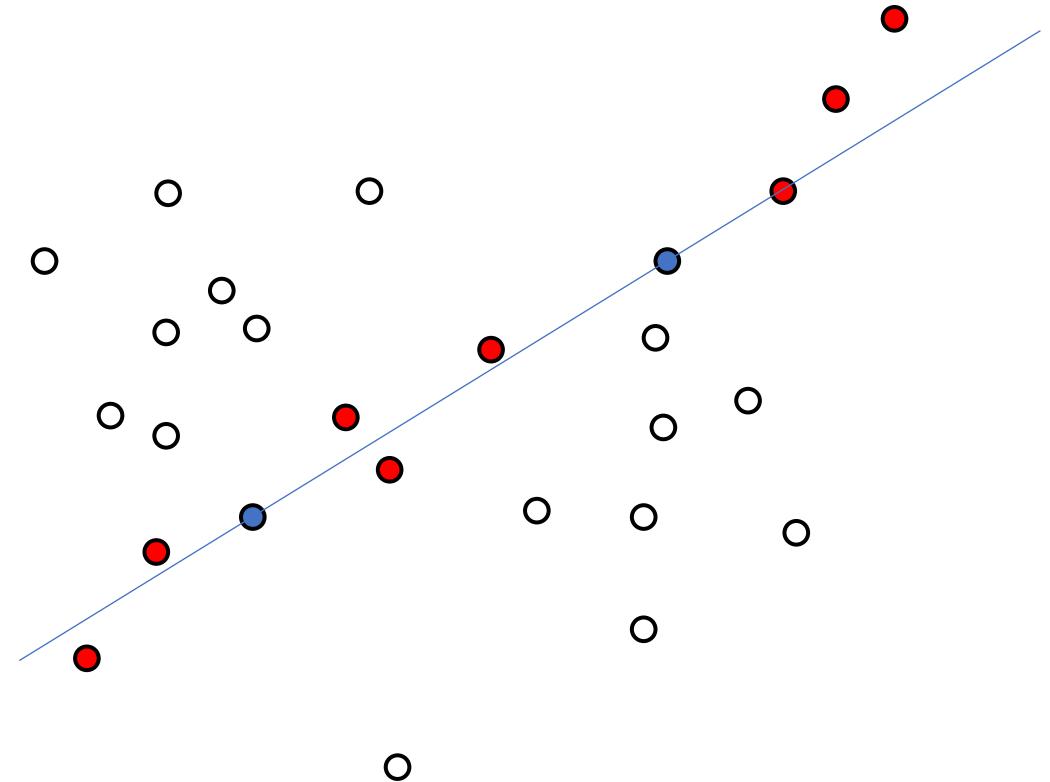
RANSAC

- See how RANSAC works.
-
- Iteration 5:



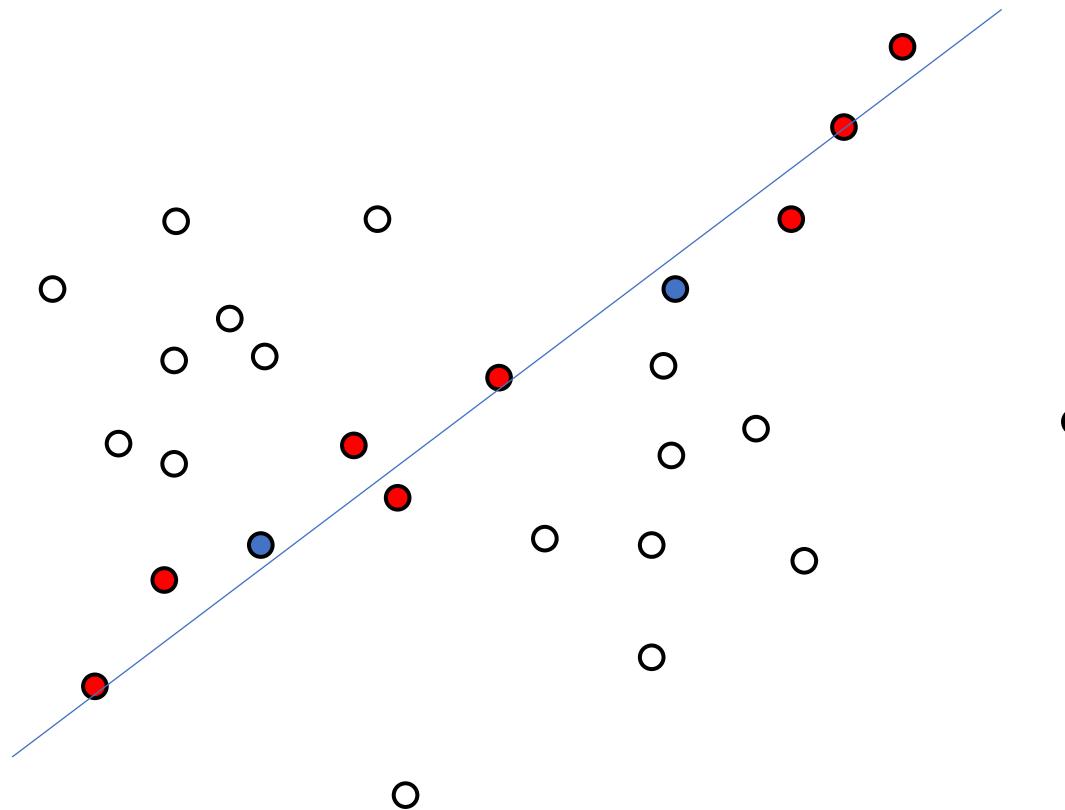
RANSAC

- See how RANSAC works.
- Iteration 5:



RANSAC

- See how RANSAC works.
- Iteration 5:

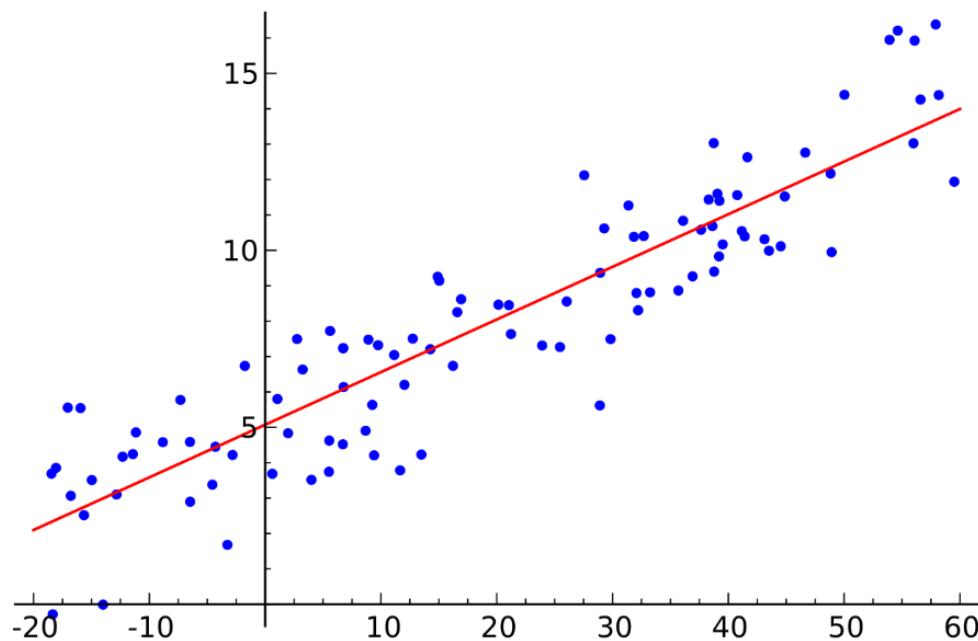


RANSAC

- **Algorithm RANSAC**
- Determine:
 - n – the smallest number of points required (e.g., for lines, $n = 2$, for circles, $n = 3$)
 - k – the number of iterations required
 - t – the threshold used to identify a point that fits well
 - d – the number of nearby points required to assert a model fits well
- Until k iterations have occurred
 - Draw a sample of n points from the data uniformly and at random
 - Fit to that set of n points
 - For each data point outside the sample
 - Test the distance (fitting error) from the point to the structure
 - If the distance is less than t , the point is close (called inlier)
 - If there are d or more points close to the structure
 - Then there is a good fit. Refit the structure using all inliers. Add the result to a collection of good fits.
- Use the best fit from this collection (using the fitting error as a criterion)

RANSAC

- RANSAC will distinguish inliers with outliers.
- A detail: how to fit multiple inliers after RANSAC?
 - Least Squares



Use Least-Squares to fit a line with much more points than minimal need (2 points)

Model Fitting

- Now we can describe the general process of model fitting:
- 1. Give a parametric model (e.g. plane, line, sphere)
- 2. Use RANSAC to find inliers from a point cloud
- 3. Use Least-Squares to fit the model to the inliers
- 4. Take the parameters and the inliers as output

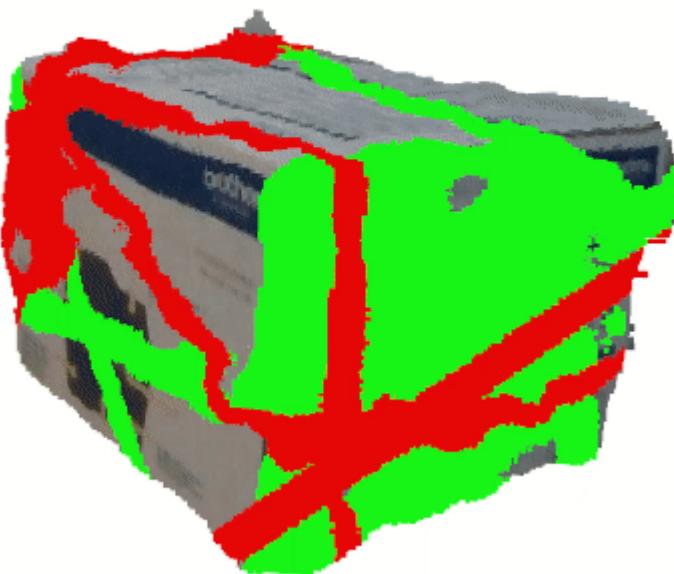
Model Fitting

- Visualization: <https://github.com/leomariga/pyRANSAC-3D/tree/Animations>
 - Red: current trials
 - Green: currently best trial
 - Circle:



Model Fitting

- Visualization: <https://github.com/leomariga/pyRANSAC-3D/tree/Animations>
 - Red: current trials
 - Green: currently best trial
 - Cuboid:



Summary

- Registration
 - PCA
 - SVD
 - ICP
- Surface Reconstruction
 - Delaunay Triangulation
 - Poisson Surface Reconstruction
- Model Fitting
 - RANSAC
 - Least Squares