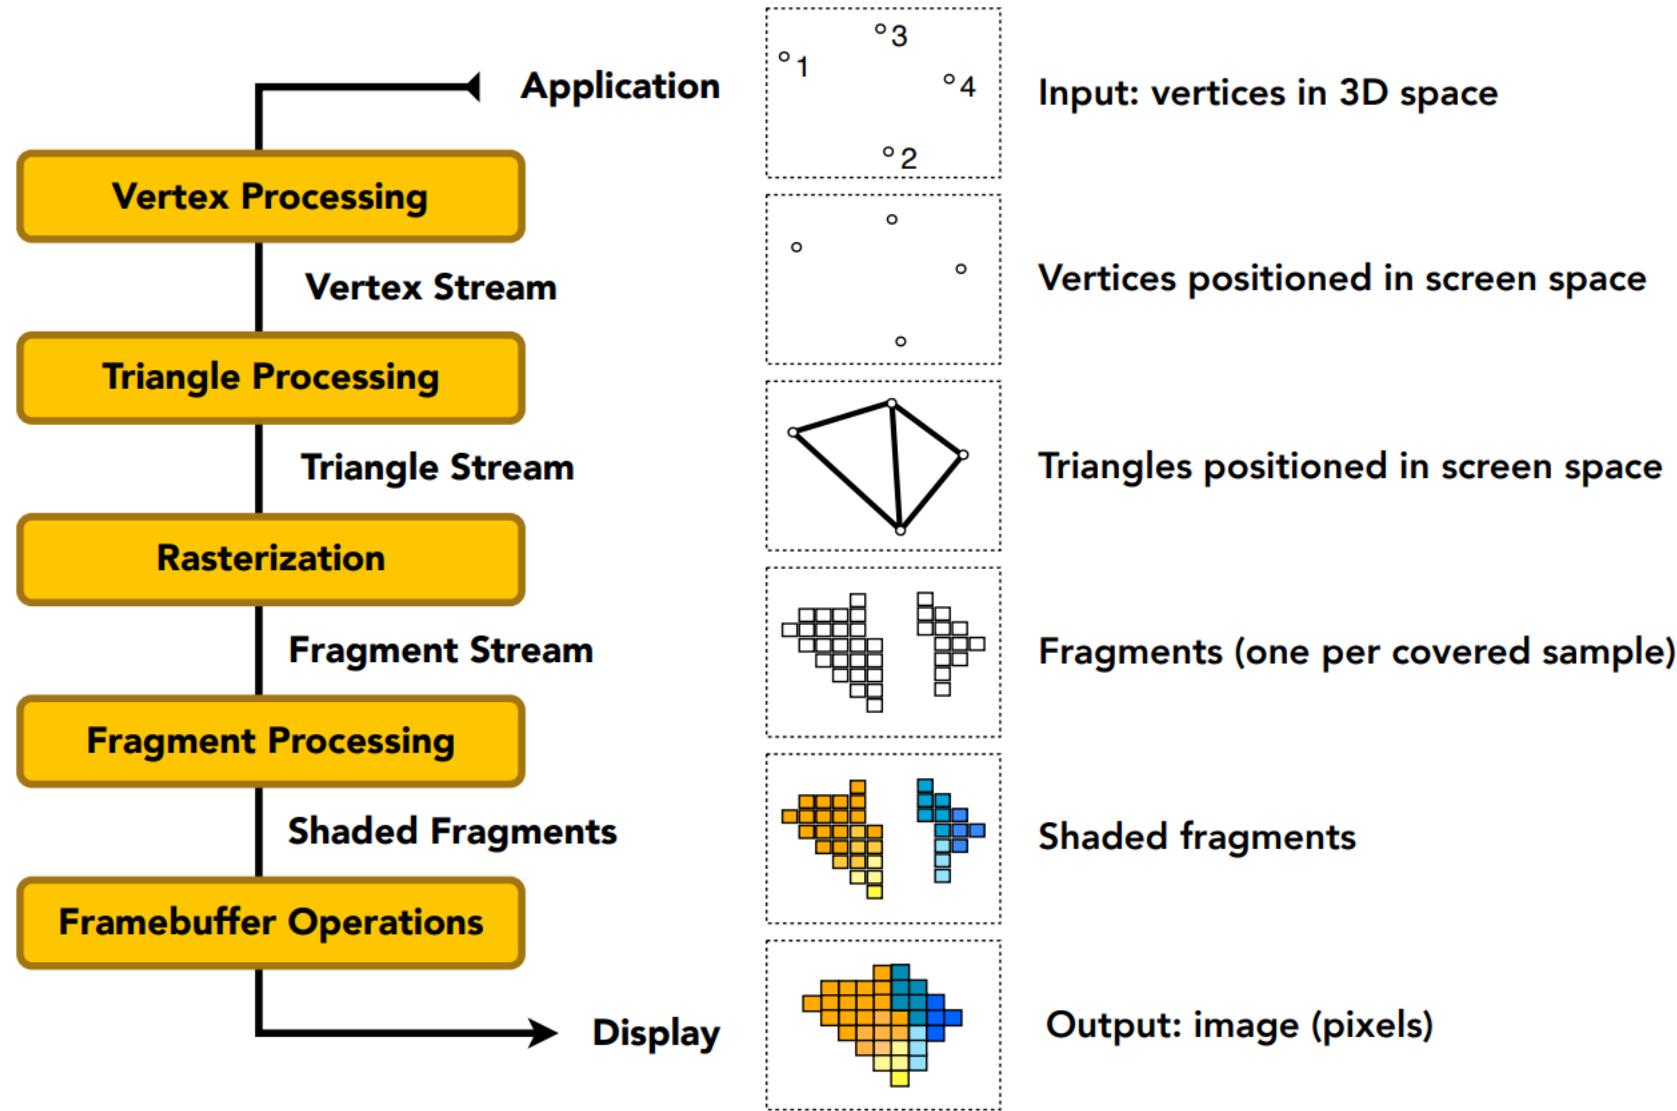


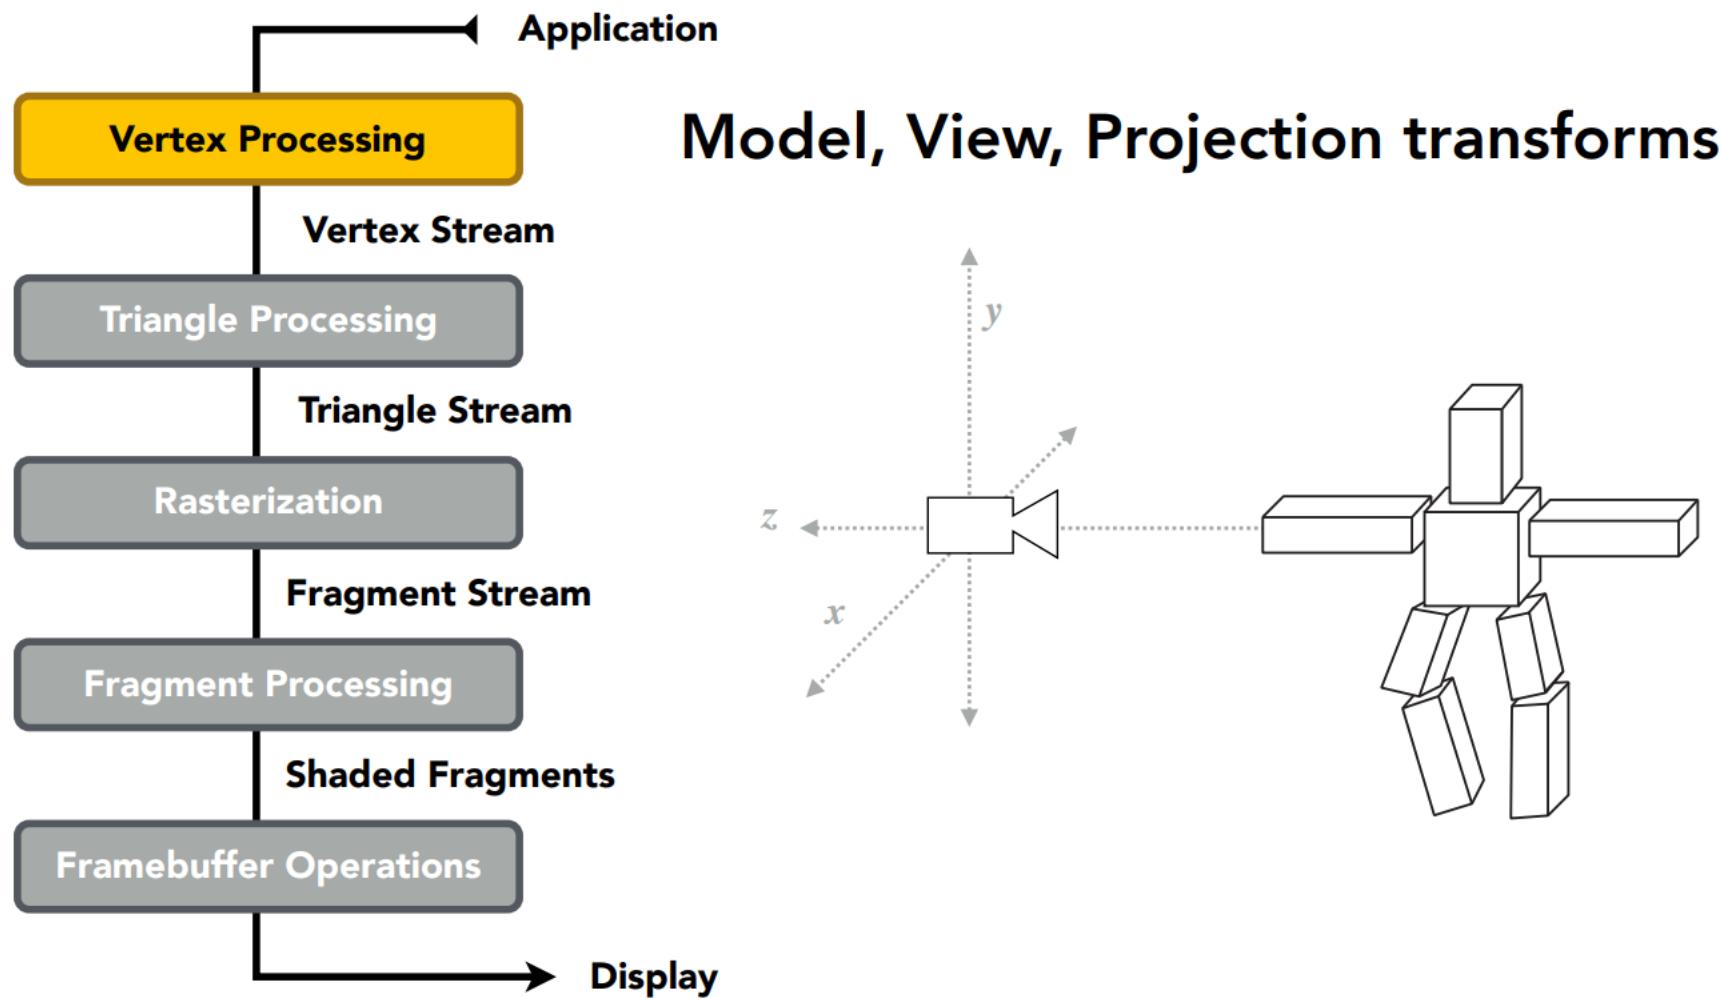
Graphics Pipeline

Raster Rendering

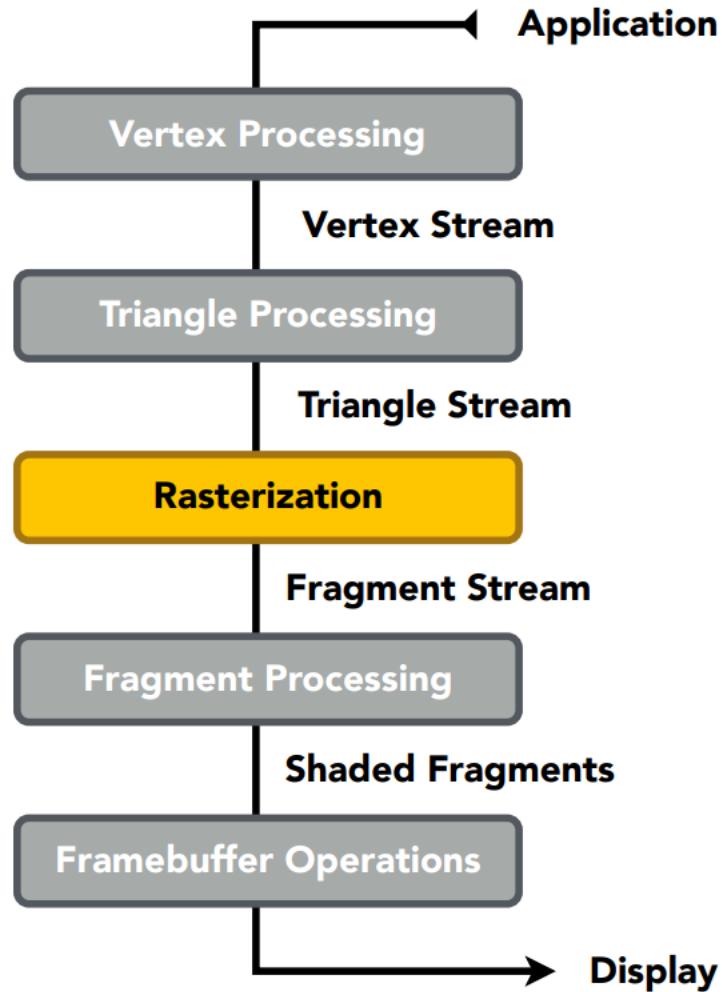
Rendering Pipeline



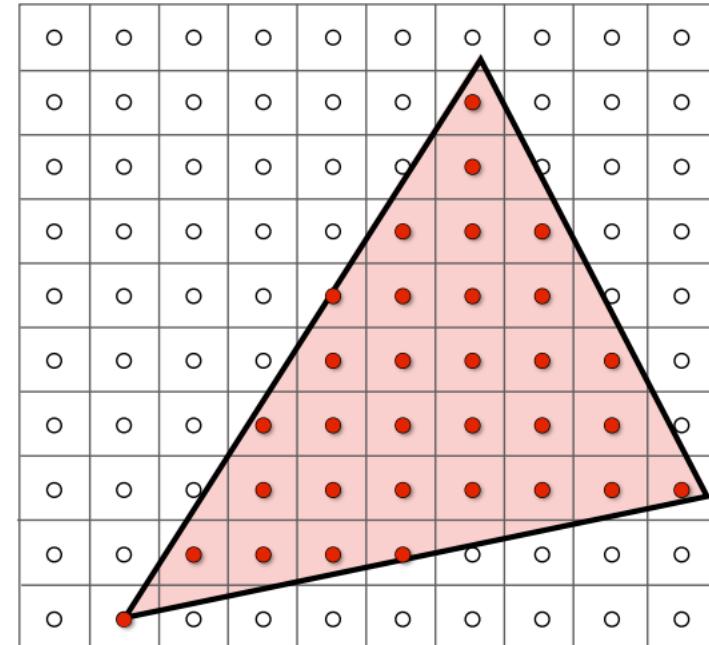
Rendering Pipeline



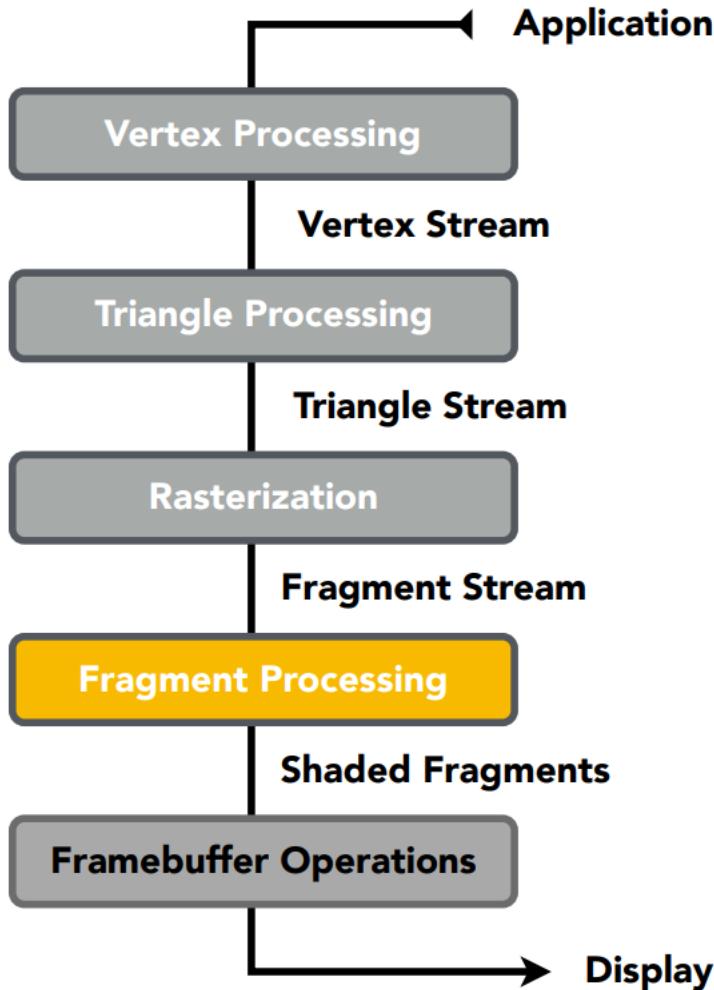
Rendering Pipeline



Sampling triangle coverage



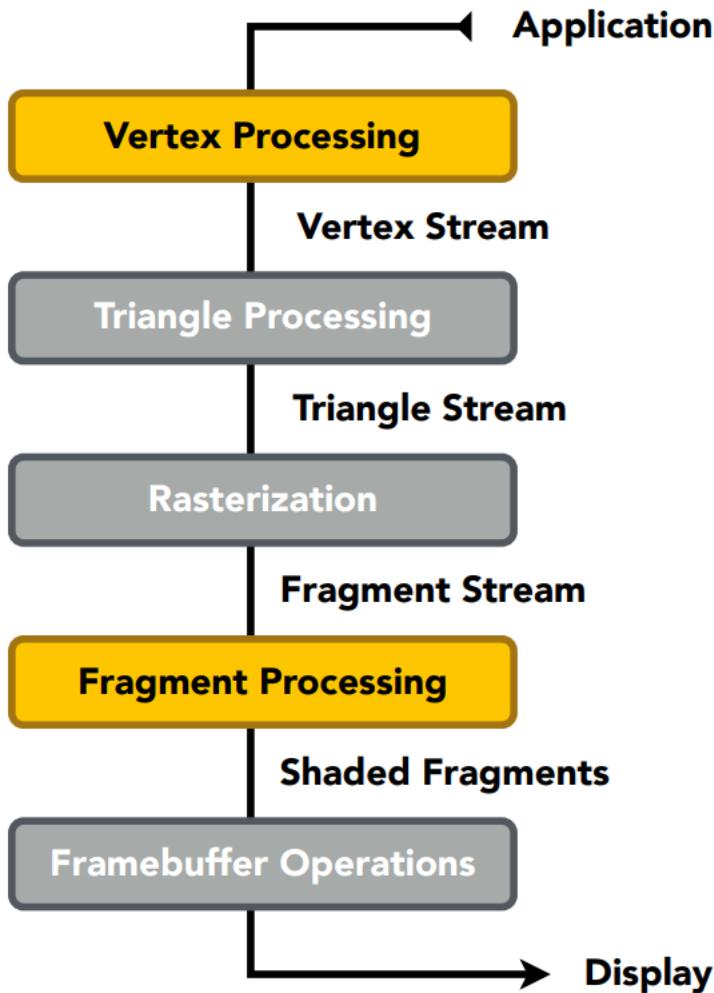
Rendering Pipeline



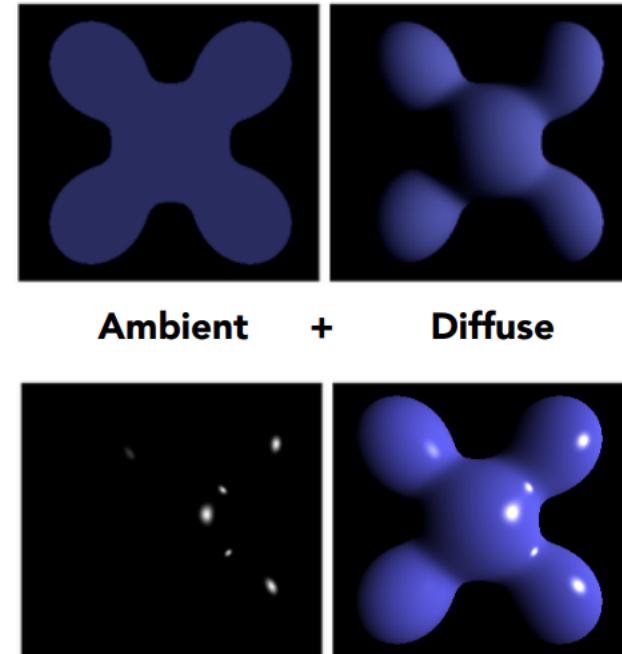
Z-Buffer Visibility Tests



Rendering Pipeline

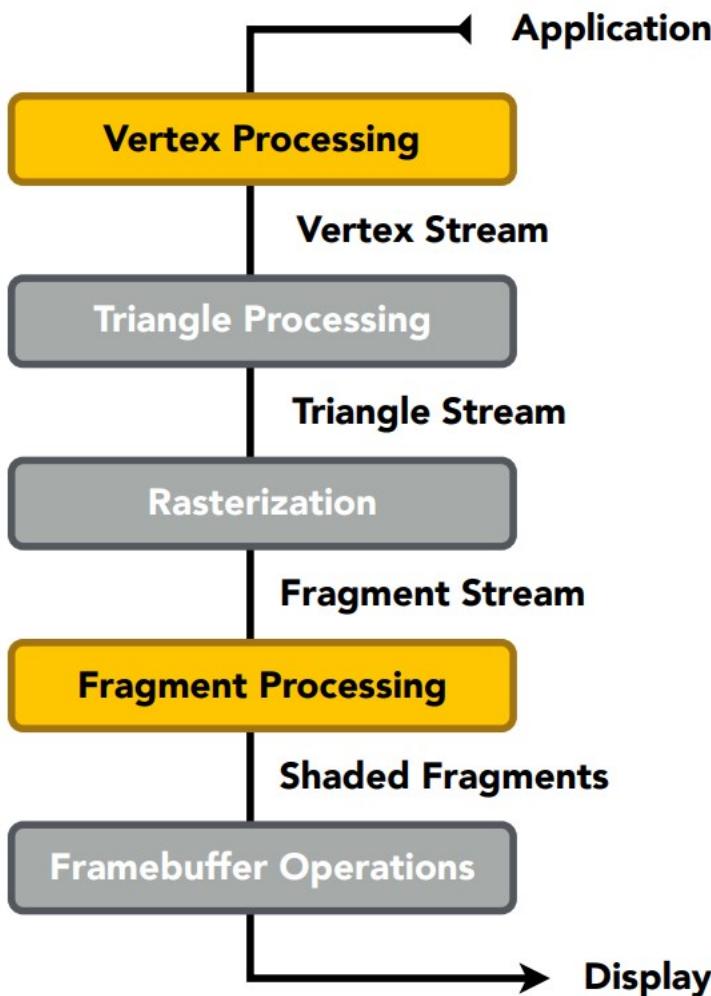


Shading

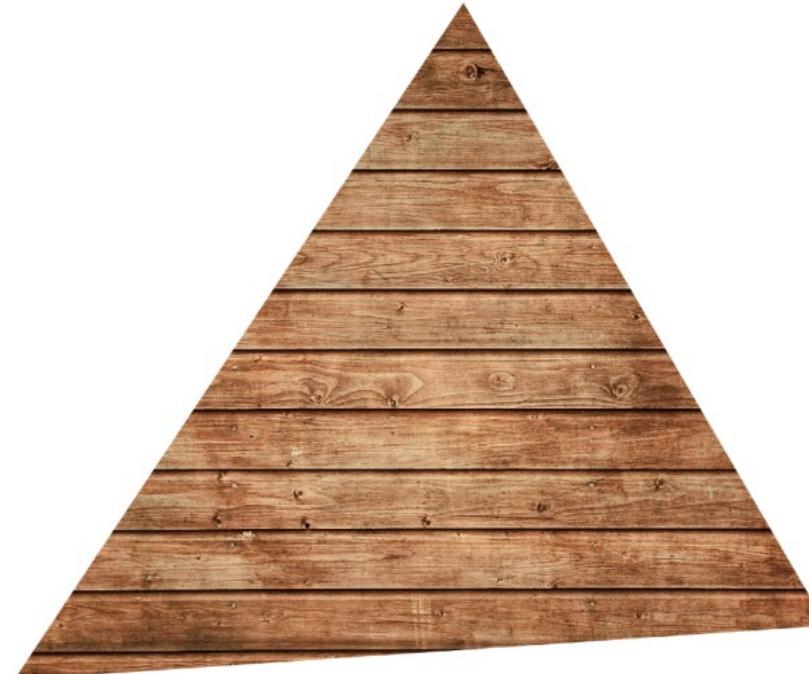


+ Specular = Blinn-Phong
Reflectance Model

Rendering Pipeline

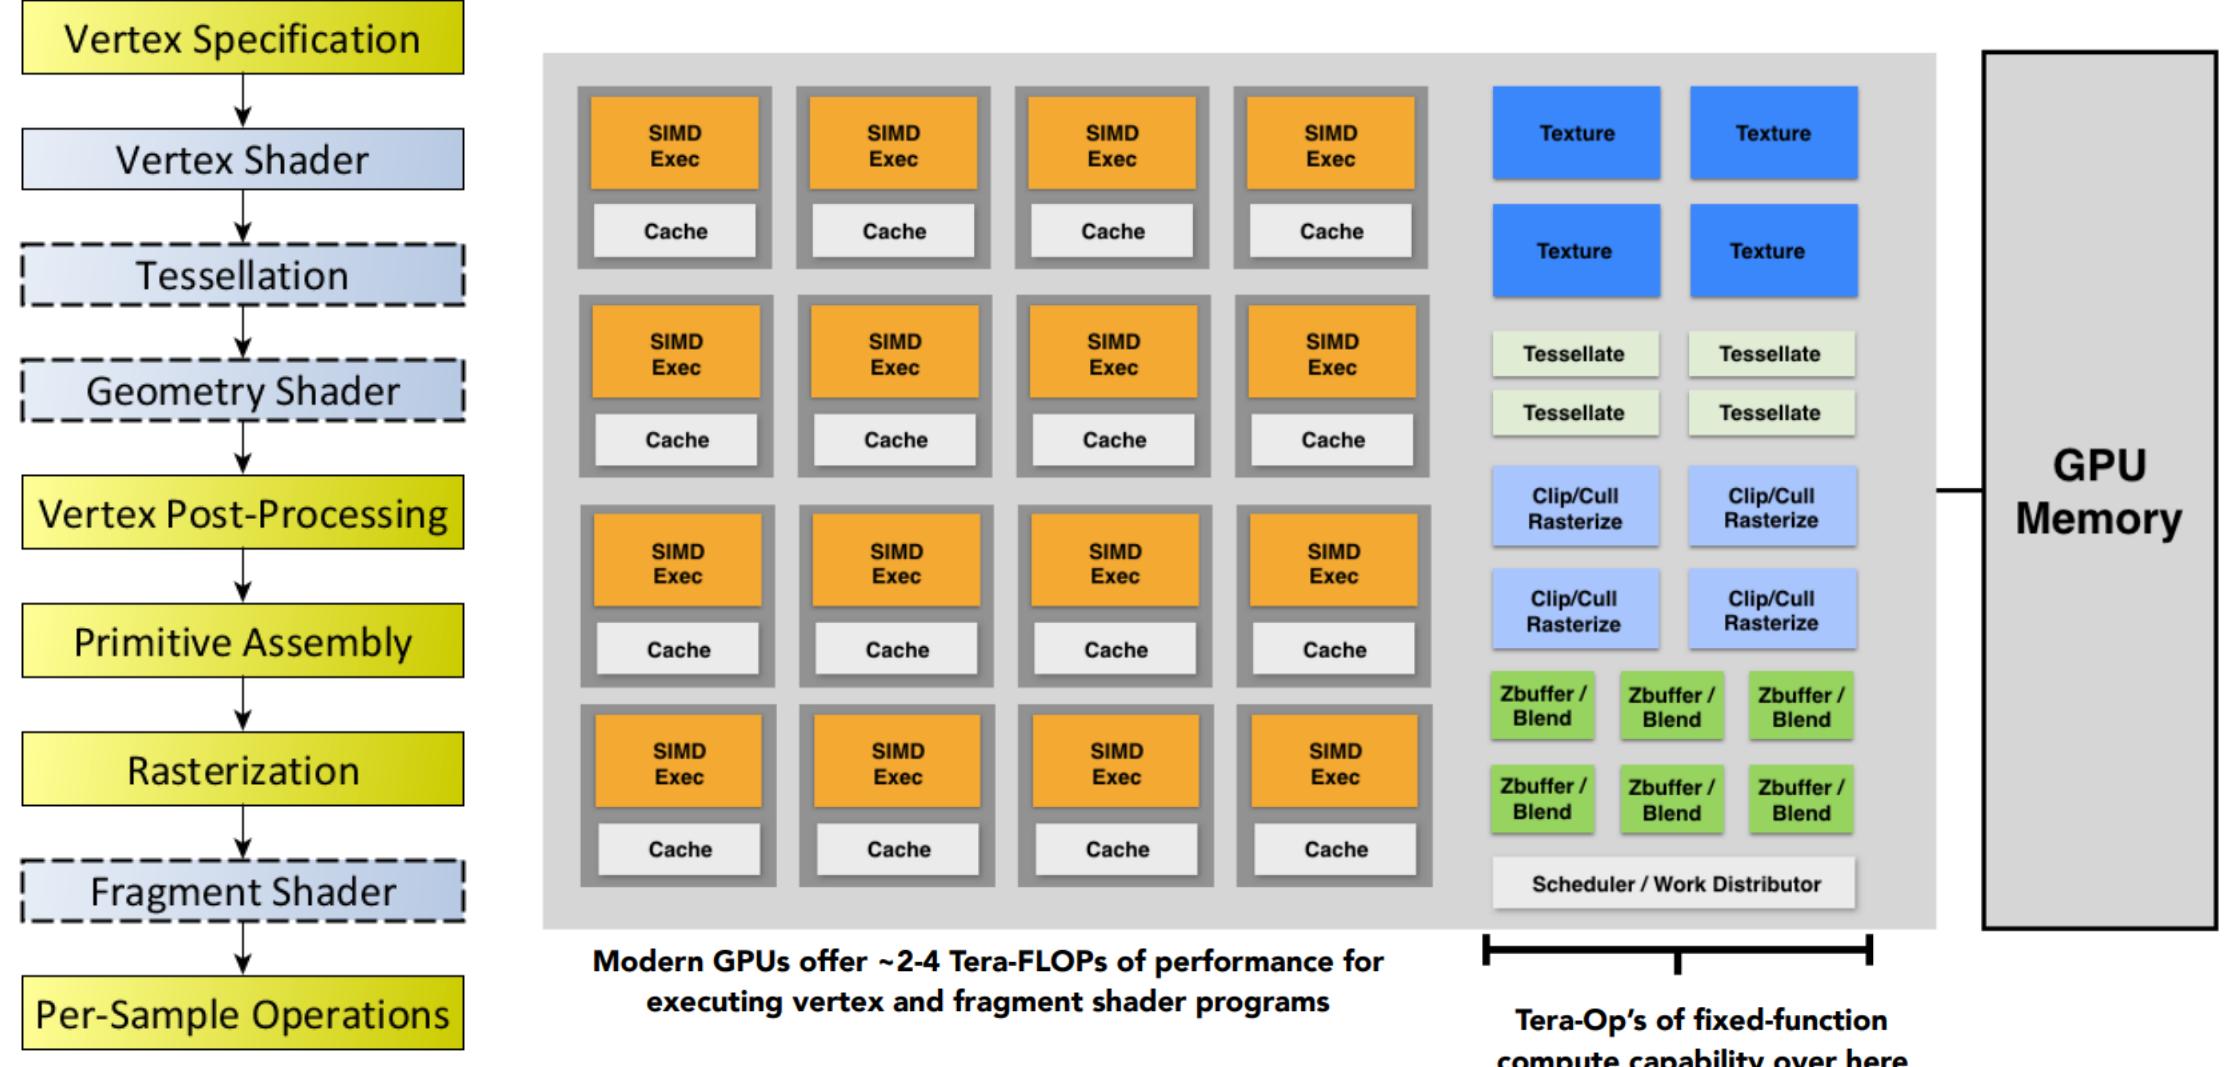


**Texture mapping
(introducing soon)**



Graphics API

Modern GPU



Programming on GPU



OpenGL, a cross-language, cross-platform API for 2D and 3D graphics.



WebGL, a JavaScript API for 2D/3D graphics within any compatible web browser without the use of plug-ins.



Metal, created by Apple



DirectX 12 Ultimate, the API that features the most advanced features of NVIDIA GPU.

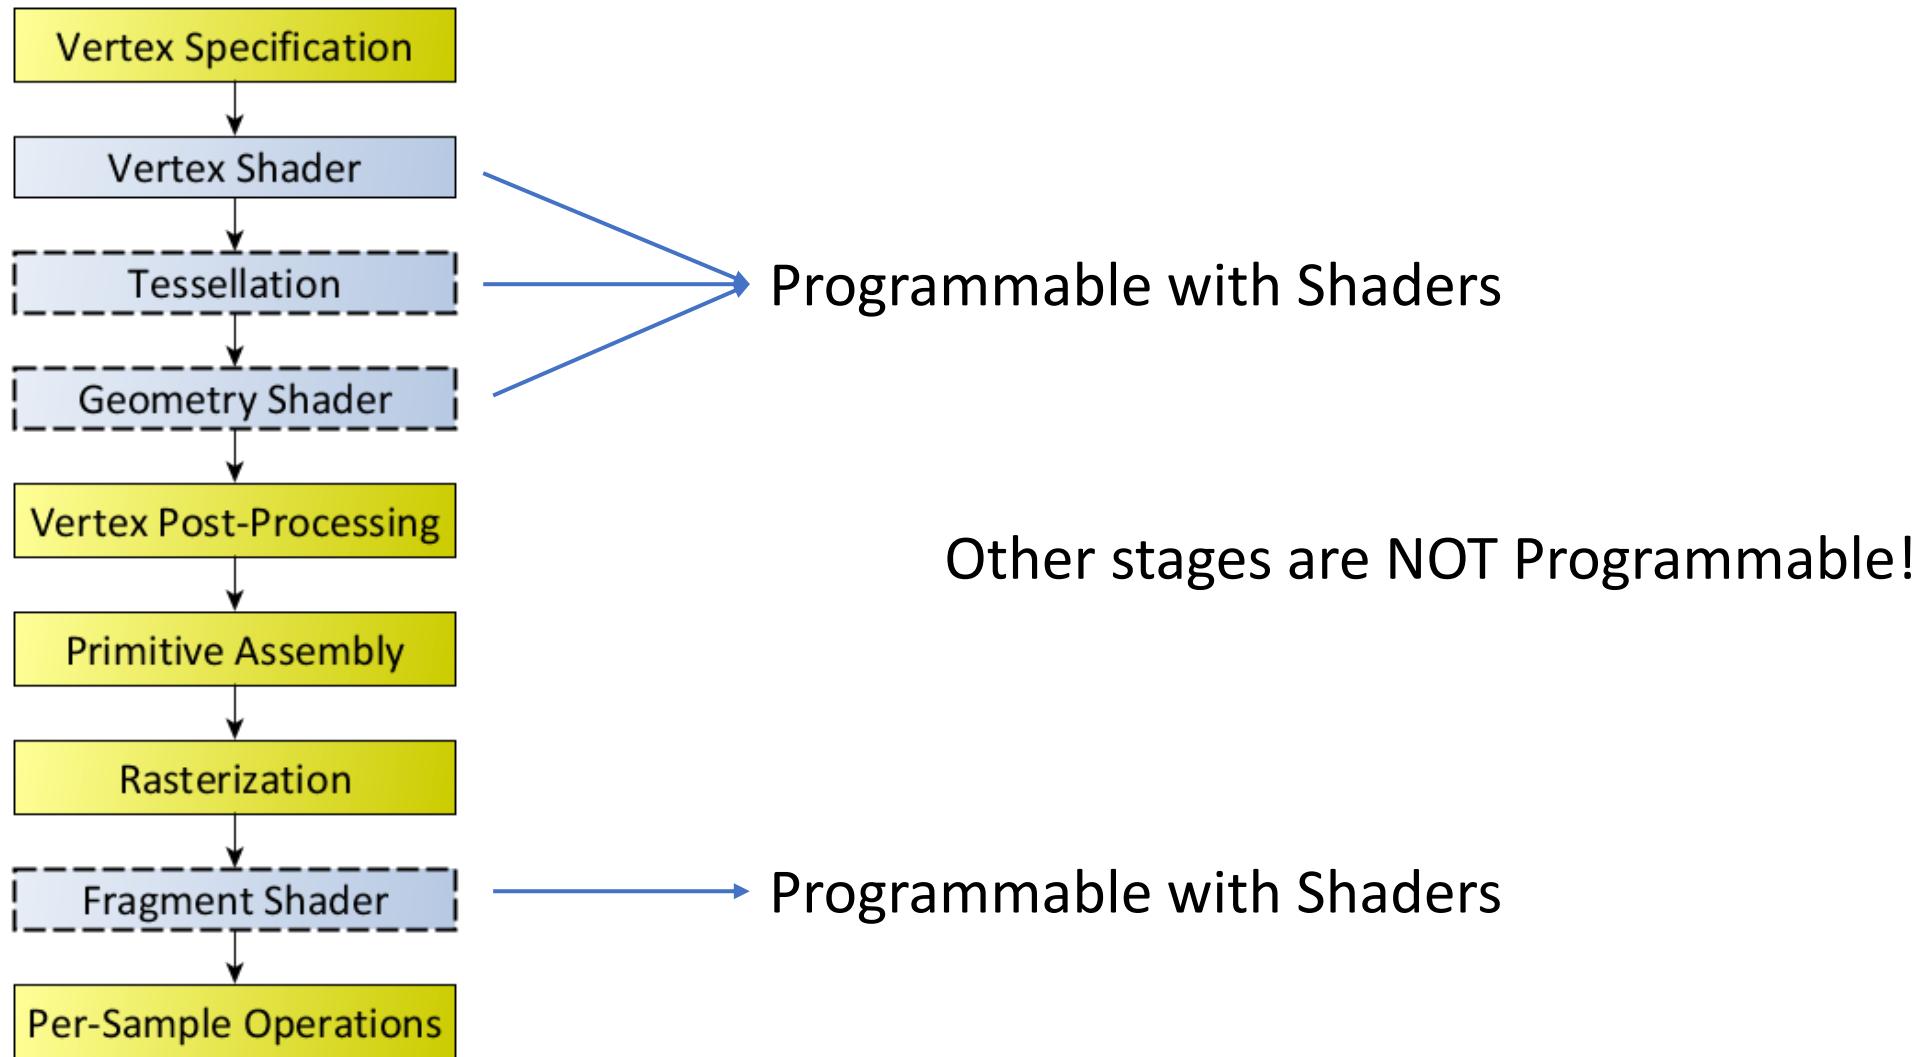


OpenGL ES, a subset of OpenGL API on embedded and mobile systems



Vulkan, built upon components of AMD's Mantle API

Programming on GPU



Shader

- Vertex shader: parallel over all vertices
 - Input some attributes of each vertex, like position, normal, color...
 - Output some attributes of each vertex, like normal, color, depth...
- Fragment shader: parallel over all pixels of all segments
 - Input values per pixel: interpolated from vertex shader output
 - Output pixel color, depth...
- Shader language: GLSL (OpenGL), HLSL (DirectX), Houdini VEX (Houdini)...

GLSL Shader Example: Phong Shading

Vertex shader

```
#version 330 core
layout (location = 0) in vec3 aPos;
layout (location = 1) in vec3 aNormal;

out vec3 FragPos;
out vec3 Normal;

uniform mat4 model;
uniform mat4 view;
uniform mat4 projection;

void main()
{
    FragPos = vec3(model * vec4(aPos, 1.0));
    Normal = mat3(transpose(inverse(model))) * aNormal;

    gl_Position = projection * view * vec4(FragPos, 1.0);
}
```

Fragment shader

```
#version 330 core
out vec4 FragColor;

in vec3 Normal;
in vec3 FragPos;

uniform vec3 lightPos;
uniform vec3 viewPos;
uniform vec3 lightColor;
uniform vec3 objectColor;

void main()
{
    // ambient
    float ambientStrength = 0.1;
    vec3 ambient = ambientStrength * lightColor;

    // diffuse
    vec3 norm = normalize(Normal);
    vec3 lightDir = normalize(lightPos - FragPos);
    float diff = max(dot(norm, lightDir), 0.0);
    vec3 diffuse = diff * lightColor;

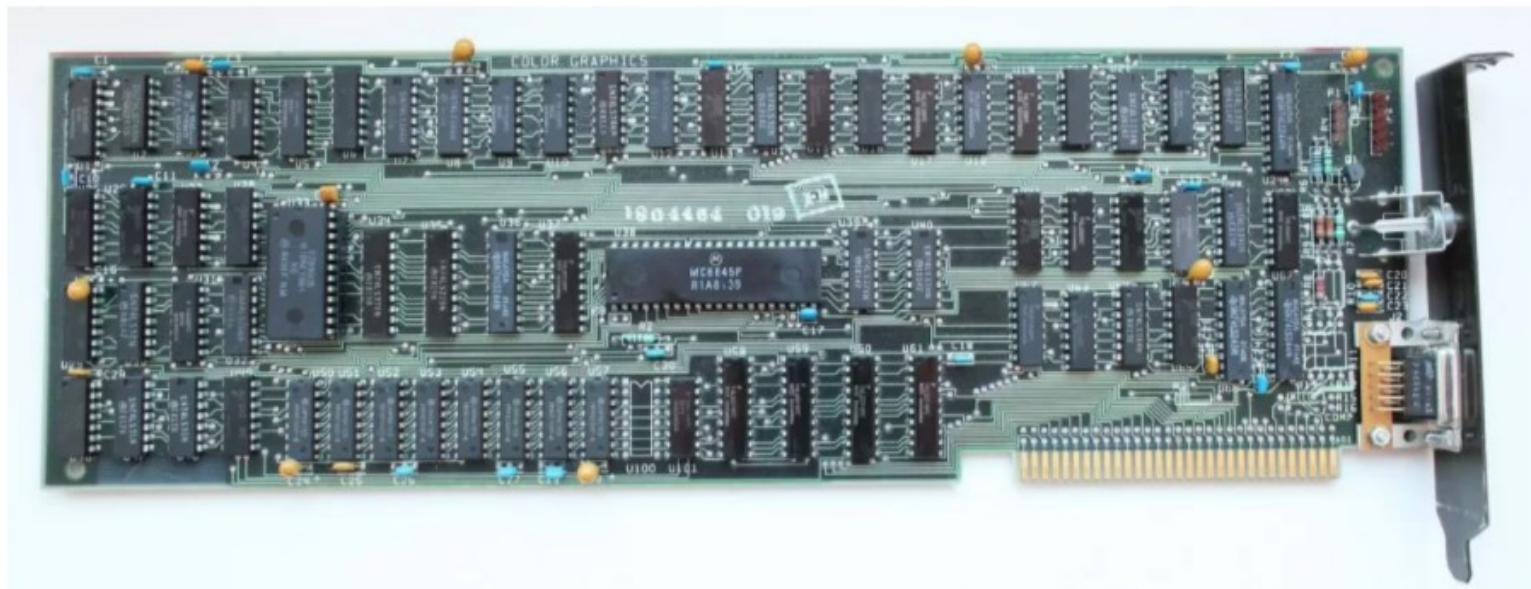
    // specular
    float specularStrength = 0.5;
    vec3 viewDir = normalize(viewPos - FragPos);
    vec3 reflectDir = reflect(-lightDir, norm);
    float spec = pow(max(dot(viewDir, reflectDir), 0.0), 32);
    vec3 specular = specularStrength * spec * lightColor;

    vec3 result = (ambient + diffuse + specular) * objectColor;
    FragColor = vec4(result, 1.0);
}
```

GPU History

- 1976-1984: only 2D hardware acceleration, 3D on CPU
- Mainly display adapters, video accelerators...

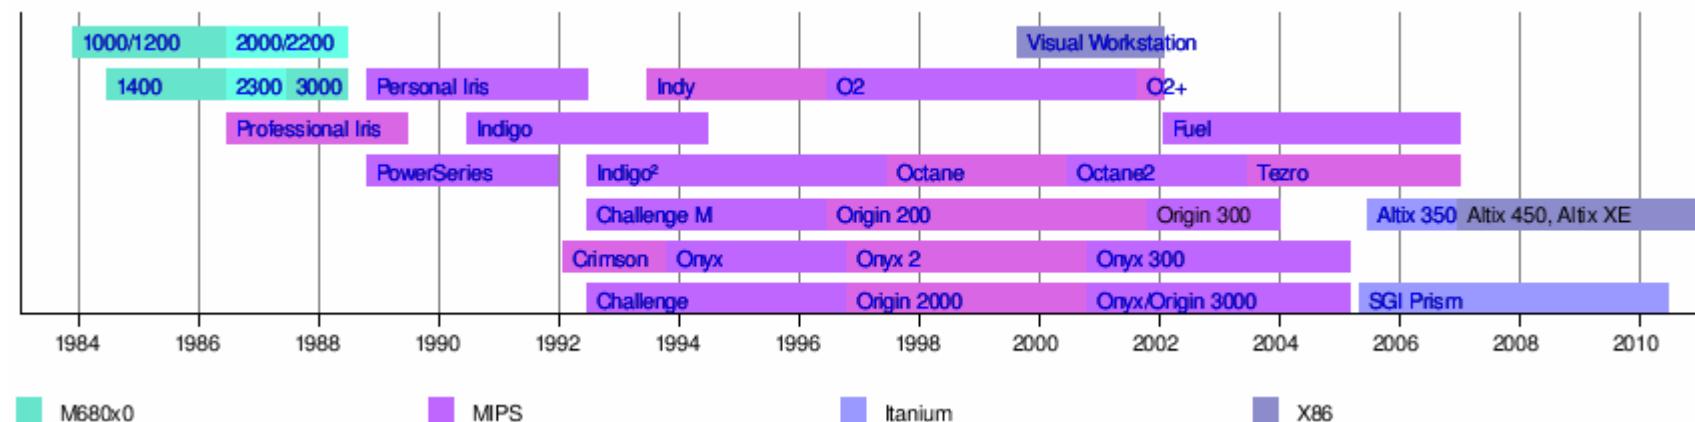
IBM PC's Monochrome Display Adapter



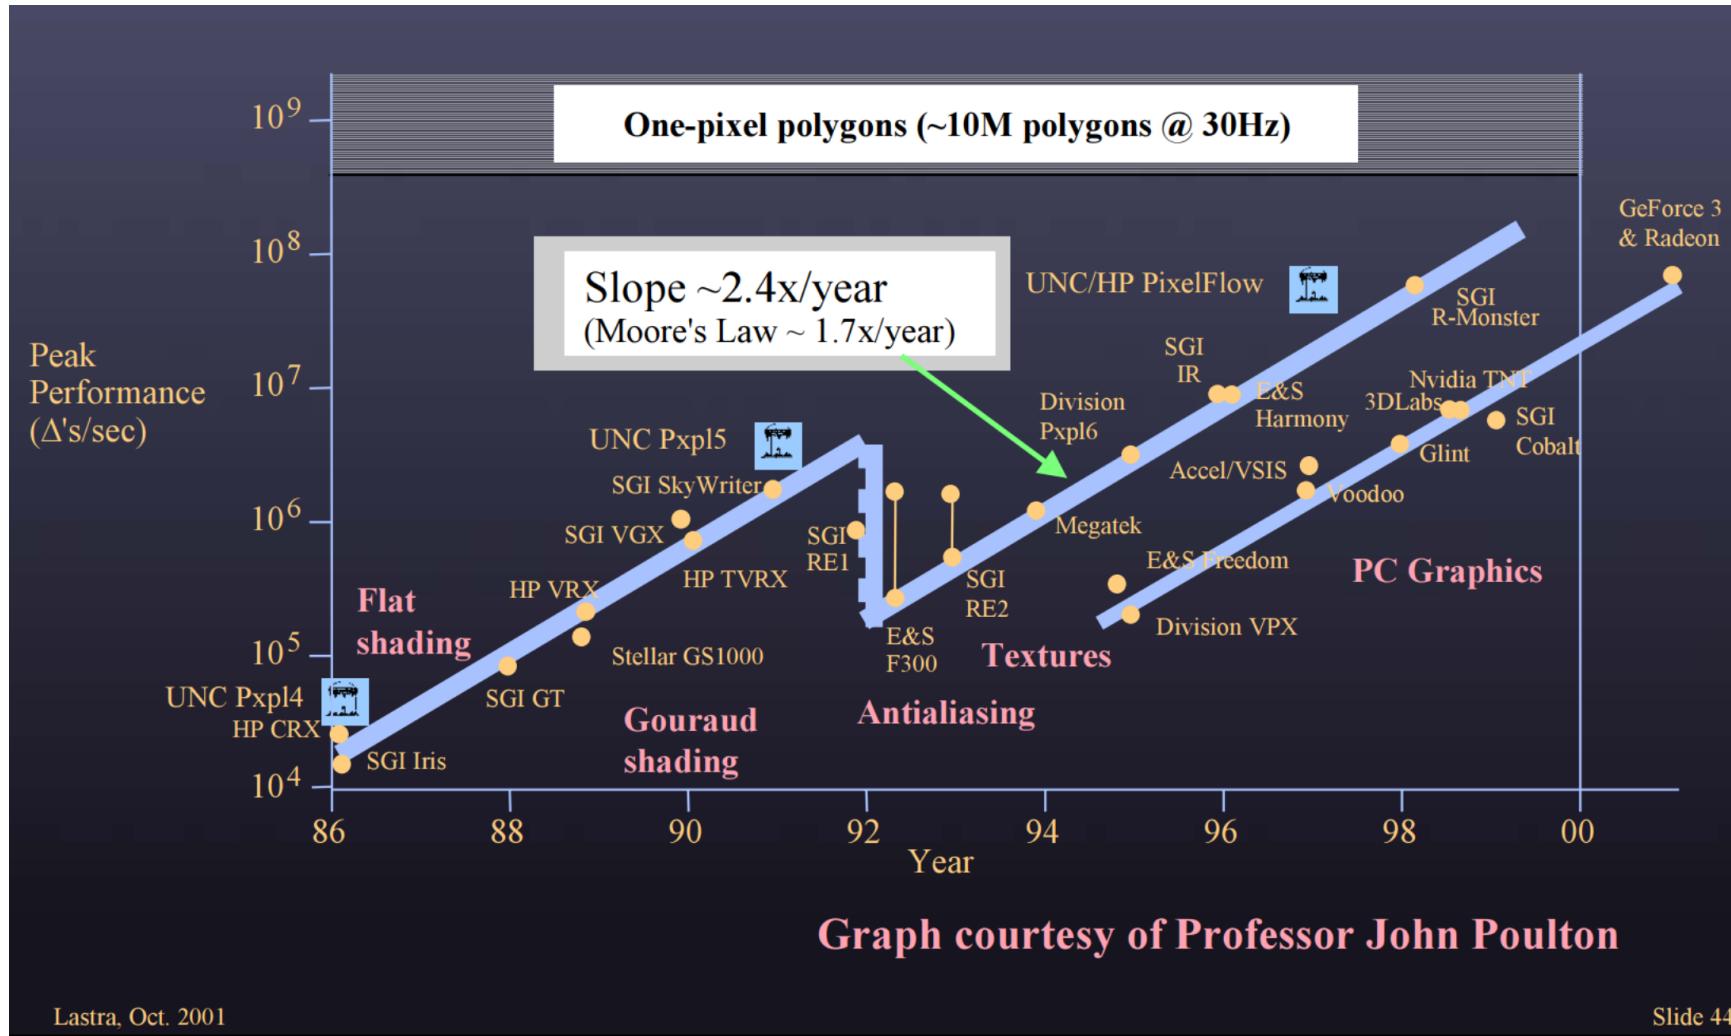
GPU History

- 1984-2005: Silicon Graphics (SGI) workstation

sgi®	
Type	Public
Traded as	NYSE: SGI OTC Pink: SGID.pk Nasdaq: SGIC
Industry	Computer hardware and software
Founded	November 9, 1981; 40 years ago Mountain View, California, U.S. ^[1]
Defunct	May 11, 2009; 12 years ago
Fate	Chapter 11 bankruptcy; assets acquired by Rackable Systems, which renamed itself Silicon Graphics International Corp.
Headquarters	Sunnyvale, California, U.S.
Key people	Jim Clark Wei Yen Kurt Akeley Ed McCracken Thomas Jermoluk Marc Hannah



Graphics Hardware



GPU History

- 1995-1999: 3Dfx VOOODOO, 3D hardware acceleration



Fast forward: GLQuake released in 1997 versus original Quake

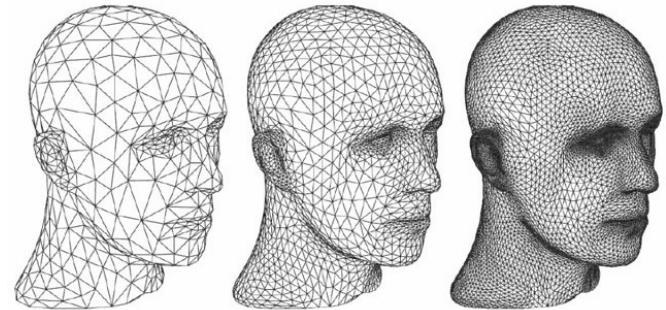
- 2000-2006: Nvidia vs ATI
- Programmable pipeline
 - vertex/fragment shader
 - 2002: HLSL by DirectX 9, GLSL by OpenGL
 - 2006: Geometry Shader
- 2006: ATI bought by AMD



Nvidia's stock GeForce 3

GPU History

- 2006-2013: general purpose GPU (GPGPU)
- 2007: CUDA by Nvidia
- 2009: tessellation shader (subdivision),
compute shader

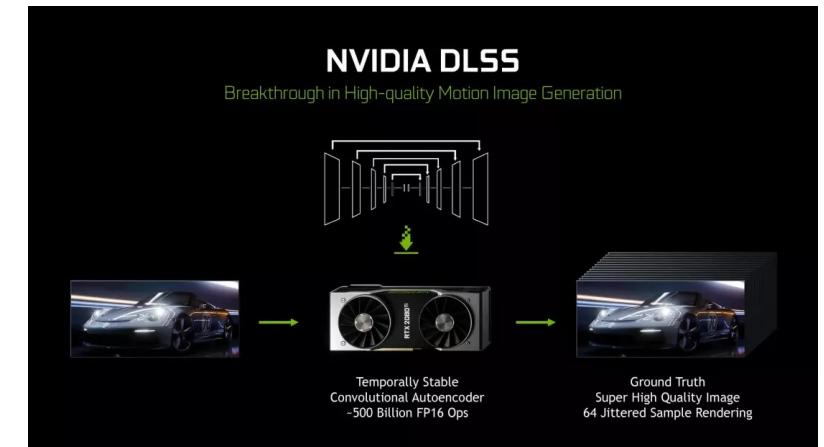
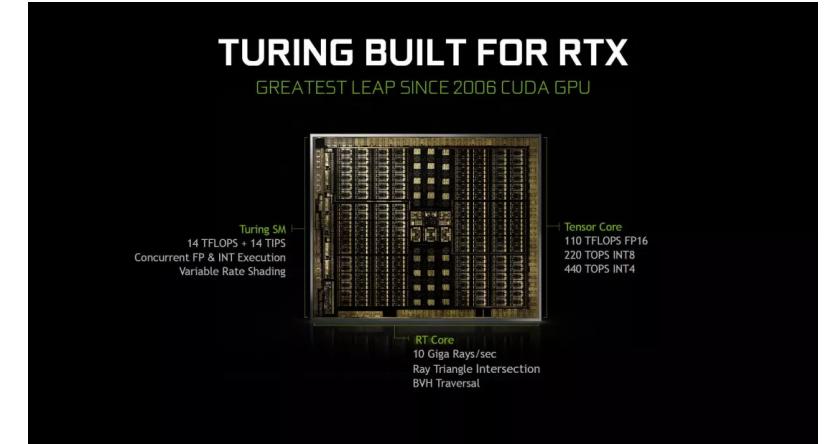


MSI's version of the GeForce 8800 GTX

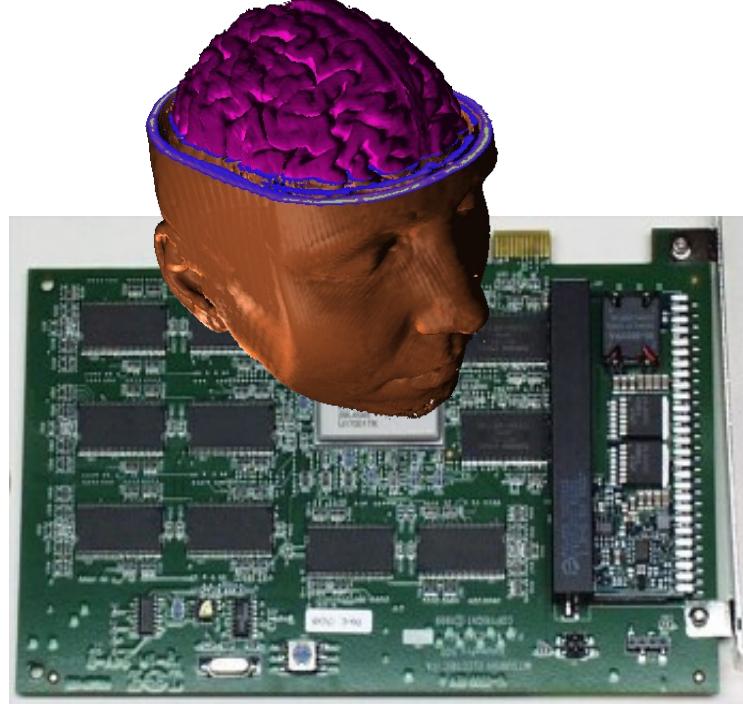
GPU History

- 2013-now: faster and faster, for gamers and AI
- Nvidia RTX: hardware ray tracing
- Nvidia DLSS: deep-learning super sampling
- GPGPU: scientific computing,

Crypto-mining, AI training...

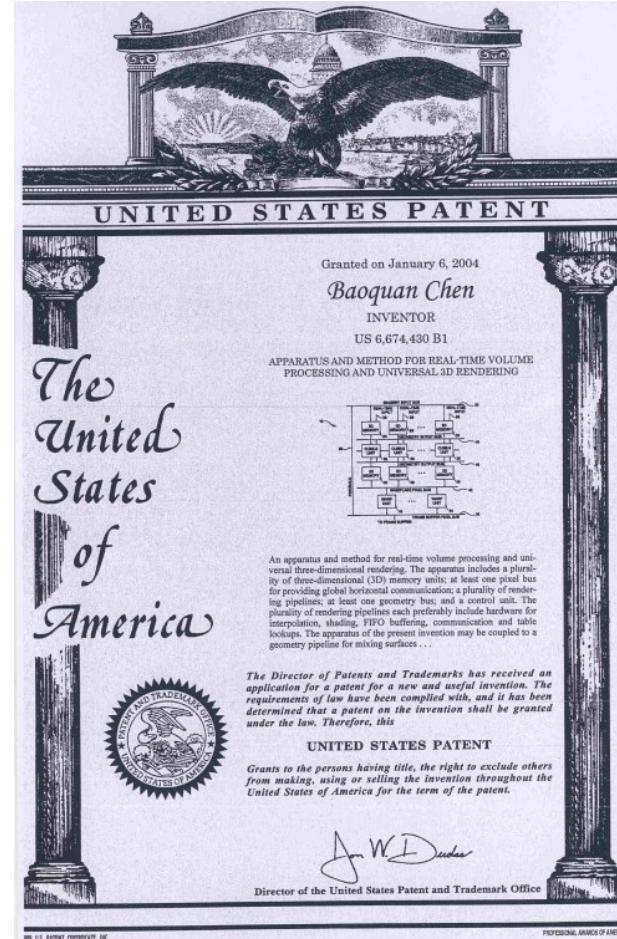


VolumePro: Special purpose hardware for volume rendering

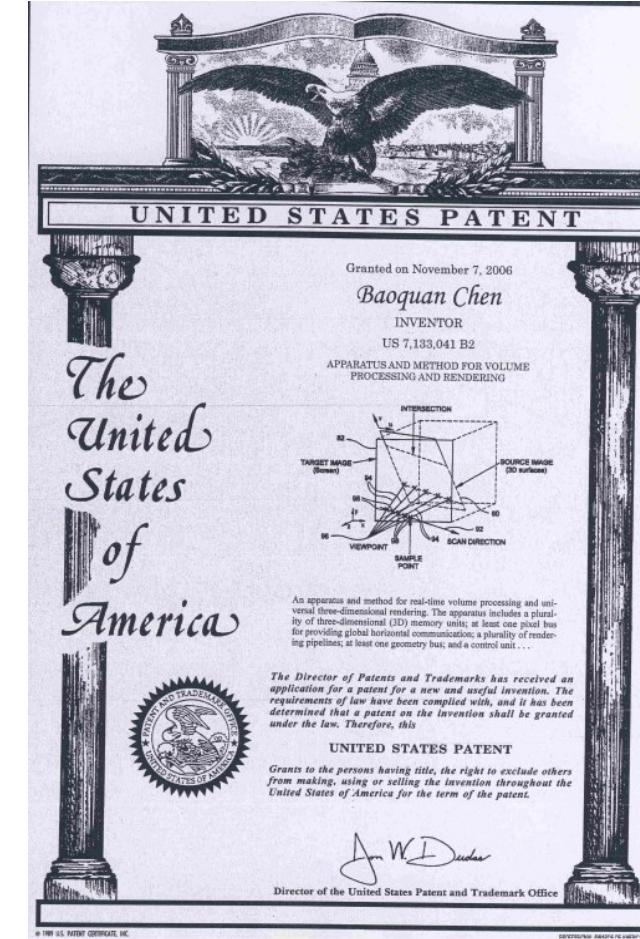


第五代Cube 芯片设计与实现

- VolumePro (, 1999)
- Cube Engine ( *Japan Radio Co., Ltd.*), 2000



Patents licensed to:



Accelerated Rendering

Forward Rendering

- Assume N triangles and M light sources in the shading stage

For N triangles:

For all pixels:

For M light sources:

$\text{color} += \text{shading}(\text{triangle}, \text{light})$

Test z-buffer and update

Complexity $O(N*M)$! Many shading results are overdrawn and discarded.

Early-Z

- Decide whether to perform shading by current z-buffer

For N triangles:

For pixels not occluded:

For M light sources:

 color += shading(triangle, light)

Test z-buffer and update

If triangles are sorted by depth, the rendering time can be significantly reduced

No effect in the worst case!

Deferred Rendering

- Convert world-space scene to screen-space G-Buffer

For N triangles:

Test depth and update G-Buffer

For all pixels:

For M light sources:

color += shading(G-Buffer, light)

Complexity O(N+M); Reduced overdrawning

High bandwidth overhead (G-Buffer can be large)

Computed G-Buffer can be used in post-processing

Deferred Rendering

Typical G-Buffer Layout

	R	G	B	A
GB0	Normal (10:10)	Smoothness		MaterialId (2)
GB1		BaseColor		MatData(5)/Normal(3)
GB2	-----	MetalMask	Reflectance	AO
GB3				Radiosity/Emissive

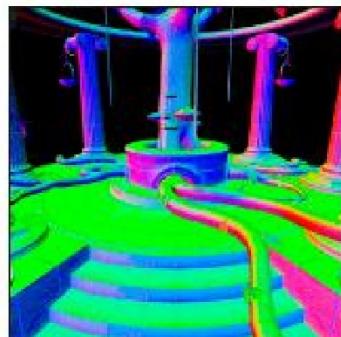
Albedo



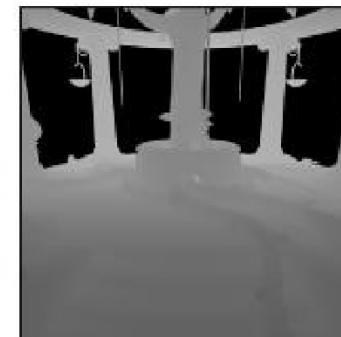
Specular



Normals



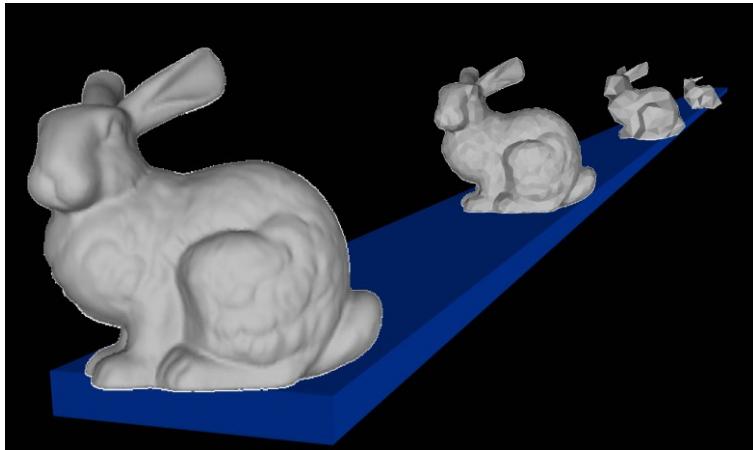
Depth



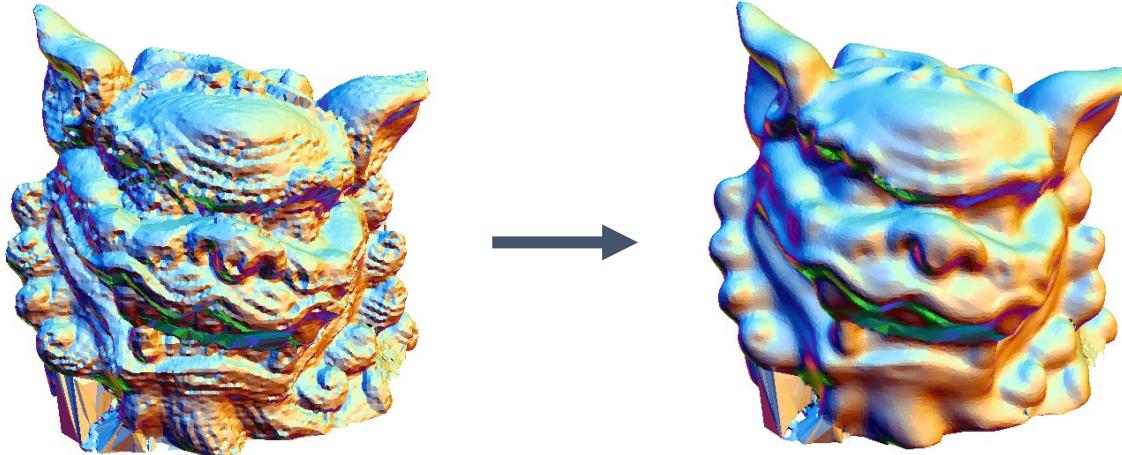
Final render



Mesh Shader



Level of Detail



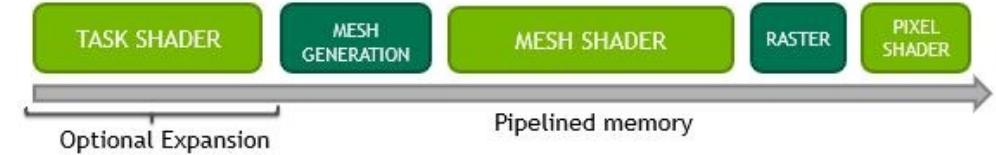
Mesh Smoothing

Image From Nvidia Blog

TRADITIONAL PIPELINE



TASK/MESH PIPELINE



- From vertex shader to mesh shader
- Allow for more flexible geometry processing

Dynamic Global Illumination



Thanks