

Global Illumination I

Whitted-Style Ray Tracing

Why Ray Tracing?

- Rasterization couldn't handle **global effects** well
 - (Soft) shadows
 - Light bounces more than once



Soft shadows




Glossy reflection



Indirect illumination

Ray Tracing by Turner Whitted



The first scene through ray tracing




Turner Whitted (right)

From Rasterization to Ray Tracing

- Simple shading (typified by OpenGL, z-buffering, and Phong illumination model) assumes:
 - direct illumination (light leaves source, bounces at most once, enters eye)
 - no shadows (except using shadow buffer)
 - opaque surfaces
 - point light sources (otherwise integration for area lights)
 - sometimes fog
- (Whitted-style) ray tracing relaxes that, simulating:
 - specular reflection
 - shadows
 - transparent surfaces (transmission with refraction)
 - sometimes indirect illumination (a.k.a. global illumination)
 - sometimes area light sources
 - sometimes fog

Ray Casting

Let's start from: Ray Casting



Ray Casting

- A very flexible visibility algorithm

- loop y, loop x

- shoot ray from eye point through pixel (x,y) into scene
- intersect with all surfaces, find first one the ray hits
- shade that surface point to compute pixel (x,y)'s color

```
Raycast()                                // generate a picture
    for each pixel x,y
        color(pixel) = Trace(ray_through_pixel(x,y))

Trace(ray)                                // fire a ray, return RGB radiance
                                          // of light traveling backward along it
    object_point = Closest_intersection(ray)
    if object_point return Shade(object_point, ray)
    else return Background_Color


Closest_intersection(ray)
    for each surface in scene
        calc_intersection(ray, surface)
    return the closest point of intersection to viewer
    (also return other info about that point, e.g., surface
    normal, material properties, etc.)

Shade(point, ray)                         // return radiance of light leaving
                                          // point in opposite of ray direction
    calculate surface normal vector
    check shadow map for light visibility
    use Phong illumination formula (or something similar)
    to calculate contributions of each light source
```


Ray Casting

- This can be easily generalized to give recursive ray tracing, that will be discussed later
- Can handle translucency (which rasterization cannot!)
- `calc_intersection (ray, surface)` is the most important operation
 - compute not only coordinates, but also geometric or appearance attributes at the intersection point


Recursive ray tracing



Recursive ray tracing




Recursive ray tracing



Recursive ray tracing


Until Later!



Ray-Surface Intersection

Ray Equation



- How to represent a ray?
 - A ray is $\mathbf{p} + t\mathbf{d}$: \mathbf{p} is ray origin, \mathbf{d} the direction
 - $t = 0$ at origin of ray, $t > 0$ in positive direction of ray
 - typically assume $\|\mathbf{d}\| = 1$
 - \mathbf{p} and \mathbf{d} are typically computed in world space



Ray-Surface Intersections


- How to represent a ray?
 - A ray is $\mathbf{p} + t\mathbf{d}$: \mathbf{p} is ray origin, \mathbf{d} the direction
 - $t = 0$ at origin of ray, $t > 0$ in positive direction of ray
 - typically assume $\|\mathbf{d}\| = 1$
 - \mathbf{p} and \mathbf{d} are typically computed in world space
- Recap: how to represent a surface?
 - Implicit functions: $f(\mathbf{x}) = 0$
 - Parametric functions: $\mathbf{x} = \mathbf{g}(u, v)$

Solve the \mathbf{x} and t for $\mathbf{x} = \mathbf{p} + t\mathbf{d}$
 $f(\mathbf{x}) = 0$



Parametric

$$\begin{aligned} \mathbf{x}(\mathbf{u}) &= r \cos(\mathbf{u}) \\ \mathbf{y}(\mathbf{u}) &= r \sin(\mathbf{u}) \end{aligned}$$



Implicit


$$F(\mathbf{x}, \mathbf{y}) = \mathbf{x}^2 + \mathbf{y}^2 - r^2$$

Ray-Surface Intersections

- Compute Intersections:
 - Substitute ray equation for $x = \mathbf{p} + t\mathbf{d}$
 - Find roots
 - Implicit: $f(\mathbf{p} + t\mathbf{d}) = 0$
 - one equation in one unknown – univariate root finding
 - Parametric: $\mathbf{p} + t\mathbf{d} - \mathbf{g}(u, v) = 0$
 - three equations in three unknowns (t,u,v) – multivariate root finding
 - For univariate polynomials, use closed form solution; otherwise, use numerical root finder

Ray-Sphere Intersection

- Ray-sphere intersection is an easy case
- A sphere's implicit function is: $x^2 + y^2 + z^2 - r^2 = 0$ if sphere at origin
- The ray equation is:
$$\begin{aligned}x &= p_x + td_x \\y &= p_y + td_y \\z &= p_z + td_z\end{aligned}$$
- Substitution gives: $(p_x + td_x)^2 + (p_y + td_y)^2 + (p_z + td_z)^2 - r^2 = 0$
- A quadratic equation in t .
- Solve the standard way: $A = d_x^2 + d_y^2 + d_z^2 = 1$ (unit vector)
$$\begin{aligned}B &= 2(p_x d_x + p_y d_y + p_z d_z) \\C &= p_x^2 + p_y^2 + p_z^2 - r^2\end{aligned}$$
- Quadratic formula has two roots: $t = (-B \pm \sqrt{B^2 - 4C})/2$
 - which correspond to the two intersection points
 - We take the smaller t (the first intersection)
 - negative discriminant means ray misses sphere




Ray Intersection With Triangle

Triangle is in a plane

- Ray-plane intersection
- Test if hit point is inside triangle

Many ways to optimize...



$$(\mathbf{p} - \mathbf{p}') \cdot \mathbf{N} = 0$$

↑
all points on plane

↑
one point
on plane

↑
normal vector

$$ax + by + cz + d = 0$$

Ray Intersection With Plane

Ray equation:

$$\mathbf{r}(t) = \mathbf{o} + t \mathbf{d}, \quad 0 \leq t < \infty$$

Plane equation:

$$\mathbf{p} : (\mathbf{p} - \mathbf{p}') \cdot \mathbf{N} = 0$$


Solve for intersection

Set $\mathbf{p} = \mathbf{r}(t)$ and solve for t

$$(\mathbf{p} - \mathbf{p}') \cdot \mathbf{N} = (\mathbf{o} + t \mathbf{d} - \mathbf{p}') \cdot \mathbf{N} = 0$$

$$t = \frac{(\mathbf{p}' - \mathbf{o}) \cdot \mathbf{N}}{\mathbf{d} \cdot \mathbf{N}}$$

Check: $0 \leq t < \infty$



Ray Intersection With Plane

Ray equation:

$$\mathbf{r}(t) = \mathbf{o} + t \mathbf{d}, \quad 0 \leq t < \infty$$

Plane equation:

$$\mathbf{p} : (\mathbf{p} - \mathbf{p}') \cdot \mathbf{N} = 0$$



Solve for intersection

Set $\mathbf{p} = \mathbf{r}(t)$ and solve for t

$$(\mathbf{p} - \mathbf{p}') \cdot \mathbf{N} = (\mathbf{o} + t \mathbf{d} - \mathbf{p}') \cdot \mathbf{N} = 0$$

$$t = \frac{(\mathbf{p}' - \mathbf{o}) \cdot \mathbf{N}}{\mathbf{d} \cdot \mathbf{N}}$$

Check: $0 \leq t < \infty$



If $t > 0$ and $\mathbf{x}(t)$ *inside*:
return Intersection point, $\mathbf{x}(t)$

Barycentric Coordinates


$$\mathbf{p} = b_0 \mathbf{x}_0 + b_1 \mathbf{x}_1 + b_2 \mathbf{x}_2$$

$$\left. \begin{aligned} b_0 &= A_0/A \\ b_1 &= A_1/A \\ b_2 &= A_2/A \end{aligned} \right\} b_0 + b_1 + b_2 = 1$$

$$A_0 = \frac{1}{2}(\mathbf{x}_1 - \mathbf{p}) \times (\mathbf{x}_2 - \mathbf{p}) \cdot \mathbf{n}$$

$$A_1 = \frac{1}{2}(\mathbf{x}_2 - \mathbf{p}) \times (\mathbf{x}_0 - \mathbf{p}) \cdot \mathbf{n}$$

$$A_2 = \frac{1}{2}(\mathbf{x}_0 - \mathbf{p}) \times (\mathbf{x}_1 - \mathbf{p}) \cdot \mathbf{n}$$




Inside: $0 < b_i < 1$ ($i = 0, 1, 2$), and coplanar

Outside: otherwise


Ray-Polygon Intersection

- Assuming we have a planar polygon
 - first, find intersection point of ray with plane
 - then check if that point is inside the polygon
- Latter step is a point-in-polygon test in 3-D:
 - inputs: a point x in 3-D and the vertices of a polygon in 3D
 - output: INSIDE or OUTSIDE
 - problem can be reduced to point-in-polygon test in 2-D (**how?**)
- Point-in-polygon test in 2-D:
 - easiest for triangles
 - easy for convex n -gons
 - harder for concave polygons
 - most common approach: subdivide all polygons into triangles
 - for optimization tips, see article by Haines in the book [Graphics Gems IV](#)




Whitted-Style Ray Tracing

Whitted-Style Ray Tracing




Ray Types

- We'll distinguish four ray types:
 - Eye rays: originating at the eye
 - Shadow rays: from surface point toward light source
 - Reflection rays: from surface point in mirror direction
 - Transmission rays: from surface point in refracted direction



Ray Tracing Algorithm


1. send ray from eye through each pixel
2. compute point of **closest intersection** with a scene surface
3. shade that point by computing shadow rays
4. **spawn reflected and refracted rays**, repeat 2-4 steps



Specular Reflection Rays


- An eye ray hits **a shiny surface**
 - We know the direction from which a specular reflection would come, based on the surface normal
 - Fire a ray in this reflected direction
 - The reflected ray is treated just like an eye ray: it hits surfaces and spawns new rays
 - Light flows in the direction opposite to the rays (towards the eye), is used to calculate shading
 - It's easy to calculate the reflected ray direction

Note: arrowheads show the direction in which we're *tracing the rays*, not the direction the light travels.



Specular Transmission Rays


- To add **transparency**:
 - Add a term for light that's coming from within the object
 - These rays are refracted (bent) when passing through a boundary between two media with different refractive indices
 - When a ray hits a transparent surface fire a *transmission ray* into the object at the proper refracted angle
 - If the ray passes through the other side of the object then it bends again (the other way)



Refraction

- Refraction:
 - The bending of light due to its different velocities through different materials
 - rays bend toward the normal when going from sparser to denser materials (e.g. air to water), away from normal in opposite case
- Refractive index:
 - Light travels at speed c/n in a material of refractive index n
 - c is the speed of light in a vacuum
 - c varies with wavelength, hence rainbows and prisms
 - Use **Snell's law** $n_1 \sin \theta_1 = n_2 \sin \theta_2$ to derive refracted ray direction
 - note: ray dir. can be computed without trig functions (only sqrts)

MATERIAL	INDEX OF REFRACTION
air/vacuum	1
water	1.33
glass	about 1.5
diamond	2.4



From a Ray Caster to a Ray Tracer

```
Trace(ray)                // fire a ray, return RGB radiance
                           // of light traveling backward along it
object_point = Closest_intersection(ray)
if object_point return Shade(object_point, ray)
else return Background_Color
```

```
Shade(point, ray)          /* return radiance along ray */
radiance = black;          /* initialize color vector */
for each light source
    shadow_ray = calc_shadow_ray(point, light)
    if !in_shadow(shadow_ray, light)
        radiance += phong_illumination(point, ray, light)
```


```
if material is specularly reflective
    radiance += spec_reflectance *
        Trace(reflected_ray(point, ray))
if material is specularly transmissive
    radiance += spec_transmittance *
        Trace(refracted_ray(point, ray))
return radiance
```

```
ray eye_ray;
eye_ray.level = 0;


reflected_ray(ray in):
    ray out;
    out.level = in.level++
    return out
```

Ray Casting vs. Ray Tracing


```
Trace(ray) // fire a ray, return RGB radiance of light traveling backward along it
    if ray.level > n return Background_Color;
    object_point = Closest_intersection(ray)
    if object_point return Shade(object_point, ray)
    else return Background_Color
```



Ray Casting -- 1 bounce




Ray Tracing -- 2 bounce



Ray Tracing -- 3 bounce

Problem with Simple Ray Tracing





Aliasing

- Ray tracing shoots one ray per pixel
- But a pixel represents an area; one ray samples only one point with the area; an area consists *infinite* number of points
 - These points may not all have the same color
 - This leads to *aliasing*
 - jaggies
 - moire patterns
- How do we fix this problem?
 - Recall antialiasing we talked earlier

Antialiasing: Supersampling

- We talked about two antialiasing methods
 - Supersampling
 - Pre-filtering (MIP-mapping)
- Here we use **supersampling**
 - Fire more than one ray for each pixel (e.g., a 3x3 grid of rays)
 - Average the results using a filter (or some kind of filter)




Antialiasing: Adaptive Supersampling

- Supersampling can be done **adaptively**
 - divide pixel into 2x2 grid, trace 5 rays (4 at corners, 1 at center)
 - if the colors are **similar** then just use their average
 - otherwise recursively subdivide each cell of grid
 - keep going until each 2x2 grid is close to uniform or limit is reached
 - filter the result
- Behavior of adaptive supersampling
 - Areas with fairly constant appearance are sparsely sampled
 - Areas with lots of variability are heavily sampled

Motion Blur

- Apply stochastic sampling to time as well as space
- Assign a time as well as an image position to each ray
- The result is still-frame motion blur and smooth animation
- This is an example of **distribution ray tracing**



Motion Blur: a classic example


- From Foley et. al. Plate III.16
- Rendered using distribution ray tracing at 4096x3550 pixels, 16 samples per pixel.
- Note motion-blurred reflections and shadows with penumbrae cast by extended light sources.




Ray Tracing Acceleration

Whitted-Style Ray Tracing


Ray-Surface Intersection!



Ray-Surface Intersection



$$\bullet \mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$$




$$\text{Solve } (\mathbf{r}(t) - \mathbf{c})^2 = R^2$$

$$a = \mathbf{d} \cdot \mathbf{d}$$

$$b = 2(\mathbf{o} - \mathbf{c}) \cdot \mathbf{d}$$

$$c = (\mathbf{o} - \mathbf{c}) \cdot (\mathbf{o} - \mathbf{c}) - R^2$$

$$t = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$



$$\text{Solve } (\mathbf{r}(t) - \mathbf{x}_0) \cdot (\mathbf{x}_{10} \times \mathbf{x}_{20}) = 0$$

$$t = \frac{\mathbf{x}_{00} \cdot \mathbf{x}_{10} \times \mathbf{x}_{20}}{\mathbf{d} \cdot \mathbf{x}_{10} \times \mathbf{x}_{20}}$$

If $t > 0$ and $\mathbf{x}(t)$ inside:
return Intersection point, $\mathbf{x}(t)$

Ray Tracing – Performance Challenges

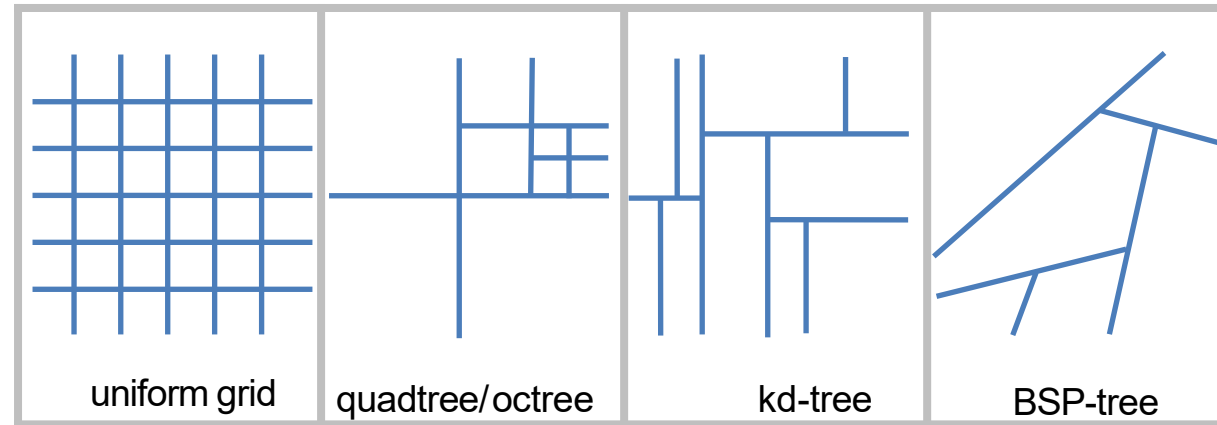


Jun Yan, Tracy Renderer

Ray Tracing – Performance Challenges

- Checking intersections with everything!

Spatial partitioning




- Checking intersections with complex geometry!

Bounding Volumes




Grid Acceleration

- 1. Find bounding box
- 2. Create grid
- 3. Store each object in overlapping cells




Grid Acceleration

- Step through grid in ray traversal order
- For each grid cell :
 - Test intersection with all objects stored at that cell




Grid Resolution?

- One cell
 - No speedup



Grid Resolution?


- Too many cells
 - Inefficiency due to extraneous grid traversal




Ray Tracing – Grid Resolution?




Spatial Partitioning



Oct-Tree


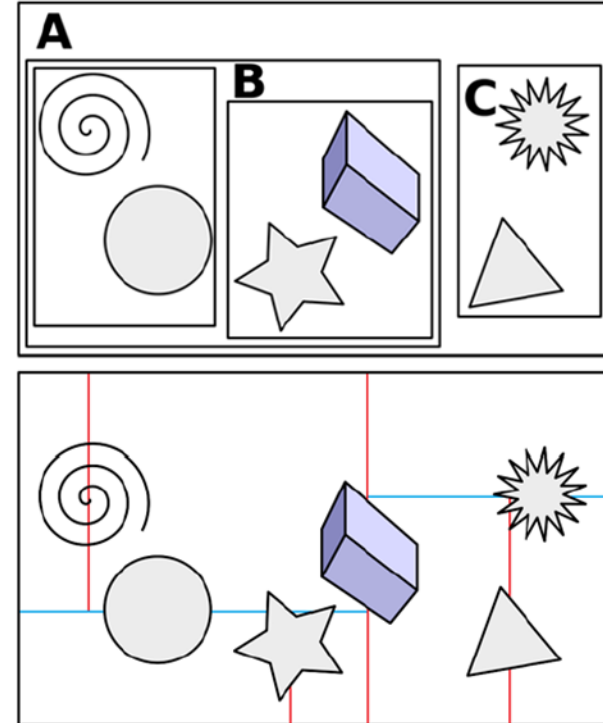



KD-Tree



BSP-Tree

Spatial Partitioning




General task:

- 1. Build the tree
- 2. For a given point, travel the root-to-leaf path and test intersections

Octrees

- Quadtree is the 2-D generalization of binary tree
 - node (cell) is a square
 - recursively split into four equal sub-squares
 - stop when leaves get “simple enough”



Octrees

- Octree is the 3-D generalization of quadtree
 - node (cell) is a cube, recursively split into eight equal sub-cubes
 - for ray tracing:
 - stop splitting when the number of objects intersecting the cell gets “small enough” or the tree depth exceeds a limit
 - internal nodes store pointers to children, leaves store list of surfaces
 - more expensive to traverse than a grid
 - but an octree adapts to nonhomogeneous, clumpy scenes better


```
trace(cell, ray) {           // returns object hit or NONE
    if cell is leaf, return closest (objects_in_cell(cell))
    for child cells pierced by ray, in order      // 1 to 4 of these
        obj = trace(child, ray)
        if obj!=NONE return obj
    return NONE
}
```

Which Data Structure is Best for Ray Tracing?

- Grids are easy to implement, but they're memory hogs (and slow) for nonhomogeneous scenes, i.e. most scenes
- Octrees are pretty good, but not as fast as grids for some scenes
- Nested grids seem to be the fastest on **static** scenes
- If scene is dynamic, the cost of regenerating or updating the data structure may become an issue
- In such cases, hierarchical bounding volumes may be best
- Hierarchical bounding volumes easy to implement if your model is naturally hierarchical (e.g. human), otherwise not
- For other visibility algorithms:
 - BSP trees useful for Painter's algorithm...

k-d Trees


- Relax the rules for quadtrees and octrees:




- first variant: *k-dimensional (k-d) tree*
 - don't always split at midpoint
 - split only one dimension at a time (i.e. x or y or z)
 - useful for clustering and choosing colormaps for color image quantization

BSP Trees

- Relax the rules for quadtrees and octrees:
- second variant: *binary space partitioning (BSP) tree*
 - permit splits with any line
 - in general, split k dimensional space with $k-1$ dimensional hyperplane
 - 2-D space split with lines (most of our examples)
 - 3-D space split with planes
 - each node corresponds to a (potentially unbounded) convex polyhedron
 - useful for Painter's algorithm



Building a BSP Tree



Building a Good Tree - the tricky part

- A naïve partitioning of n polygons will yield $O(n^3)$ polygons!
- Algorithms exist to find partitionings that produce $O(n^2)$.
 - For example, try all remaining polygons and add the one which causes the fewest splits (greedy algorithm!)
 - Fewer splits -> larger polygons -> better polygon fill efficiency
- Also, we want a balanced tree.
 - More important for ray casting than scan conversion.
- These goals conflict.
- *note: in the examples we've shown, the geometric objects being stored are planar, and we split using the planes of these objects, but that needn't be so – could theoretically split with any plane*

Uses for Binary Space Partitioning (BSP) Trees

- Painter's algorithm rendering
 - good for
 - static 3-D scenes with moving viewpoint (flight simulators)
 - architectural scenes with a small number of polygons (DOOM)
 - if you don't have z-buffer hardware
 - Add a few monsters and such after the environment is drawn
- Ray tracing
- Solid modeling with polyhedra
- History:
 - BSP trees first used by Naylor, Fuchs, et al. for Painter's algorithm ~1980
 - theoreticians scoffed at their worst-case performance
 - considered unpromising
 - revived by John Carmack, author of Quake, and the PC game community
 - out of necessity: no z-buffer hardware for PC's at the time